

Semantics and invariance proof method for weakly consistent parallelism

Jade Alglave (MSR-Cambridge, UCL, UK)
Patrick Cousot (NYU, Emer. ENS, PSL)

Dagstuhl Seminar 16471

<http://www.dagstuhl.de/16471>

Concurrency with Weak Memory Models: Semantics, Languages,
Compilation, Verification, Static Analysis, and Synthesis
November 20 – 25 , 2016

Weakly consistent parallel programs

Weakly consistent parallel programs

```
var x1,...,xm; // shared variables  
P0; // prelude initializing x1,...,xm  
[P1 || P2 || ... || Pn]
```

- P₁, P₂, ..., P_n are the **processes** modifying the shared variables and their **local registers R, ...**
- The execution of a **write x := E** to a shared variable and the **read R := x of a shared variable is not instantaneous** (as in sequential consistency)

Example (lb, load buffer)

- **Algorithm A:**

```
0:{ x = 0; y = 0; }  
P0          ||| P1  
1:r[] r1 x  ||| 11:r[] r2 y;  
2:w[] y 1    ||| 12:w[] x 1 ;  
3:           ||| 13:      ;
```

- **Specification S_{inv} :**

$$\text{at 3} \wedge \text{at 13} \Rightarrow \neg(r1=1 \wedge r2=1)$$

Example (Peterson)

- **Algorithm A:**

```
0:{ w F1 false; w F2 false; w T 0; }

P0:
1:w[] F1 true
2:w[] T 2
3:do {i}
4:    r[] R1 F2
5:    r[] R2 T
6:while R1 ∧ R2 ≠ 1
7:skip (* CS1 *)
8:w[] F1 false
9:

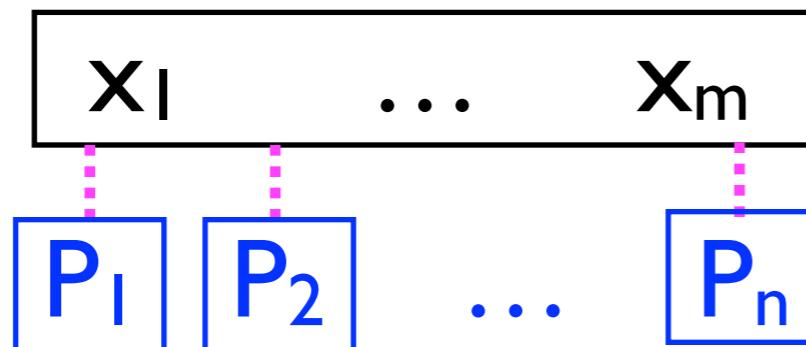
P1:
10:w[] F2 true;
11:w[] T 1;
12:do {j}
13:    r[] R3 F1;
14:    r[] R4 T;
15:while R3 ∧ R4 ≠ 2;
16:skip (* CS2 *)
17:w[] F2 false;
18:
```

- **Specification S_{inv} :**

1: {true}	10: {true}
...	...
7: {¬at{16}}	16: {¬at{7}}
...	...
9: {true}	18: {true}

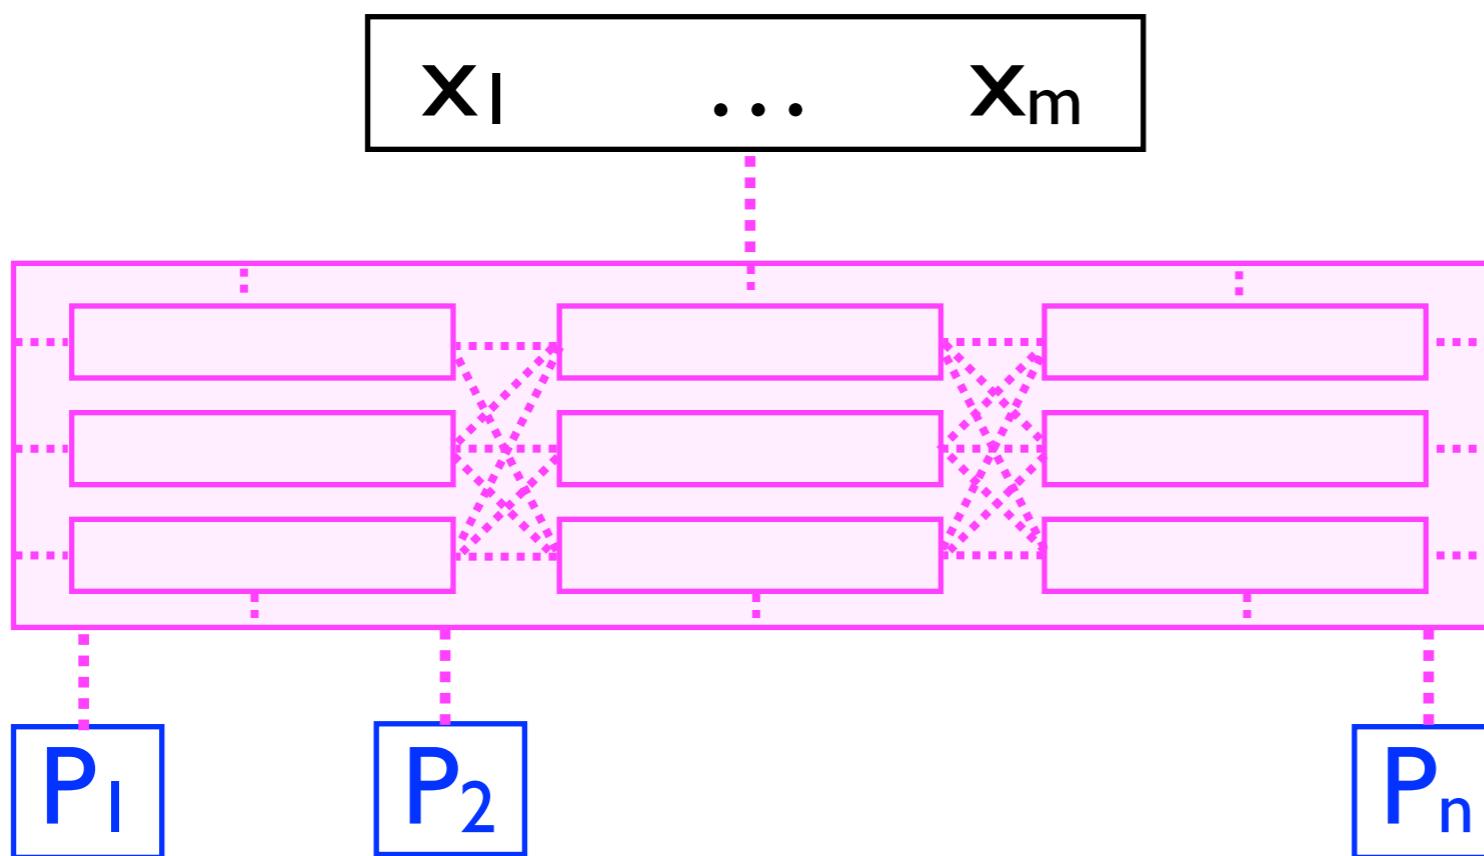
Weak memory/consistency models

- Sequential consistency:



atomic
instantaneous
communications

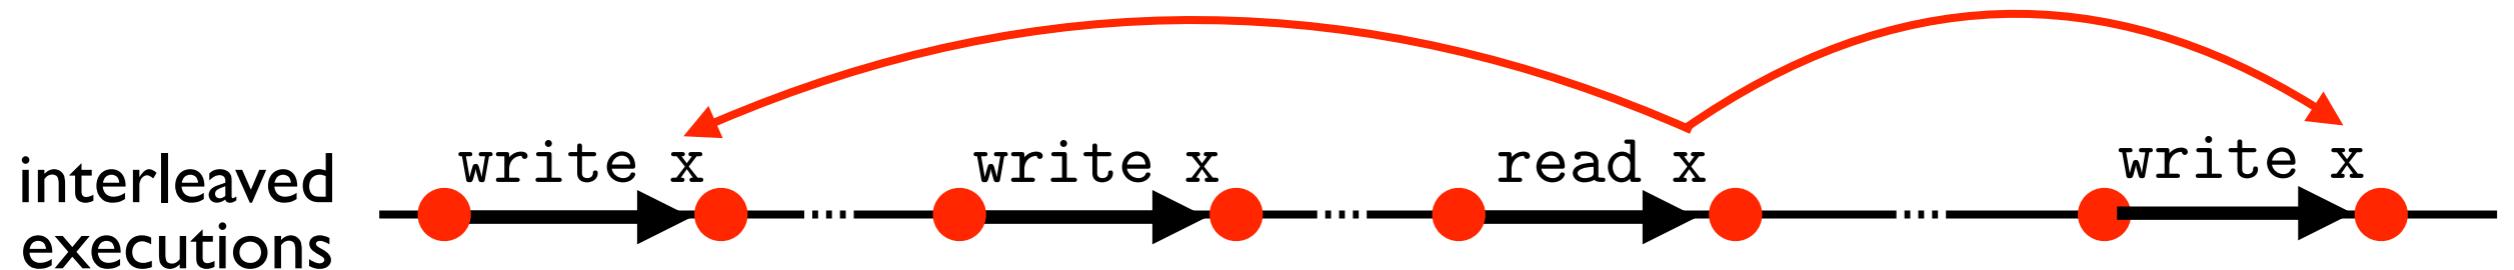
- Weak memory models:



communication
network
(anticipations,
delays, shuffles...)

Read/write matching

- In the worst case a read x can read from any past or future write x of any process (including for the reading process)



Example (lb, incorrect)

- ```
0:{ x = 0; y = 0; }
P0
1:r[] r1 x
2:w[] y 1
3:
```

```
P1
 ;
11:r[] r2 y;
12:w[] x 1 ;
13: ;
```
- at 3  $\wedge$  at 13  $\wedge$  r1=1  $\wedge$  r2=1
- This erroneous behavior can be observed on TSO machines

# Example: Peterson (incorrect)

- Can read the wrong flags

|                                      |                                   |
|--------------------------------------|-----------------------------------|
| 0:{ w F1 false; w F2 false; w T 0; } |                                   |
| P0:                                  |                                   |
| 1:w[] F1 true                        |                                   |
| 2:w[] T 2                            |                                   |
| 3:do                                 |                                   |
| 4: r[] R1 F2                         |                                   |
| 5: r[] R2 T                          |                                   |
| 6:while R1 $\wedge$ R2 $\neq$ 1      |                                   |
| 7:skip (* CS1 *)                     |                                   |
| 8:w[] F1 false                       |                                   |
| 9:                                   |                                   |
|                                      | P1:                               |
|                                      | 10:w[] F2 true;                   |
|                                      | 11:w[] T 1;                       |
|                                      | 12:do                             |
|                                      | 13: r[] R3 F1;                    |
|                                      | 14: r[] R4 T;                     |
|                                      | 15:while R3 $\wedge$ R4 $\neq$ 2; |
|                                      | 16:skip (* CS2 *)                 |
|                                      | 17:w[] F2 false;                  |
|                                      | 18:                               |

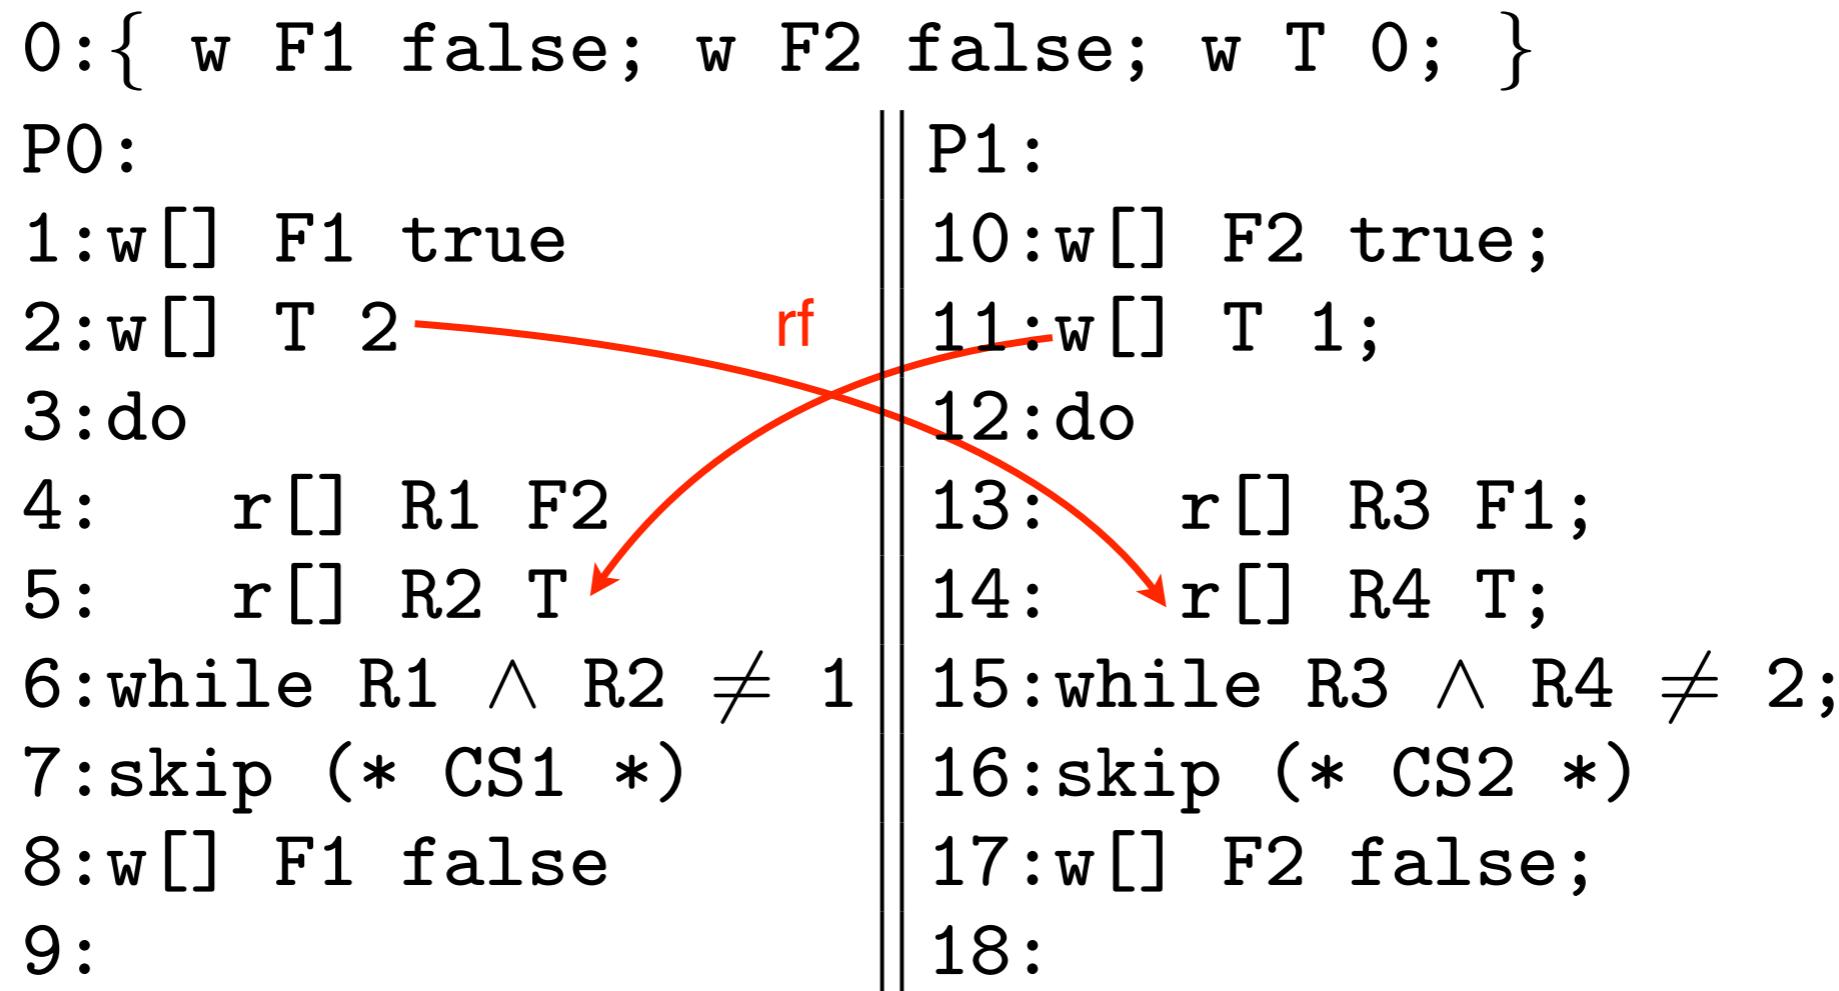
Red annotations: A red arrow points from line 4 to line 13, labeled 'rf' (Race Flag). Another red arrow points from line 2 to line 11.

at 6  $\wedge$  at 16:  $\neg R1 \wedge R2=1 \wedge \neg R3 \wedge R4=2$  holds

$\Rightarrow$  both processes simultaneously enter their critical section

# Example: Peterson (incorrect)

- Can read the wrong turns



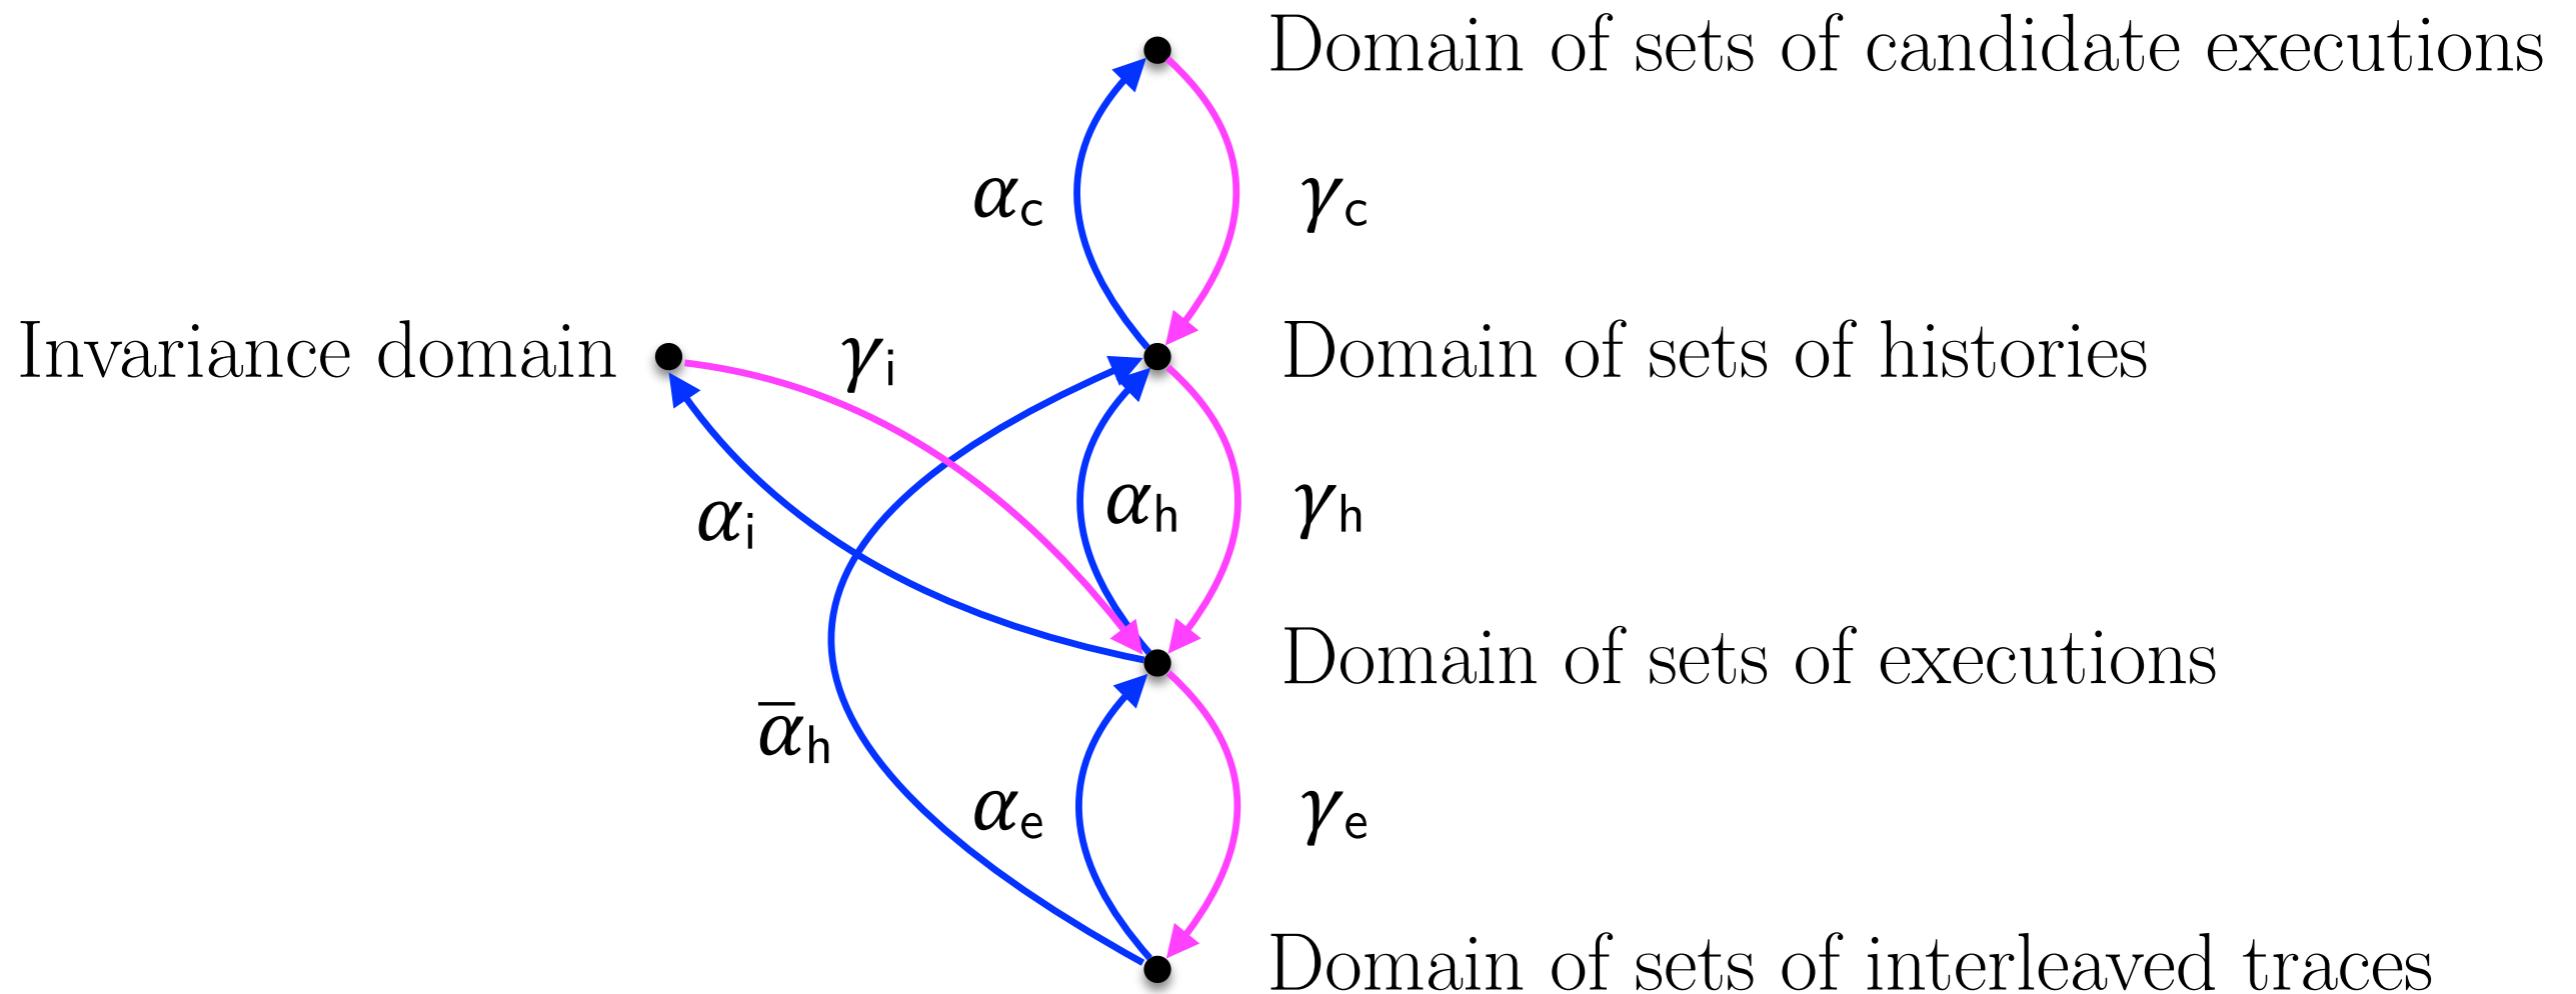
at 6  $\wedge$  at 16:  $\neg R1 \wedge R2=1 \wedge \neg R3 \wedge R4=2$  holds

$\Rightarrow$  both processes simultaneously enter their critical section

# A hierarchy of semantics of weakly consistent parallelism

# Hierarchy of semantics

- Hierarchy of semantic domains:



- Induces a hierarchy of semantics

# Sets of interleaved traces

- **Traces:** maximal finite or infinite sequence of states separated by events generated by computation and communication steps —→ *global time*
- **States:** shared memory assigning values to global variables, store buffers, ... program point of each process, assignment to local registers
- **Events**  $e$ :  $P(e)$ : process executed,  $A(e)$ : labelled action executed,  $X(e)$ : shared variable involved,  $V(e)$ : value involved, ...
- No restriction on who can read which write on the same shared variable!

# Example of interleaved trace for 1b

- 0:{ x = 0; y = 0; }

|            |              |
|------------|--------------|
| P0         | P1 ;         |
| 1:r[] r1 x | 11:r[] r2 y; |
| 2:w[] y 1  | 12:w[] x 1 ; |
| 3:         | 13: ;        |

- start  $\xrightarrow{0: \overbrace{x = 0; y = 0}^{w_x^0 \quad w_y^0}}$   $\langle \{x \leftarrow w_x^0, y \leftarrow w_y^0\}, 1:\{r1 \leftarrow 0\}, 11:\{r2 \leftarrow 0\} \rangle$   $\xrightarrow{1:r[] r1 x^{r_x^1}} \langle \{x \leftarrow w_x^{12}, y \leftarrow w_y^0\}, 2:\{r1 \leftarrow 1\}, 11:\{r2 \leftarrow 0\} \rangle$   $\xrightarrow{11:r[] r2 y^{r_y^{11}}} \langle \{x \leftarrow w_x^{12}, y \leftarrow w_y^2\}, 2:\{r1 \leftarrow 1\}, 12:\{r2 \leftarrow 1\} \rangle$   $\xrightarrow{12:w[] x 1^{w_x^{12}}} \langle \{x \leftarrow w_x^{12}, y \leftarrow w_y^2\}, 2:\{r1 \leftarrow 1\}, 13:\{r2 \leftarrow 1\} \rangle$

# Sets of truly parallel execution traces

- project **traces per process**  $\longrightarrow$  *local time on computations*
- get rid of shared memory states using a **read-from relation rf**  $\longrightarrow$  *no time on communications*

$$\langle r, w \rangle \in \text{rf} \iff \tau = \tau_0 \langle \nu, \dots \rangle \xrightarrow{r} \langle \nu', \dots \rangle \tau_1 \quad \wedge \quad \nu'(\mathbf{X}(r)) = w$$

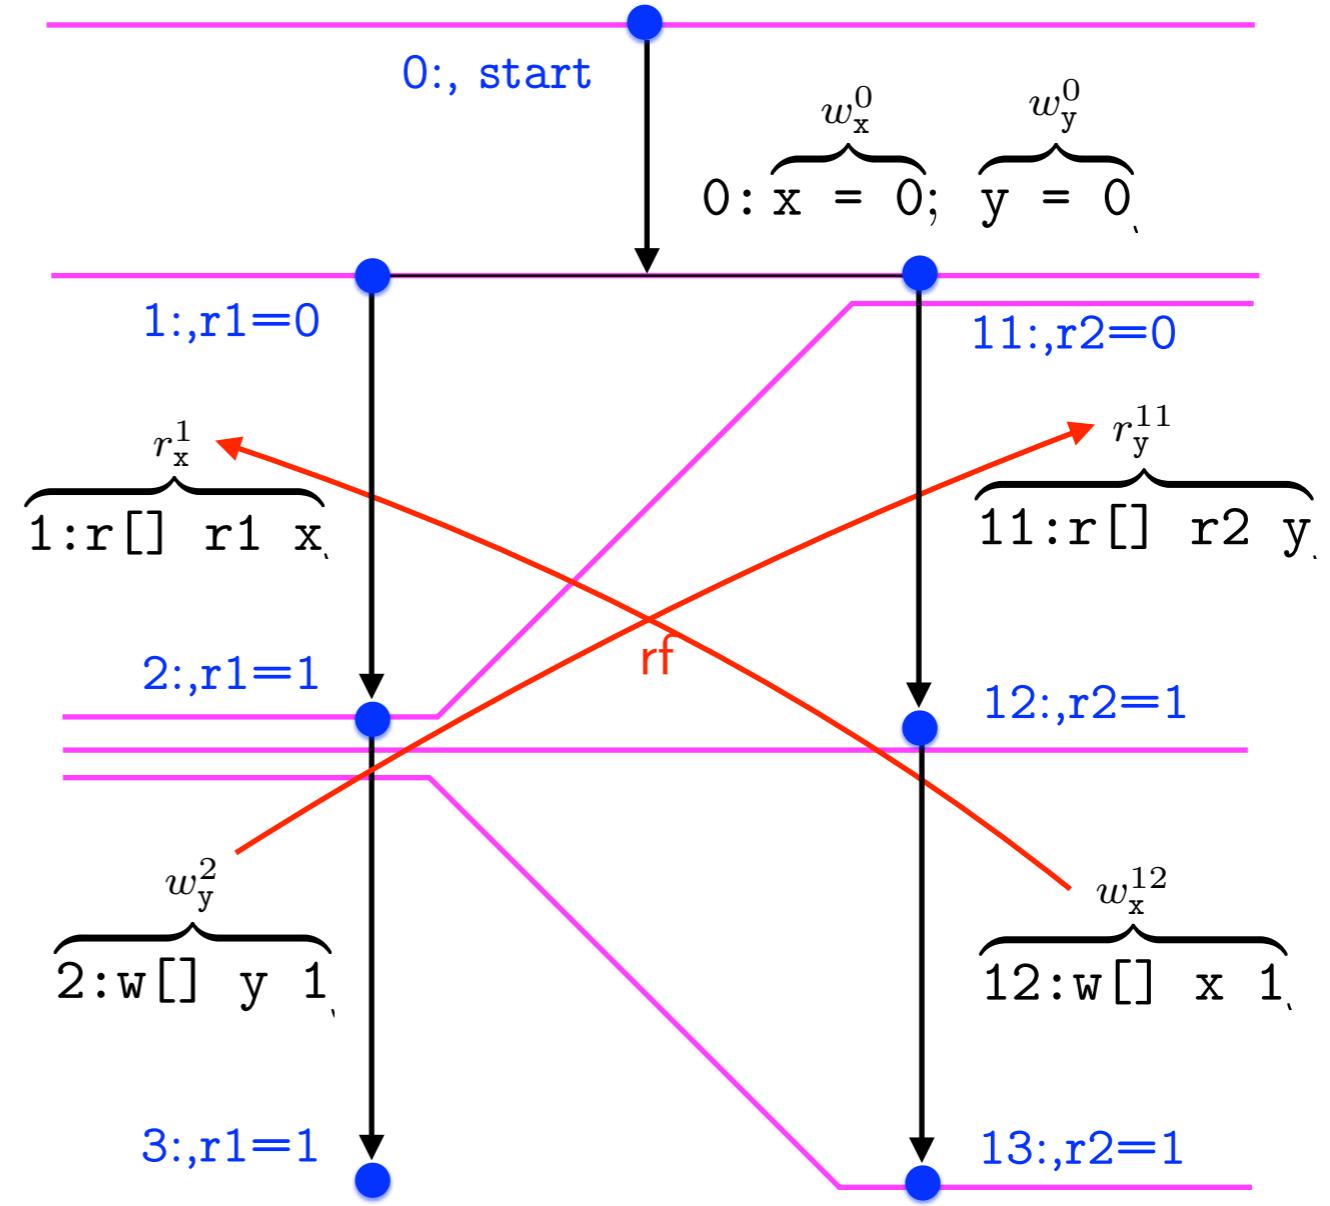
- keep **local states** on process control points and values of registers
- keep computation progress information using **cuts** of parallel traces  $\longrightarrow$  *global time*

# Example of truly parallel execution for lb

```

0:{ x = 0; y = 0; }
P0 P1
1:r[] r1 x || 11:r[] r2 y;
2:w[] y 1 || 12:w[] x 1 ;
3: 13: ;

```



# Sets of histories

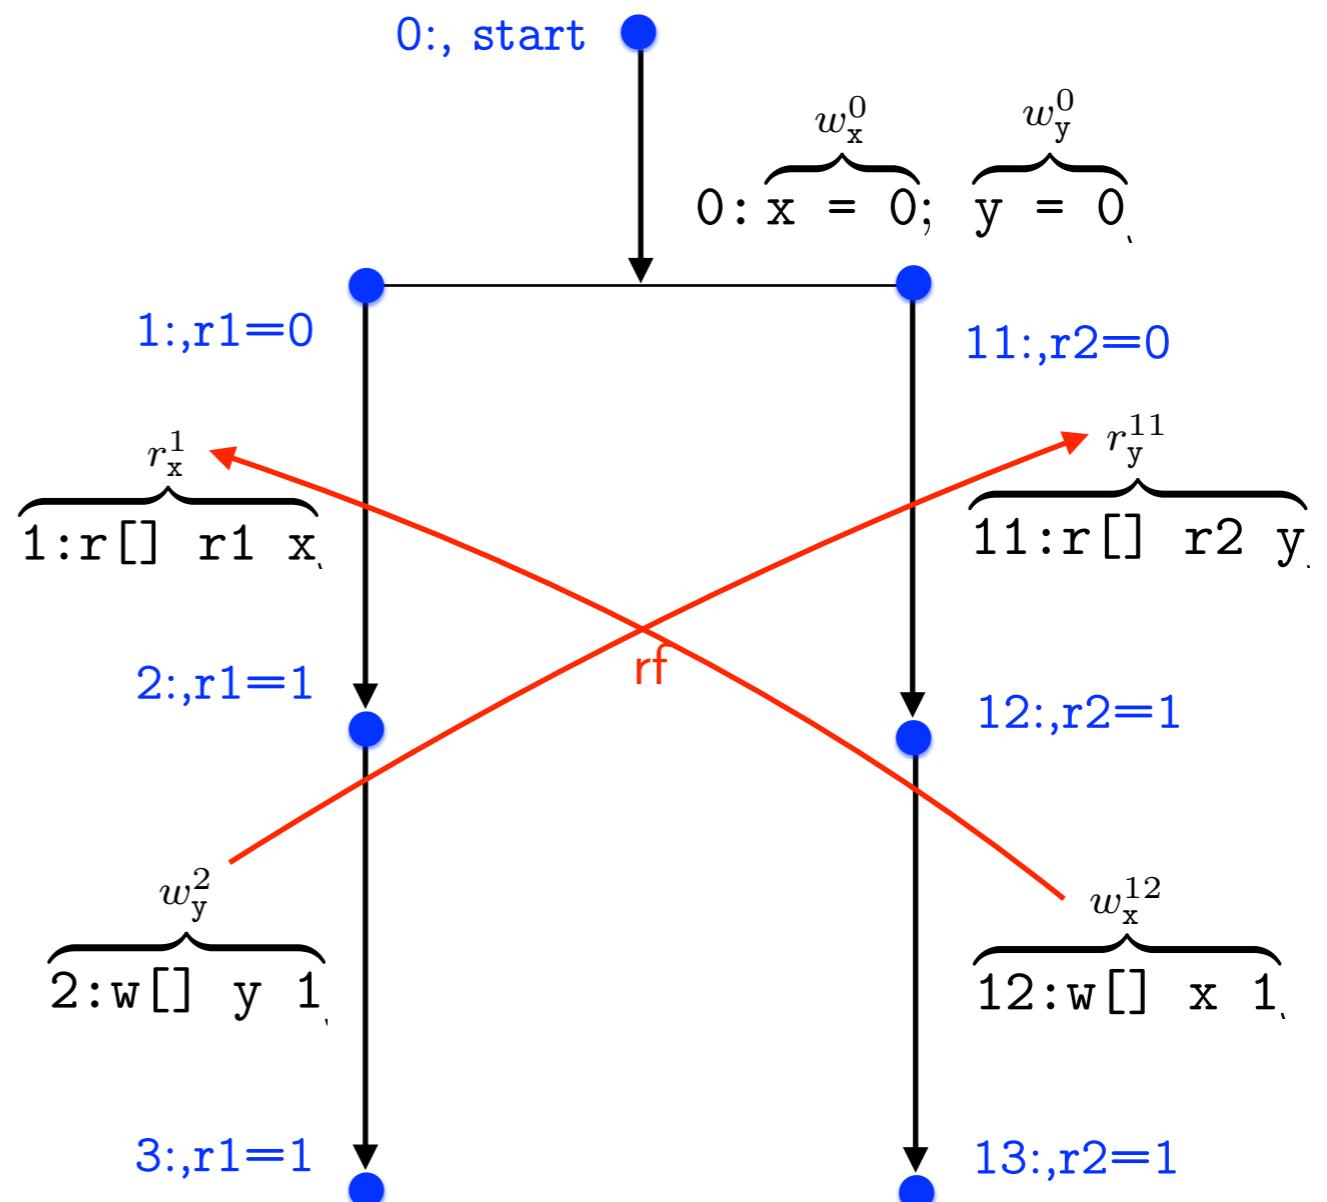
- Get rid of cuts —→ *no global time*
- A processor cannot know where the others parallel processors are in their computations

# Example of history for 1b

```

0:{ x = 0; y = 0; }
P0 P1
1:r[] r1 x || 11:r[] r2 y;
2:w[] y 1 || 12:w[] x 1 ;
3: || 13: ;

```



# Sets of candidate executions

- Keep the **set of events**
- Keep the **read-from relation rf**
- Represent process traces  $\tau_0 \prod_{i=1}^n \tau_i, \text{rf}$  by
  - the set of **initial writes IW** in  $\tau_0$
  - the **program order po**
$$\langle e, e' \rangle \in \text{po} \iff \tau_i = \tau_i' \xrightarrow{e} \tau_i'' \xrightarrow{e'} \tau_i'''$$

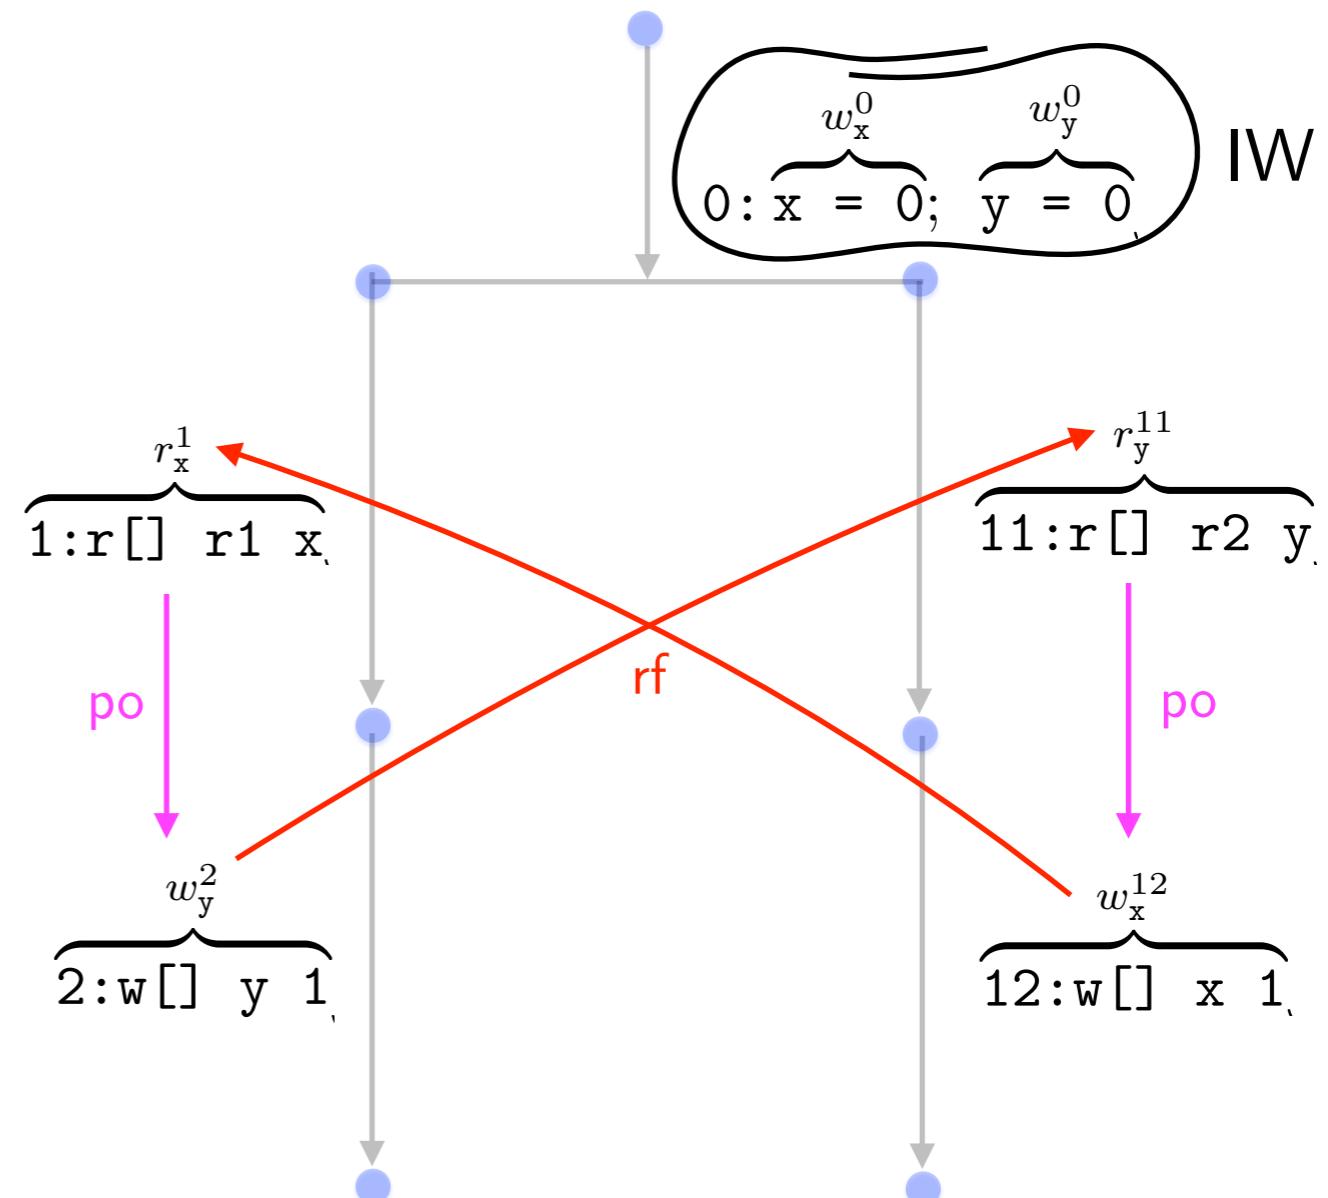
$\longrightarrow$  *relational on events*
- Get rid of states

# Example of candidate execution for 1b

```

0:{ x = 0; y = 0; }
P0 P1
1:r[] r1 x || 11:r[] r2 y;
2:w[] y 1 || 12:w[] x 1 ;
3: || 13: ;

```



# Auxiliary relations

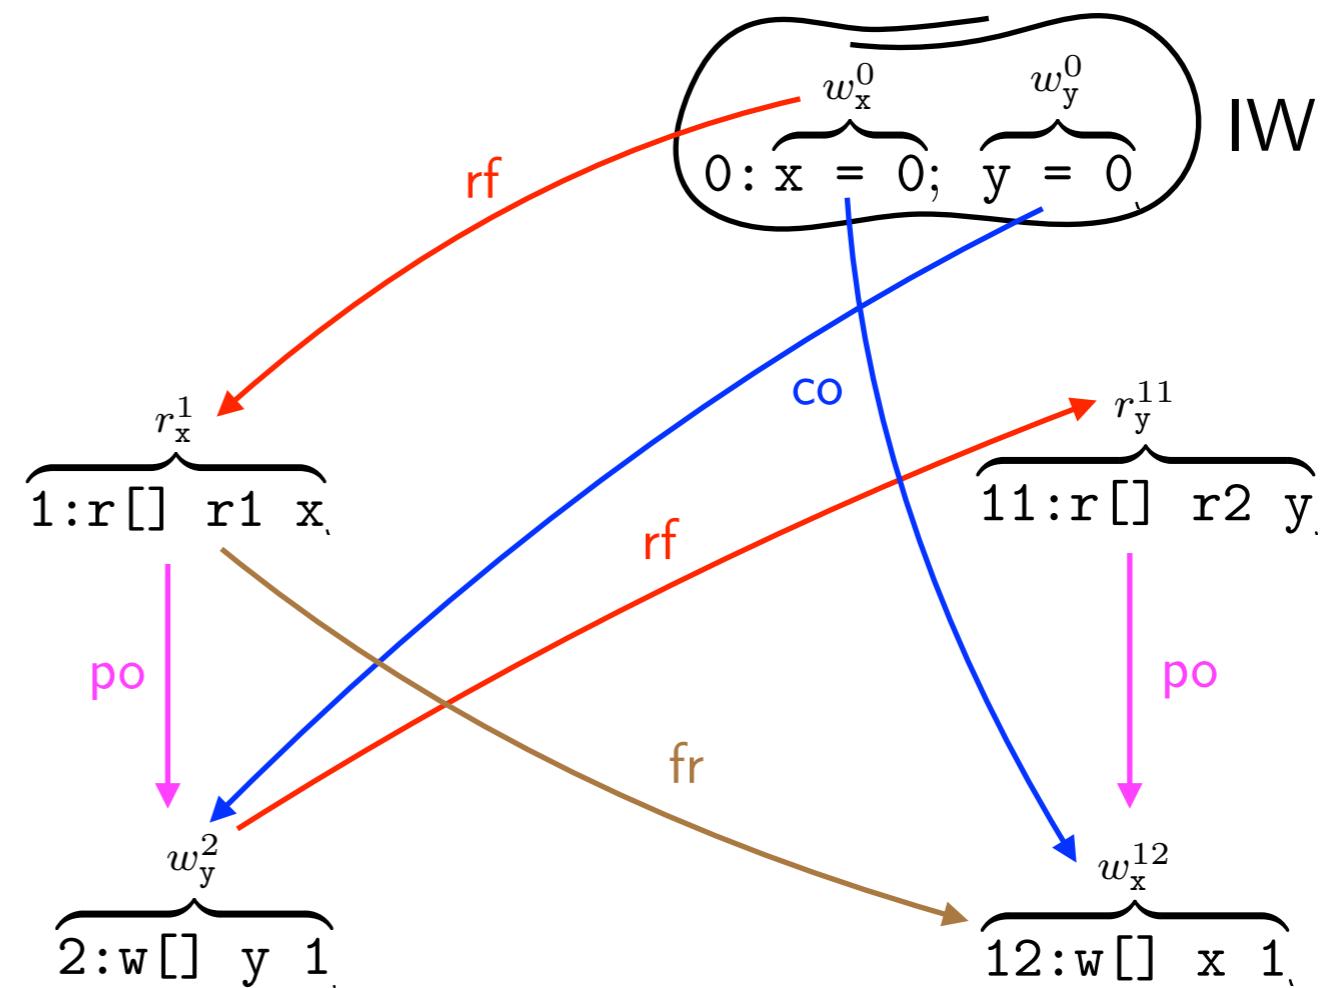
- **loc**: between events on the same shared variable
- **ext**: between events on different processes
- coherence order **co**: between a write and the later ones on the same shared variable
- from-read **fr**: between a read reading from a write and the later writes to the same shared variable  
$$\text{fr} = \text{rf}^{-1} ; \text{co}$$

# Auxiliary relations

```

0:{ x = 0; y = 0; }
P0 P1
1:r[] r1 x || 11:r[] r2 y;
2:w[] y 1 || 12:w[] x 1 ;
3: || 13: ;

```



# co in cat

"co.cat"

```
let fold f =
 let rec fold_rec (es,y) = match es with
 || {} -> y
 || e ++ es -> fold_rec (es, f(e,y))
 end in
 fold_rec

let map f = fun es -> fold (fun (e,y) -> f e ++ y) (es,{})

let rec cross S = match S with
 || {} -> { 0 }
 || S1 ++ S ->
 let yss = cross S in
 fold
 (fun (e1,r) -> map (fun t -> e1 | t) yss | r)
 (S1,{}) end
```

```
let co0 = loc & (IW * (W\IW))
let makeCo(s) = linearisations(s,co0)
let same-loc-writes = loc & (W*W)
let allCoL = map makeCo (classes (same-loc-writes))
let allCo = cross allCoL
```

with co from allCo

**Example of specification of  
weakly consistent parallelism  
in the semantic hierarchy:  
sequential consistency**

# Sequential consistency

- **Interleaved semantics:** a read can only read from the last past write

- $|b|$ :

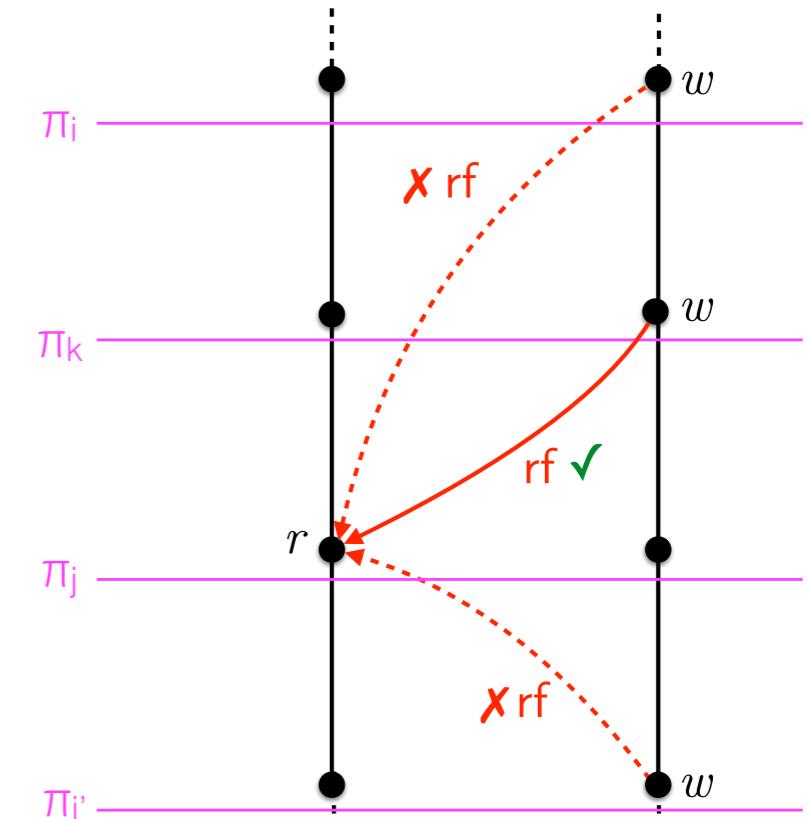
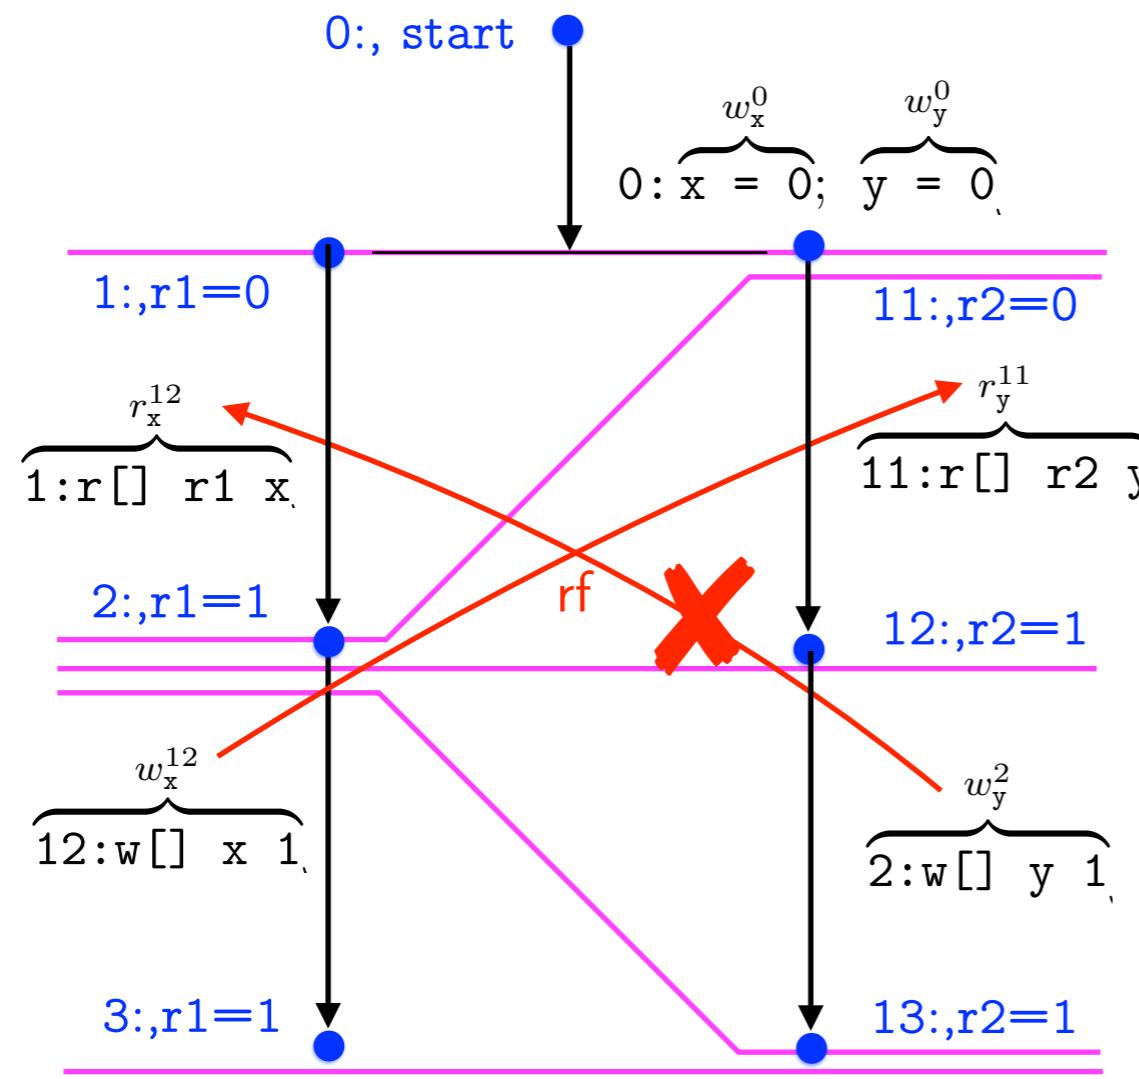
$$\begin{array}{l} \text{start} \xrightarrow[0: x = 0; y = 0]{w_x^0 \quad w_y^0} \langle \{x \leftarrow w_x^0, y \leftarrow w_y^0\}, 1:\{r1 \leftarrow 0\}, 11:\{r2 \leftarrow 0\} \rangle \xrightarrow[r_x^{12}]{1:r[] \quad r1 \quad x} \langle \{x \leftarrow w_x^{12}, y \leftarrow w_y^0\}, 2:\{r1 \leftarrow 1\}, 11:\{r2 \leftarrow 0\} \rangle \xrightarrow[r_y^{11}]{11:r[] \quad r2 \quad y} \\ \langle \{x \leftarrow w_x^{12}, y \leftarrow w_y^2\}, 2:\{r1 \leftarrow 1\}, 12:\{r2 \leftarrow 1\} \rangle \xrightarrow[w_x^{12}]{12:w[] \quad x \quad 1} \langle \{x \leftarrow w_x^{12}, y \leftarrow w_y^2\}, 2:\{r1 \leftarrow 1\}, 13:\{r2 \leftarrow 1\} \rangle \xrightarrow[w_y^2]{2:w[] \quad y \quad 1} \langle \{x \leftarrow w_x^{12}, y \leftarrow w_y^2\}, 3:\{r1 \leftarrow 1\}, 11:\{r2 \leftarrow 1\} \rangle \end{array}$$

The diagram illustrates a sequence of state transitions. It starts at a 'start' state with initial values  $w_x^0$  and  $w_y^0$ . A transition labeled  $0: x = 0; y = 0$  leads to a state where  $x \leftarrow w_x^0$  and  $y \leftarrow w_y^0$ , with local variables 1 and 11. This is followed by a read operation  $r_x^{12}$  on  $x$ , resulting in a state where  $x \leftarrow w_x^{12}$  and  $y \leftarrow w_y^0$ , with local variables 2 and 11. A subsequent write operation  $r_y^{11}$  on  $y$  leads to a state where  $x \leftarrow w_x^{12}$  and  $y \leftarrow w_y^2$ , with local variables 2 and 12. Finally, a read operation  $w_x^{12}$  on  $x$  leads to a state where  $x \leftarrow w_x^{12}$  and  $y \leftarrow w_y^2$ , with local variables 3 and 11. The values  $w_x^{12}$  and  $w_y^2$  are crossed out in red.

# Example: sequential consistency for lb

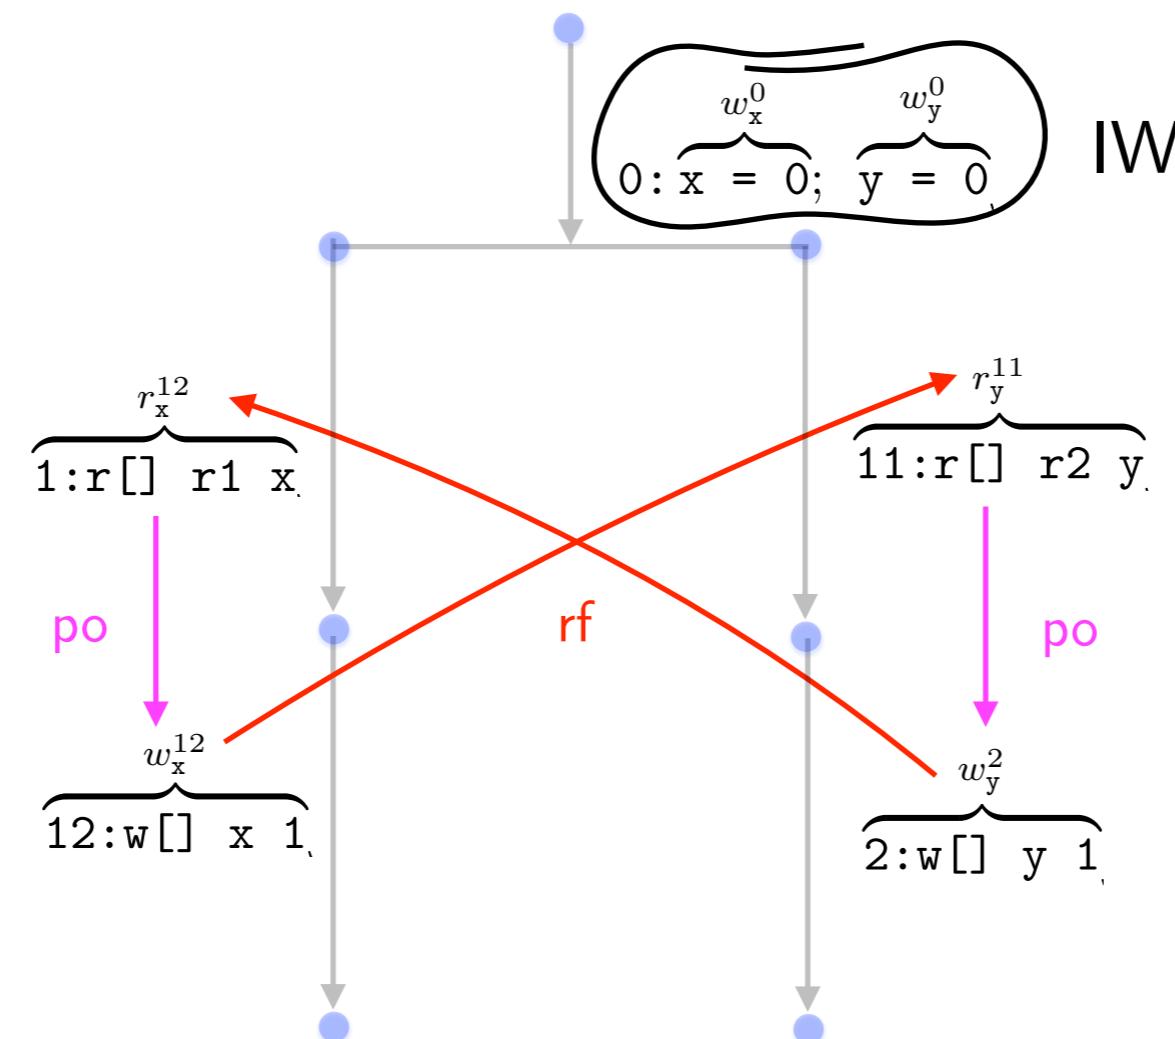
- Parallel executions with cuts: a read can read only the last write before its cut

- lb:



# Example: sequential consistency for lb

- **Parallel histories:** abstract to candidate execution and check it is allowed
- **Candidate executions:** irreflexive po ; rf ; po; rf



# Analytic semantics of weakly consistent parallelism

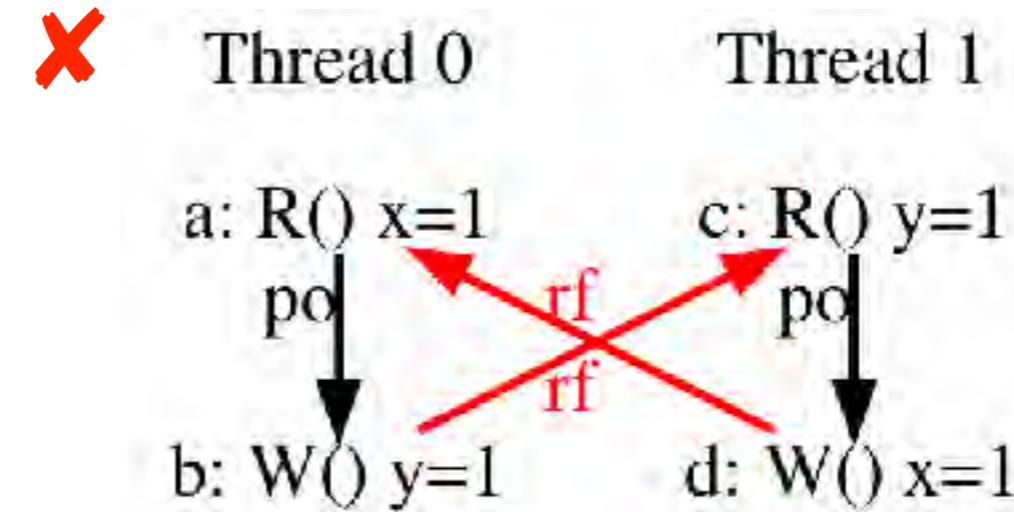
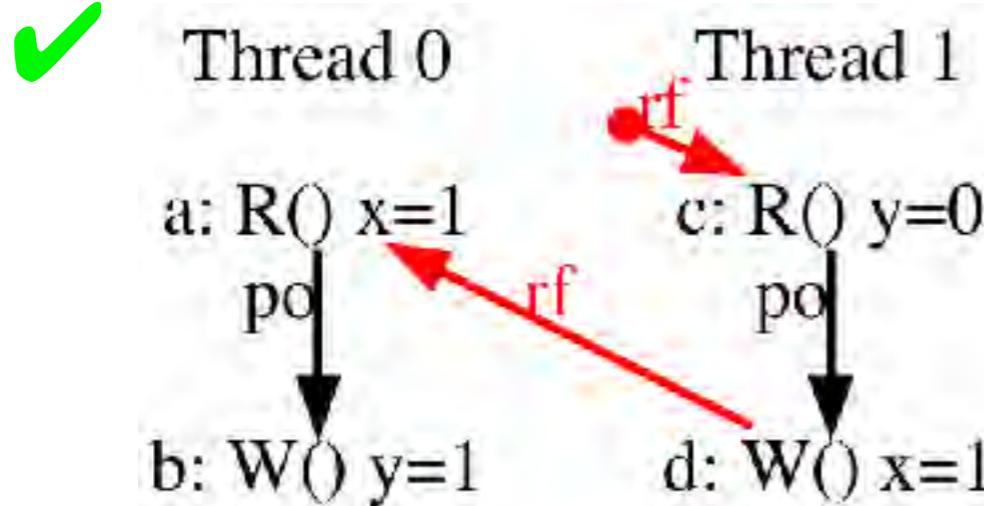
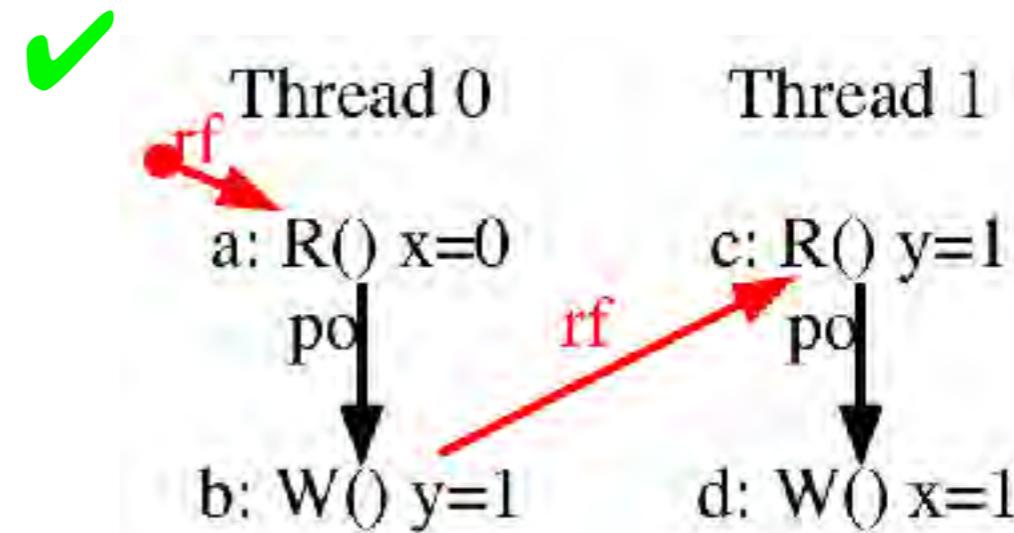
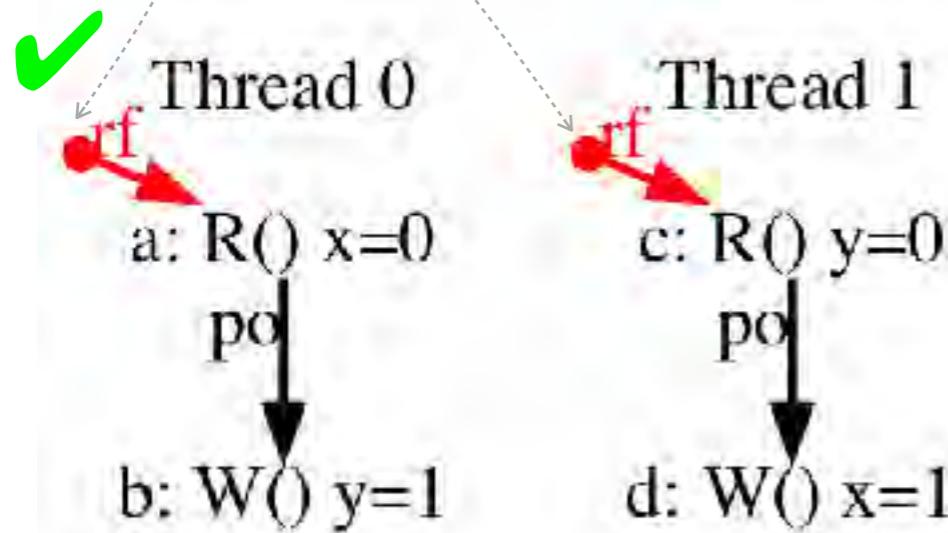
# Analytic semantics

- **Anarchic semantics:** all possible executions with cuts/histories with no restriction on rf (any read can read any value from any write to the same shared variable)
- **Communication consistency:** requirements on rf specified on an abstraction to a candidate execution
- **Analytic semantics:** all executions with cuts/histories which rf satisfies the consistency requirements

# Example of anarchic semantics: LB

read from  
initial write

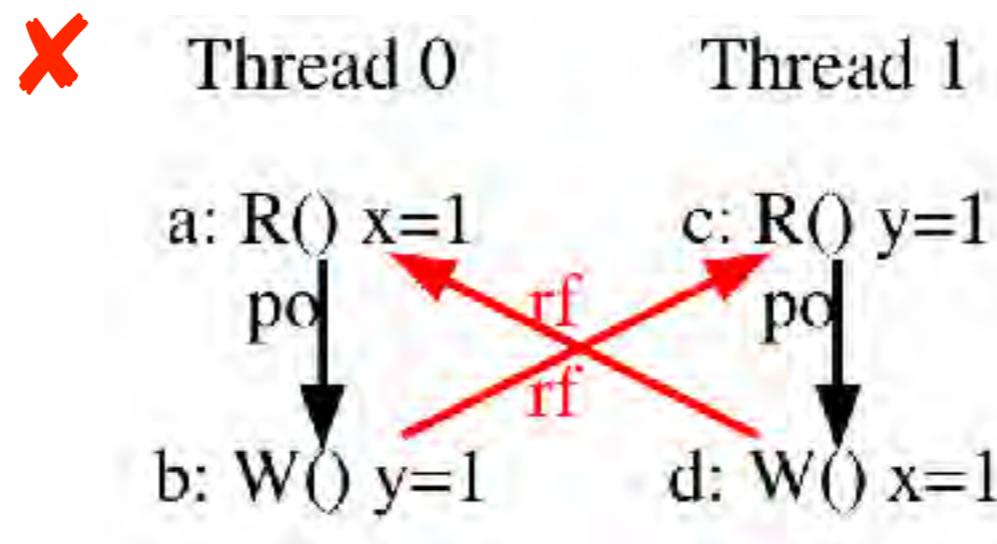
```
{ x = 0; y = 0; }
P0 | P1 ;
r[] r1 x | r[] r2 y ;
w[] y 1 | w[] x 1 ;
```



# Example of communication specification in the cat language for LB

irreflexive  $(po \mid rf)^+$

Rejects only the anarchic execution:



---

J. Alglave. *A Shared Memory Poetics*. PhD thesis, Université Paris 7, 2010.

J. Alglave, P. Cousot, and L. Maranget. Syntax and semantics of the cat language. *HSA Foundation*, Version 1.1:38 p., 16 Oct 2015b. URL <http://www.hsafoundation.com/?ddownload=5382>.

# Examples of architecture specification

- SC (sequential consistency):

```
let co = (IW*W) & loc
let fr = (rf^-1;co)
acyclic po | rf | co | fr as sc
```

- TSO:

```
let co = (IW*W) & loc
let fr = (rf^-1;co)
let po-loc = po & loc
acyclic po-loc | rf | co | fr as scpv
let ppo = po \ (W*R)
let rfe = rf & ext
acyclic ppo | rfe | co | fr as tso
```

- For lb:

```
acyclic (po | rf) as lb
```

**sc  $\Rightarrow$  lb, tso  $\not\Rightarrow$  lb**

# Fence specification:

- In Lisa:  
$$\begin{array}{l} \{ x = 0; y = 0; \} \\ P0 \quad | \quad P1 \quad ; \\ r[] \ r1 \ x \ | \ r[] \ r2 \ y \ ; \\ f[dep] \ | \ f[lw] \ ; \\ w[] \ y \ 1 \ | \ w[] \ x \ 1 \ ; \end{array}$$
- Implementation with dependencies and fences in TSO:

{ x = 0; y = 0; }

P0 | P1 ;

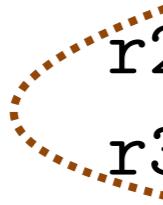
r[] r1 x | r[] r2 y ;

r2 = xor r1 r1 | mfence ;

r3 = r2+1 |

w[] y r3 | w[] x r3 ;

data dependency



# cat

- Handles one history at a time
- For each execution relies on:
  - the set  $E(\_)$  of events of the execution (partitionned into initial writes  $IW$ , writes  $W$ , read  $R$ , fences  $F$ , ...)
  - the program order  $po$  of events per process
  - the read-from relation  $rf$  per variable
- Has predefined relations  $loc$ ,  $ext$ , ...
- Can define new relations e.g.  $\ast, ;, |, \&, \backslash, +, ^{-1}, \dots$
- Accepts/eliminates the execution by defining relations  $r$  and checking irreflexive  $r$ , acyclic  $r$ , empty  $r$ , not empty  $r$

# ARM in cat

```
let fr = rf^-1;co
acyclic po-loc | rf | co | fr as scpv

let deps = addr | data
let rdw = po-loc & (fre;rfe)
let detour = po-loc & (coe ; rfe)

let ii0 = deps | rfi | rdw
let ic0 = 0
let ci0 = ctrlcfence(ISB) | detour
let cc0 = deps | ctrl | (addr;po)

let rec ii = ii0 | ci | (ic;ci) | (ii;ii)
and ic = ic0 | ii | cc | (ic;cc) | (ii;ic)
and ci = ci0 | (ci;ii) | (cc;ci)
and cc = cc0 | ci | (ci;ic) | (cc;cc)

let ppo = ii & R*R | ic & R*W

let dmb = fencerel(DMB)
let dsb = fencerel(DSB)
let fences = dmb|dsb
let A-cumul = rfe;fences

let hb = ppo | fences | rfe
acyclic hb as no-thin-air

let prop-base = (fences | A-cumul);hb*
let prop = (prop-base & W*W)| (com*; prop-base*; fences; hb*)

irreflexive fre;prop;hb* as observation
acyclic co | prop as propagation
```

# Invariance proof method for weakly consistent parallelism

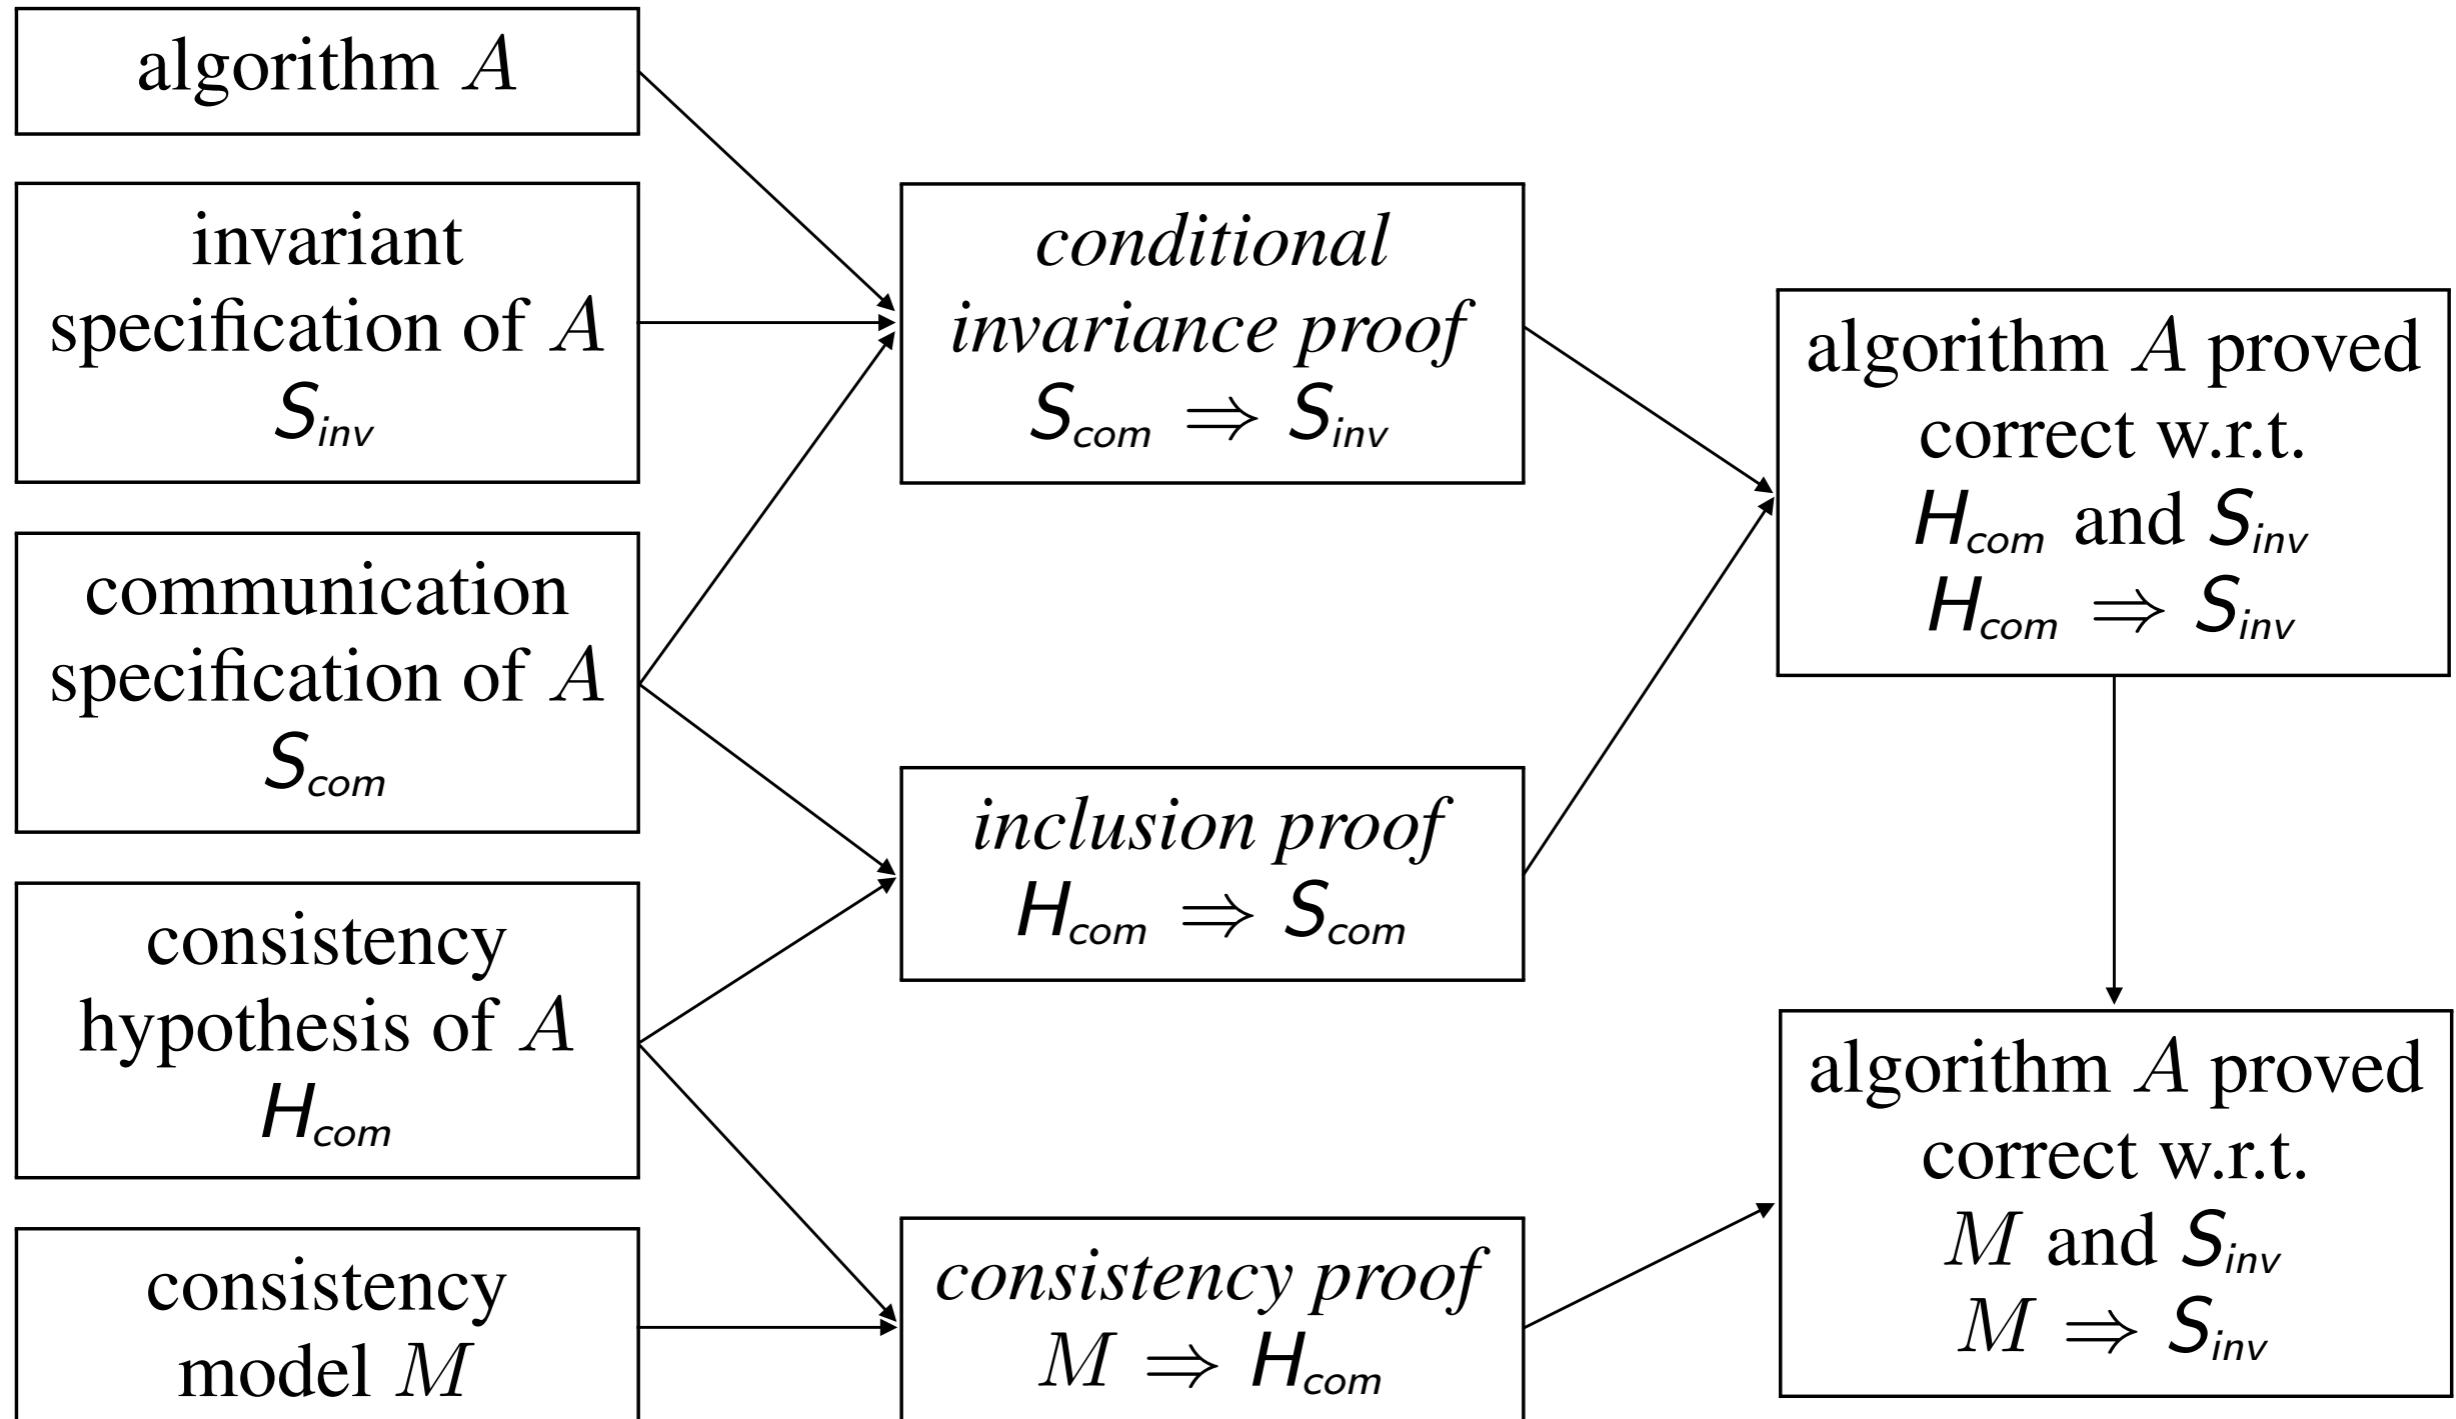
# Difficulties

- There is no longer a notion of instantaneous value of the shared variables:
  - ⇒ **pythia variables** (denoting values of variables when read)
  - ⇒ **communications rf** (keeping track of which writes events the pythia variables take their values from)
  - ⇒ **stamps** (keeping track of events to distinguish different instruction executions)

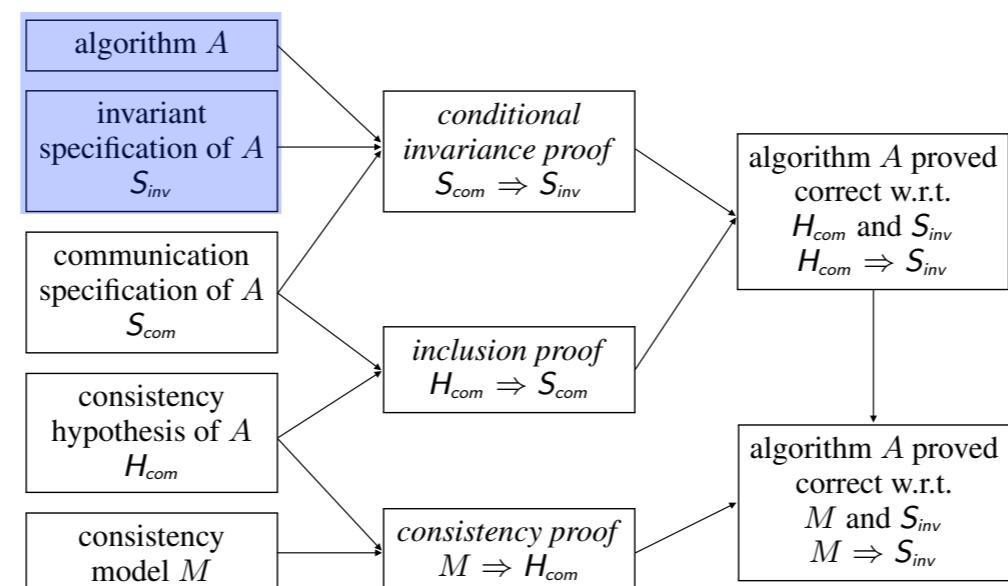
# Difficulties

- We have to make hypotheses on how communications do happen:  
⇒ communication specification  $S_{com}$
- We have to show that the communication specification is correctly implemented on an architecture:  
⇒ a way to mix invariant  $S_{com}$  and cat specifications

# Methodology



# Invariant



# Pythia variables

- Unique name given to communicated values during execution (using stamps)

```
0:{ w F1 false; w F2 false; w T 0; }
```

```
P0:
```

```
1:w[] F1 true
```

```
2:w[] T 2
```

```
3:repeat {i}
```

```
4: r[] R1 F2 { \rightsquigarrow F24i}
```

```
5: r[] R2 T { \rightsquigarrow T5i}
```

```
6:until $\neg R1 \vee R2 = 1$ {iend}
```

```
7:skip (* CS1 *)
```

```
8:w[] F1 false
```

```
9:
```

```
P1:
```

```
10:w[] F2 true;
```

```
11:w[] T 1;
```

```
12:repeat {j}
```

```
13: r[] R3 F1; { \rightsquigarrow F113j}
```

```
14: r[] R4 T; { \rightsquigarrow T14j}
```

```
15:until $\neg R3 \vee R4 = 2$; {jend}
```

```
16:skip (* CS2 *)
```

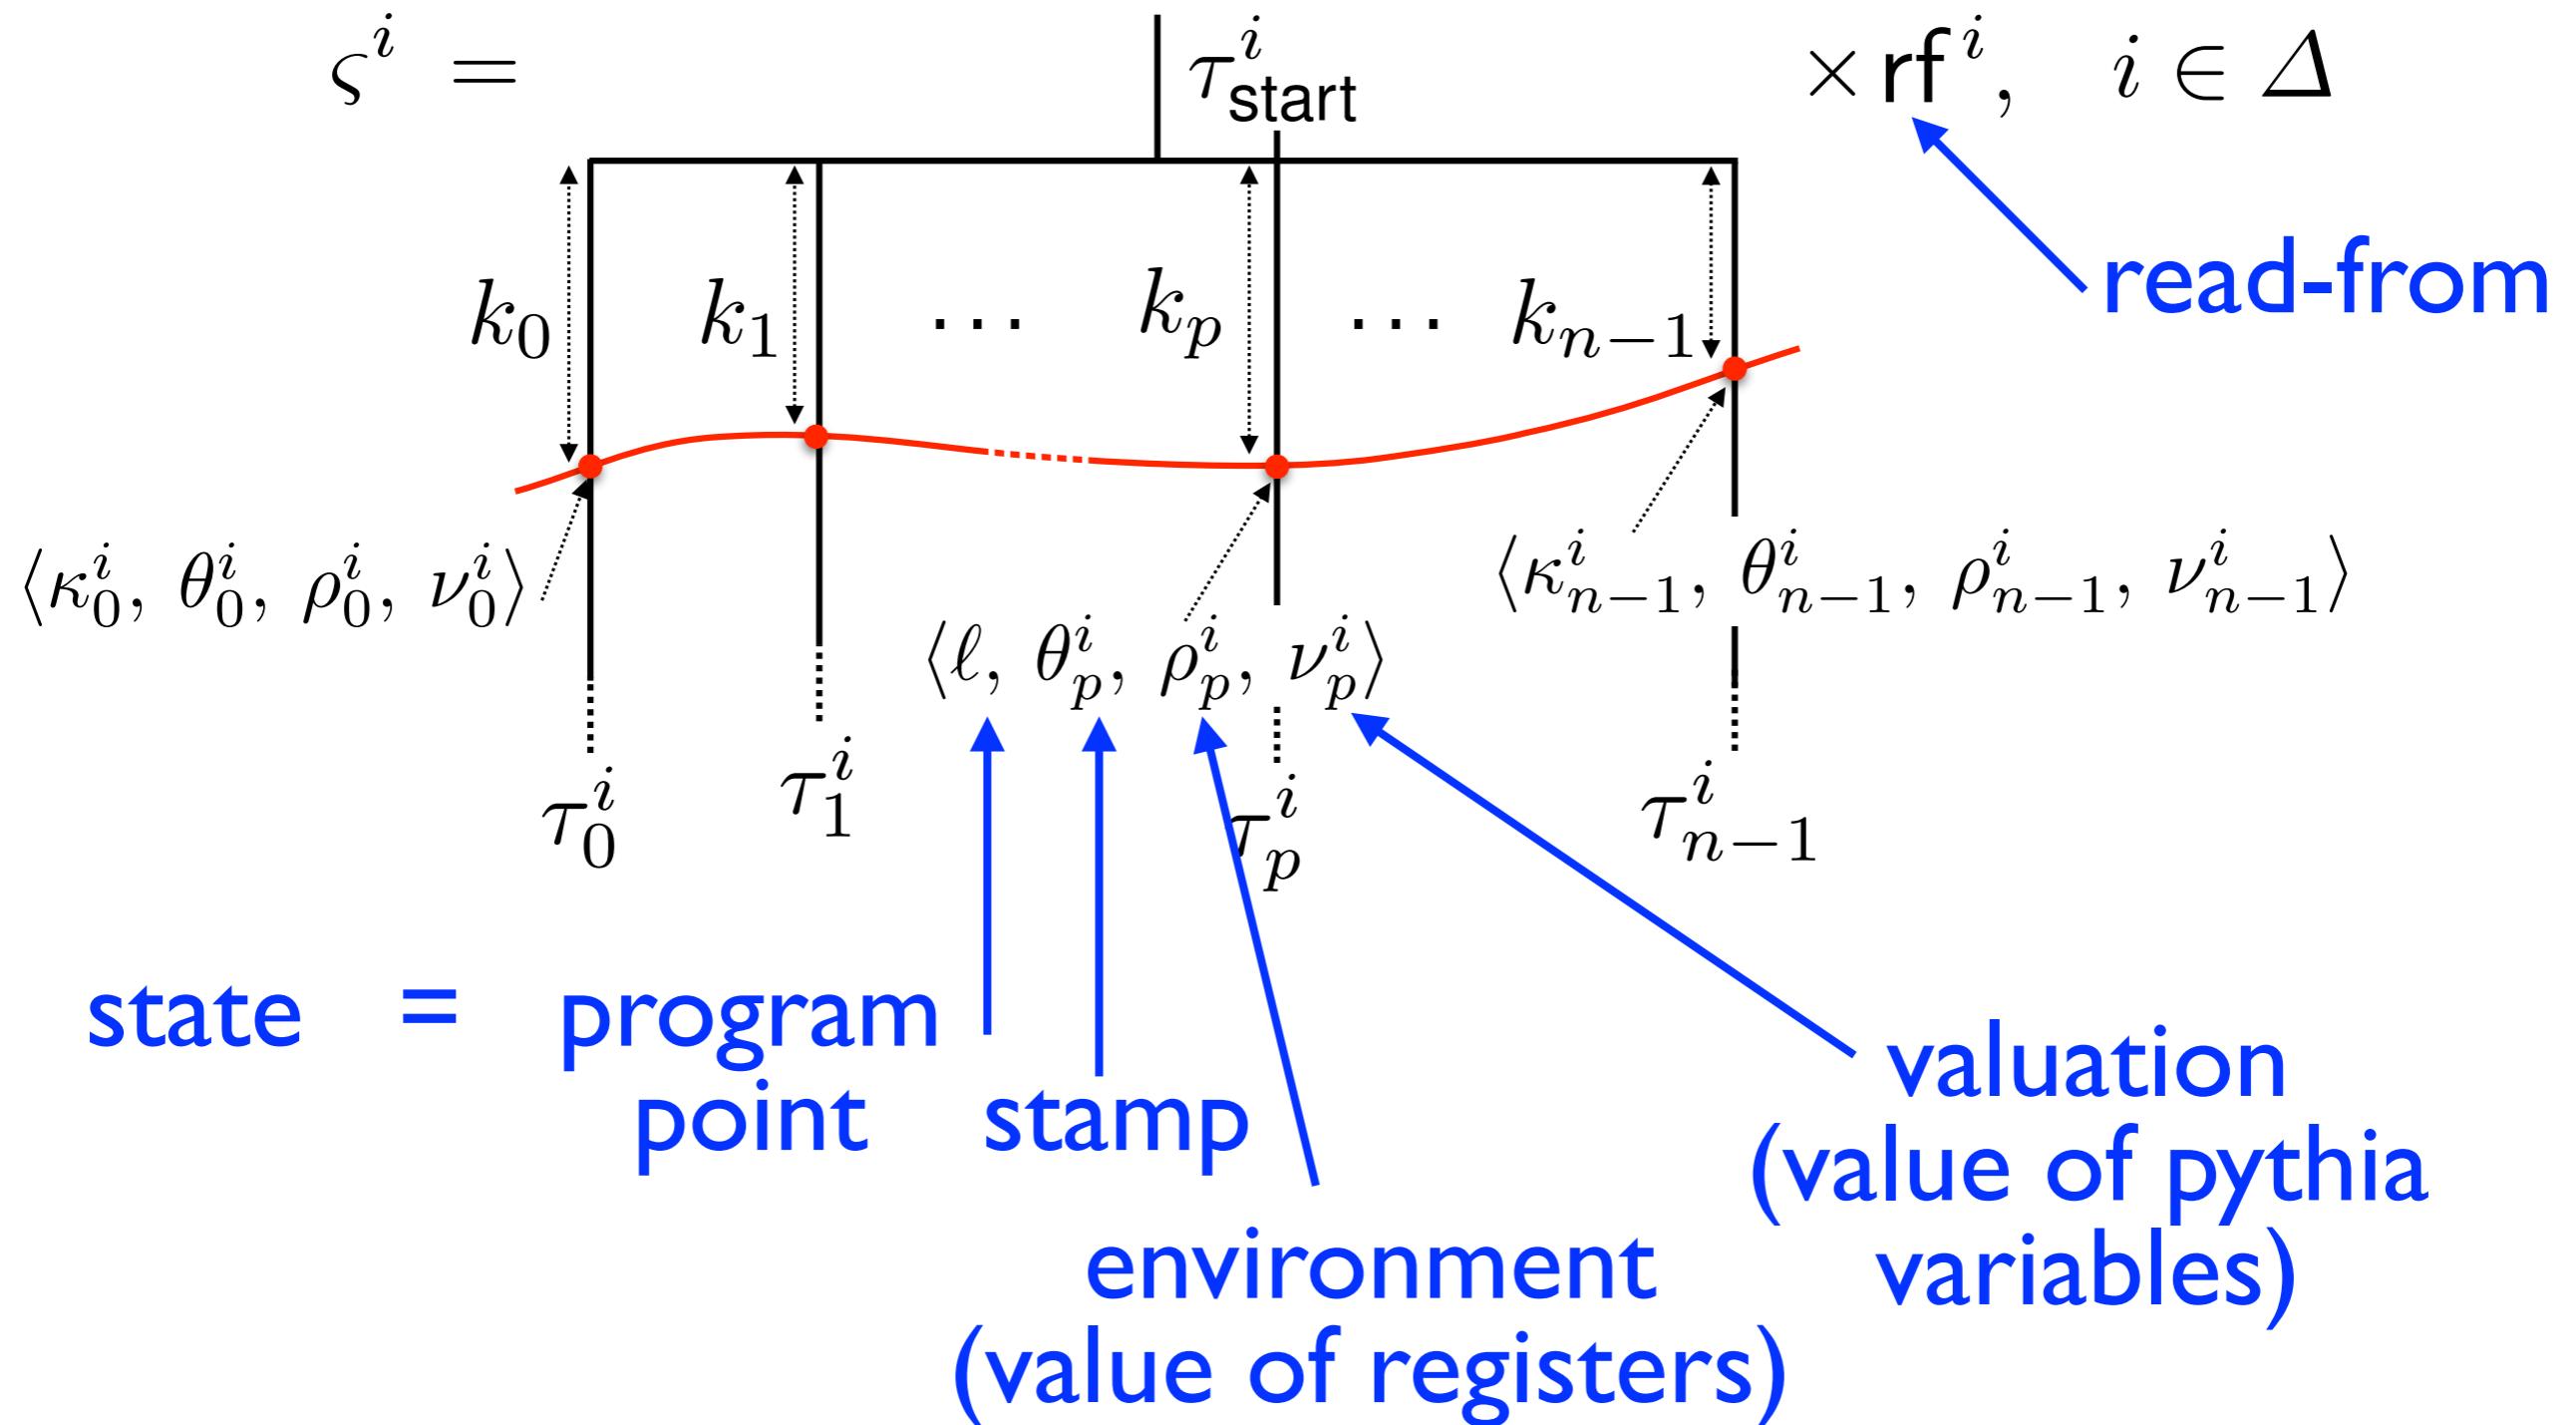
```
17:w[] F2 false;
```

```
18:
```

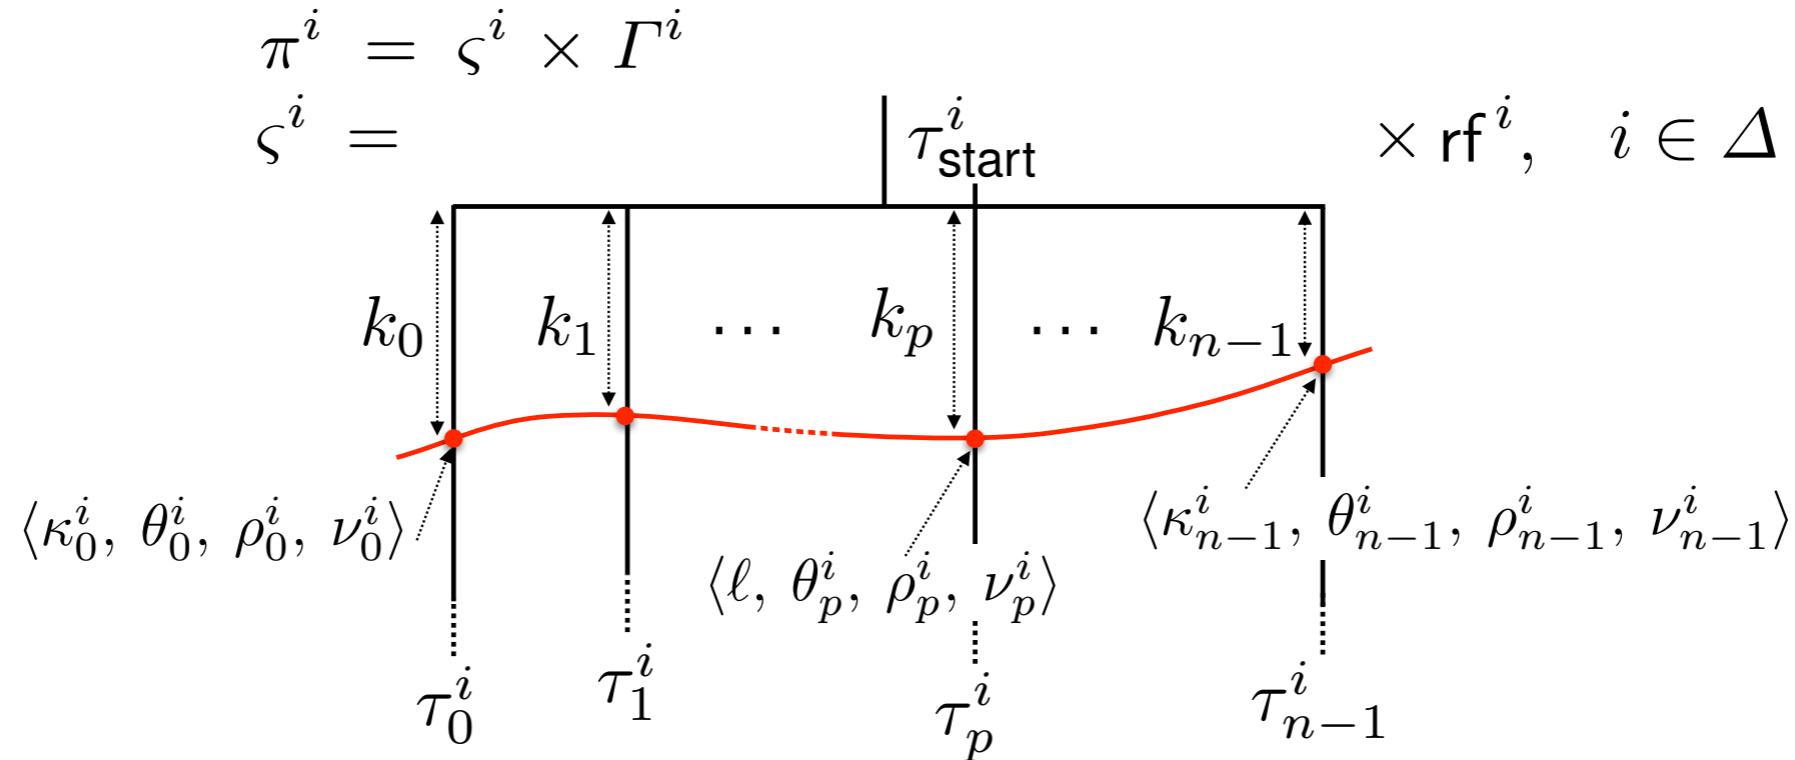
Stamp: label,counter

Pythia variables

# Invariance abstraction



# Invariance abstraction



$$\alpha_a(\{\pi^i \mid i \in \Delta\}) \triangleq \prod_{p \in \mathbb{P}} \prod_{\ell \in \mathbb{L}(p)} \bigcup_{i \in \Delta} \{ \langle \kappa_{0,k_0}^i, \theta_{0,k_0}^i, \rho_{0,k_0}^i, \nu_{0,k_0}^i, \dots,$$

$$\nu_{p-1,k_{p-1}}^i, \theta_{p,k_p}^i, \rho_{p,k_p}^i, \nu_{p,k_p}^i, \kappa_{p+1,k_{p+1}}^i, \dots, \kappa_{n-1,k_{n-1}}^i, \theta_{n-1,k_{n-1}}^i,$$

$$\rho_{n-1,k_{n-1}}^i, \nu_{n-1,k_{n-1}}^i, \text{rf}^i \rangle \mid \forall q \in [0, n[ \setminus \{p\} \cdot \underline{\tau_q^i}_{k_q} =$$

$$\mathfrak{s} \langle \kappa_{q,k_q}^i, \theta_{q,k_q}^i, \rho_{q,k_q}^i, \nu_{q,k_q}^i \rangle \wedge \underline{\tau_p^i}_{k_p} = \mathfrak{s} \langle \ell, \theta_{p,k_p}^i, \rho_{p,k_p}^i, \nu_{p,k_p}^i \rangle \}.$$

# Invariant

- An **invariant**  $S_{inv}(p)$  at point  $p$  of process  $P_i$  is a statement relating
  - the **program points**  $p_1, \dots, p_{i-1}, p_{i+1}, \dots, p_m$  of the other processes
  - the **pythia variables** (forbidden to mention of shared variables)
  - the **local registers** of all processes
  - the **communications** ( $rf$ )which always holds when at the cut where execution reaches point  $p$  of process  $P_i$  and the other processes are at  $p_1, \dots, p_{i-1}, p_{i+1}, \dots, p_m$

# Example (Peterson)

```

0: { w F1 false; w F2 false; w T 0; }
 {F1=false ∧ F2=false ∧ T=0} }

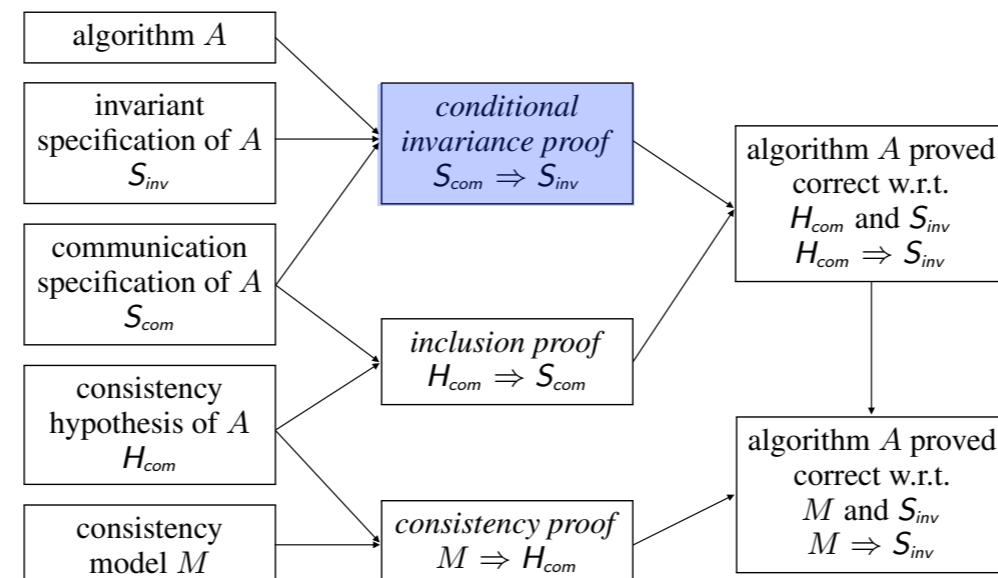
1: {R1=0 ∧ R2=0}
 w[] F1 true
2: {R1=0 ∧ R2=0}
 w[] T 2
3: {R1=0 ∧ R2=0}
 do {i}
4: {(i=0 ∧ R1=0 ∧ R2=0) ∨
 (i>0 ∧ R1=F24i-1 ∧ R2=T5i-1)}
 r[] R1 F2 {~> F24i}
5: {R1=F24i ∧ (i=0 ∧ R2=0) ∨
 (i>0 ∧ R2=T5i-1)}
 r[] R2 T {~> T5i}
6: {R1=F24i ∧ R2=T5i}
 while R1 ∧ R2≠1 {i_end}
7: {¬F24i_end ∨ T5i_end=1}
 skip (* CS1 *)
8: {¬F24i_end ∨ T5i_end=1}
 w[] F1 false
9: {¬F24i_end ∨ T5i_end=1}

10: {R3=0 ∧ R4=0}
 w[] F2 true;
11: {R3=0 ∧ R4=0}
 w[] T 1;
12: {R3=0 ∧ R4=0}
 do {j}
13: {(j=0 ∧ R3=0 ∧ R4=0) ∨
 (j>0 ∧ R3=F113j-1 ∧ R4=T14j-1)}
 r[] R3 F1 {~> F113j};
14: {R3=F113j ∧ (j=0 ∧ R4=0) ∨
 (j>0 ∧ R4=T14j-1)}
 r[] R4 T; {~> T14j}
15: {R3=F113j ∧ R4=T14j}
 while R3 ∧ R4≠2 {j_end} ;
16: {¬F113j_end ∨ T14j_end=2}
 skip (* CS2 *)
17: {¬F113j_end ∨ T14j_end=2}
 w[] F2 false;
18: {¬F113j_end ∨ T14j_end=2}

```

(these invariants are for the anarchic semantics, so all communications are possible, no constraints on rf)

# Invariance proof



# Verification conditions

- Sequential proof
- Absence of interference proof
- Communication proof

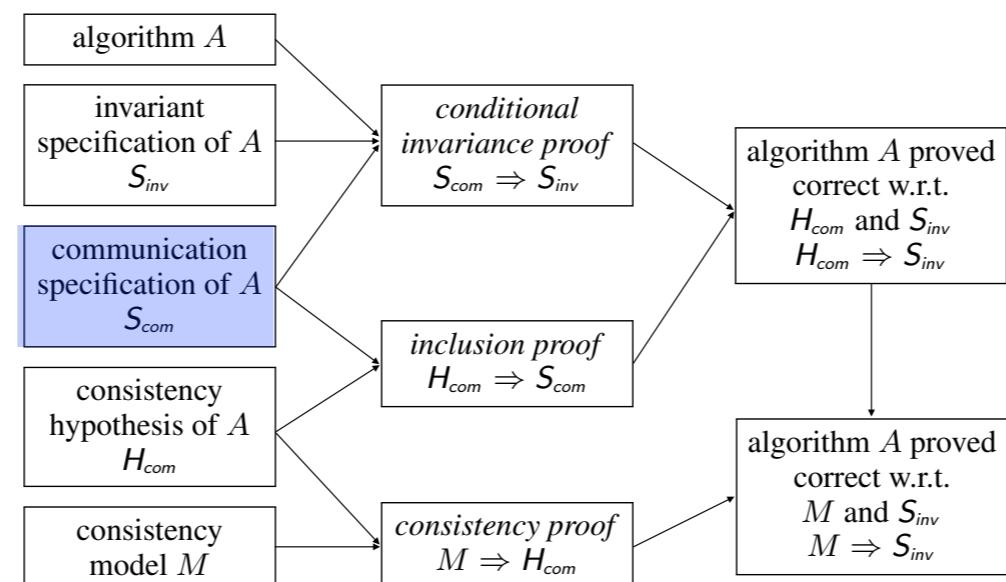
Examples:

- $\{ P(R, \dots, rf) \wedge \langle w(x, v), r(\theta, x) \rangle \in rf \}$   
read  $x \ R \ \{ \rightsquigarrow x_0 \}$  *communication*  
 $\{ P[R \leftarrow x_0, x_0 \leftarrow v, \dots, rf] \}$
- $\{ P \}$  fence  $\{ P \}$  (fences are markers in the execution)
- $\{ P \}$  write  $R \ x \ { P }$  (a write has no local effect)

# Communication proof

- The communications  $\text{rf}$  must be checked to be well-formed (none allowed by  $H_{cm}$  should miss, see later)
- If  $\langle w(P, p, \theta, x, v), r(P', p', \theta', x, x_{\theta'}) \rangle \in \text{rf}$  then:
  - The read instruction of at point  $p'$  process  $P'$  must read from an *initial* or a *reachable* write
  - A read event (for a given stamp  $\theta'$ ) must read from a *unique write event* with the same variable  $x$
  - The *value assigned* to the read pythia variable  $x_{\theta'}$  must be that of  $v$  the matching write

# Communication specification $S_{com}$



# Communication specification

- The algorithm  $A$  is often incorrect for the anarchic semantics
- The allowable communications are specified by a communication specification  $S_{com}$  (i.e. an invariant constraining the allowed communications  $rf$ )
- This communication specification can often be calculated from the anarchic invariant and the inductive invariant  $S_{ind}$

# Example (Peterson)

at 7  $\wedge$  at 16  
 $\Rightarrow (\neg F2_4^{i_{\text{end}}} \vee T_5^{i_{\text{end}}} = 1) \wedge (\neg F1_{13}^{j_{\text{end}}} \vee T_{14}^{j_{\text{end}}} = 2) \}$   
    *{i.e. the invariants at lines 7: and 16: hold}*  
 $\Rightarrow \neg S_{\text{com}} \quad \{ \text{since by taking } i = i_{\text{end}} \text{ and } j = j_{\text{end}}, \text{ we have}$   
 $(F2_4^i = \text{false} \vee T_5^i = 1) \wedge (F1_{13}^j = \text{false} \vee T_{14}^j = 2) \}$

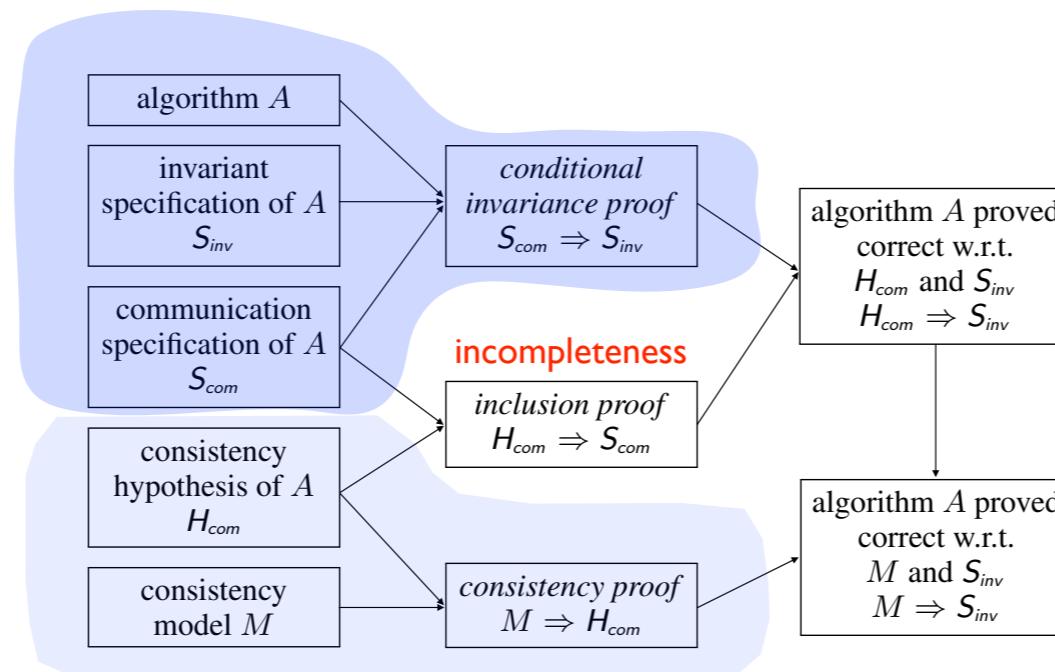
so that Peterson has been proved correct under the hypothesis that the communication specification  $S_{\text{com}}$  holds:

$$\begin{aligned} S_{\text{com}} \triangleq & \neg [\exists i, j. [\text{rf}\langle F2_4^i, \langle 0:, \text{false} \rangle \rangle \vee \text{rf}\langle F2_4^i, \langle 17:, \text{false} \rangle \rangle \\ & \vee \text{rf}\langle T_5^i, \langle 11:, 1 \rangle \rangle] \wedge [\text{rf}\langle F1_{13}^j, \langle 0:, \text{false} \rangle \rangle \\ & \vee \text{rf}\langle F1_{13}^j, \langle 8:, \text{false} \rangle \rangle \vee \text{rf}\langle T_{14}^j, \langle 2:, 2 \rangle \rangle]] \end{aligned}$$

(preventing the incorrect case)

# Soundness and completeness

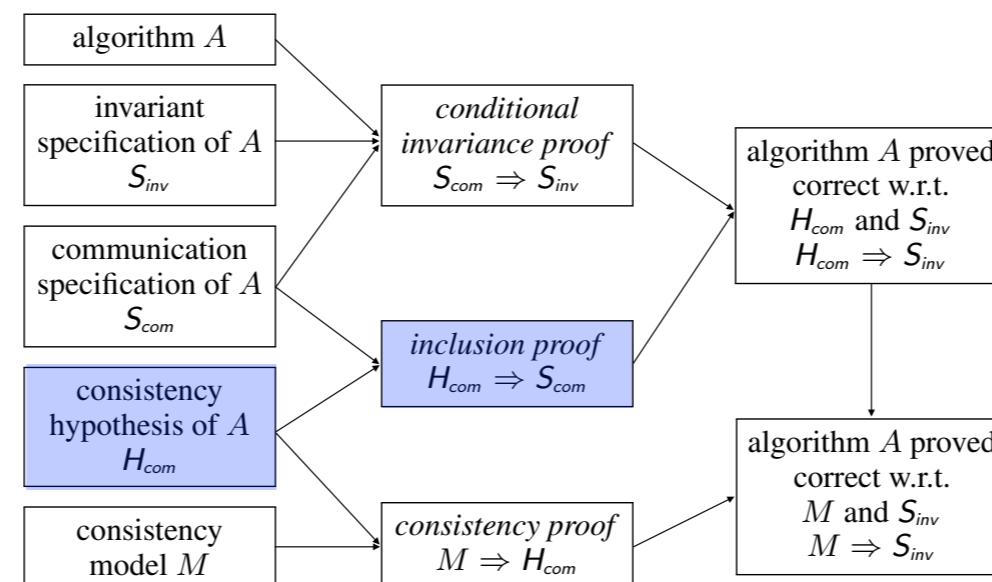
- The invariance proof method is derived from the truly parallel semantics with cuts by **calculational design**  
⇒ **soundness** and (relative) **completeness**



- A consistency specification  $H_{com}$  may be less expressive than  $S_{com} \Rightarrow incompleteness$  (\*)

(\*) e.g. hardware cannot restrict a read to input from writes writing odd numbers.

# Consistency hypothesis and inclusion proof



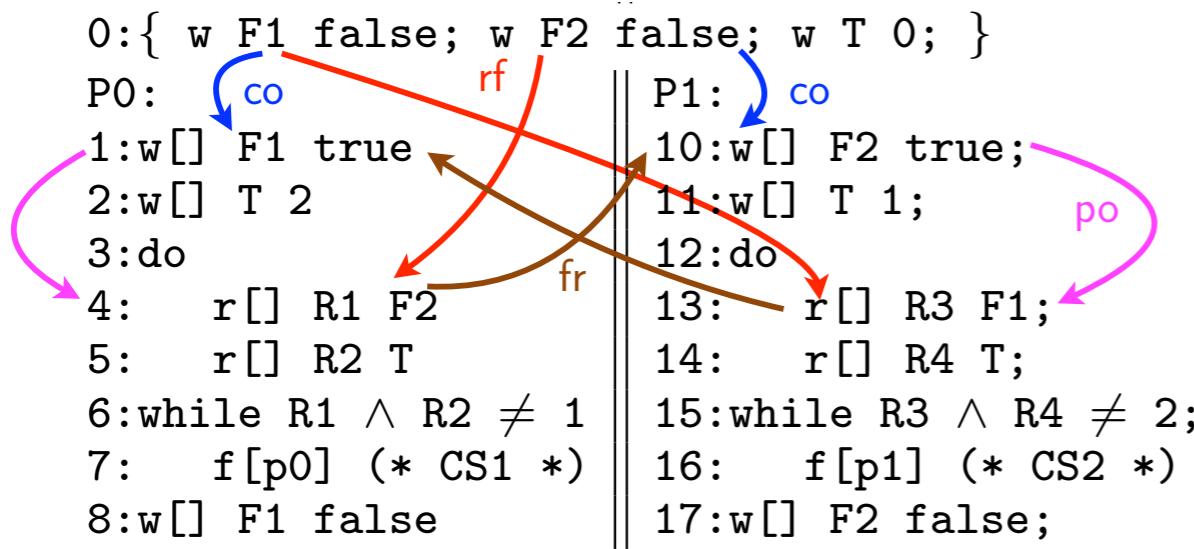
# Consistency hypothesis

- The communication specification  $S_{com}$  is useful to reason on invariance, but not on machine architecture
- We express  $S_{com}$  as a consistency hypothesis  $H_{com}$  expressed in the cat language
- $H_{com}$  is derived from  $S_{com}$  by calculations design while doing the inclusion proof

# Inclusion proof

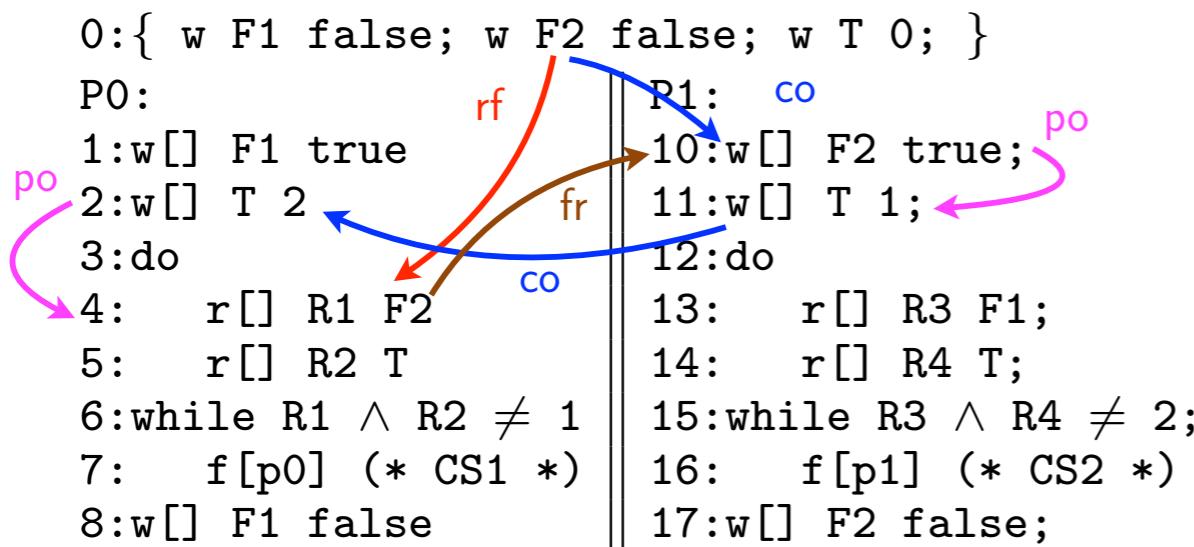
- Inclusion proof:  $\neg S_{com} \Rightarrow \neg H_{com}$
- Calculational design of  $H_{com}$ :
  - Calculate all possible execution scenarios violating  $S_{com}$ 
$$S_{com} \triangleq \neg [\exists i, j. [\text{rf}\langle F2_4^i, \langle 0:, \text{false} \rangle \rangle \vee \text{rf}\langle F2_4^i, \langle 17:, \text{false} \rangle \rangle \vee \text{rf}\langle T_5^i, \langle 11:, 1 \rangle \rangle] \wedge [\text{rf}\langle F1_{13}^j, \langle 0:, \text{false} \rangle \rangle \vee \text{rf}\langle F1_{13}^j, \langle 8:, \text{false} \rangle \rangle \vee \text{rf}\langle T_{14}^j, \langle 2:, 2 \rangle \rangle]]$$
  - Prevent each of them by a cat specification
  - $H_{com}$  is their conjunction

# Example: Peterson



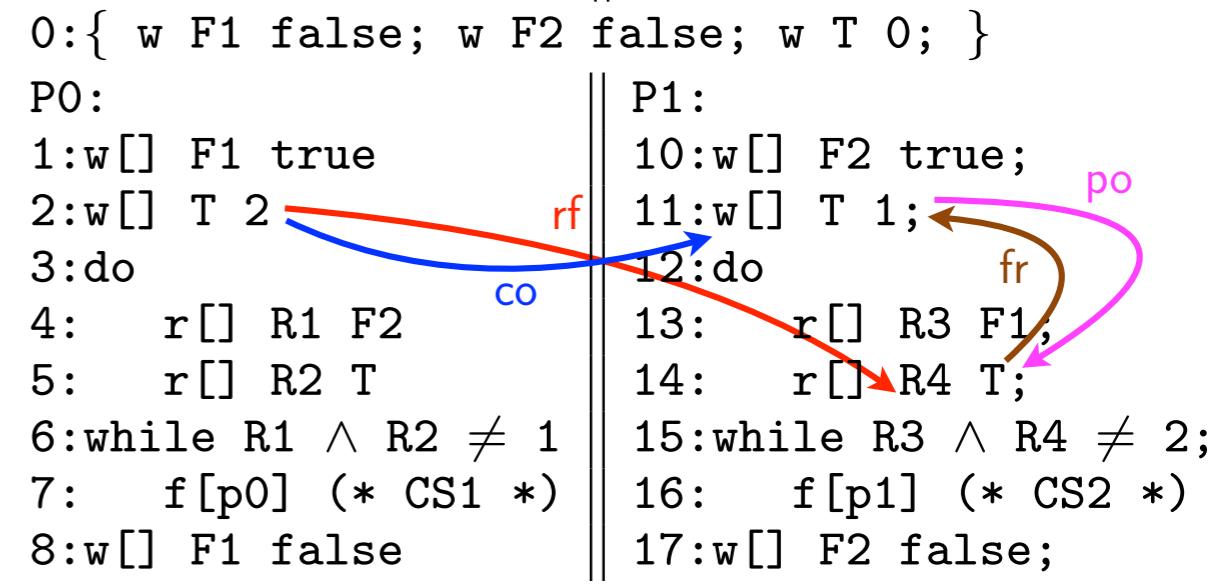
$4 \xrightarrow{\text{fr}} 10 \xrightarrow{\text{po}} 13 \xrightarrow{\text{fr}} 1 \xrightarrow{\text{po}} 4$

case 1: 0:F2,0:F1



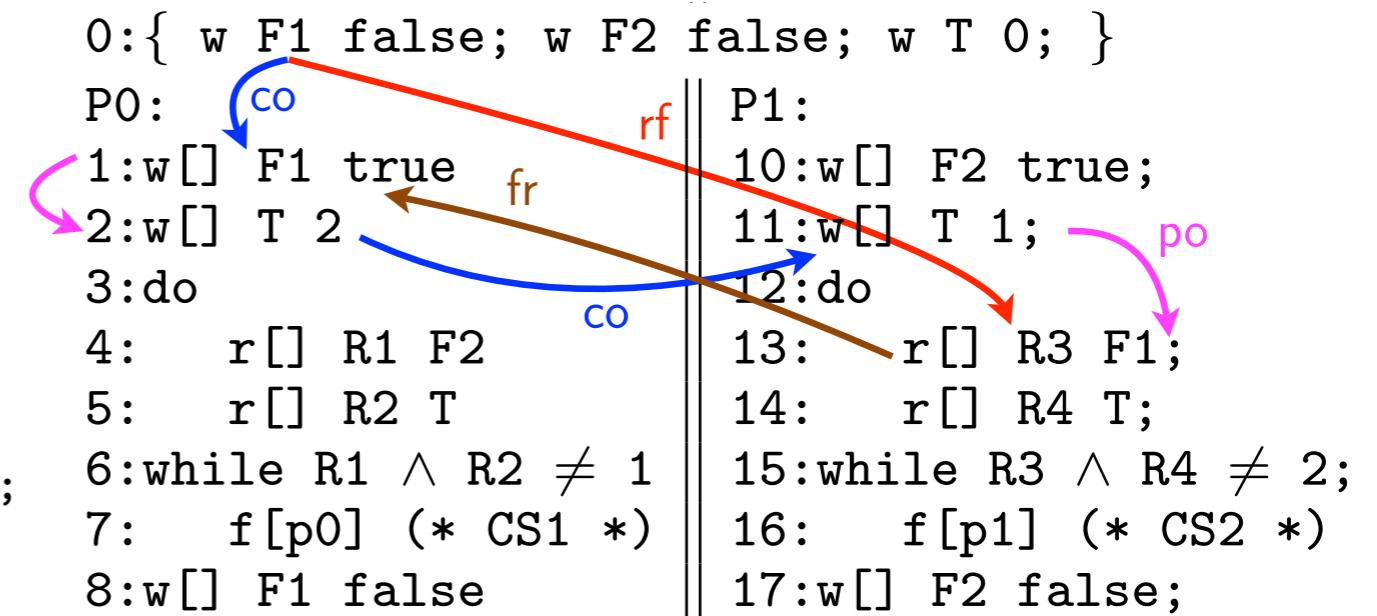
$11 \xrightarrow{\text{co}} 2 \xrightarrow{\text{po}} 4 \xrightarrow{\text{fr}} 10 \xrightarrow{\text{po}} 11$

case 2b: 0:F2,1:F1 ( $11 \xrightarrow{\text{co}} 2$ )



$14 \xrightarrow{\text{fr}} 11 \xrightarrow{\text{po}} 14$

case 2a: 0:F2,1:F1 ( $2 \xrightarrow{\text{co}} 11$ )



$2 \xrightarrow{\text{co}} 11 \xrightarrow{\text{po}} 13 \xrightarrow{\text{fr}} 1 \xrightarrow{\text{po}} 2$

case 3a: 10:F2,0:F1 ( $2 \xrightarrow{\text{co}} 11$ )

# Example: Peterson

```

0:{ w F1 false; w F2 false; w T 0; }
P0:
1:w[] F1 true
2:w[] T 2
3:do
4: r[] R1 F2
5: r[] R2 T
6:while R1 ∧ R2 ≠ 1
7: f[p0] (* CS1 *)
8:w[] F1 false

```

5 —fr→ 2 —po→ 5

case 3b: 10:F2,0:F1 (11 —co→ 2)

```

0:{ w F1 false; w F2 false; w T 0; }
P0:
1:w[] F1 true
2:w[] T 2
3:do
4: r[] R1 F2
5: r[] R2 T
6:while R1 ∧ R2 ≠ 1
7: f[p0] (* CS1 *)
8:w[] F1 false

```

5 —fr→ 2 —po→ 5

case 4b: 10:F2,1:F1 (11 —co→ 2)

```

P1:
10:w[] F2 true;
11:w[] T 1;
12:do
13: r[] R3 F1;
14: r[] R4 T;
15:while R3 ∧ R4 ≠ 2;
16: f[p1] (* CS2 *)
17:w[] F2 false;

```

11 —co→ 2

```

0:{ w F1 false; w F2 false; w T 0; }
P0:
1:w[] F1 true
2:w[] T 2
3:do
4: r[] R1 F2
5: r[] R2 T
6:while R1 ∧ R2 ≠ 1
7: f[p0] (* CS1 *)
8:w[] F1 false

```

10:F2,1:F1 (2 —co→ 11)

```

P1:
10:w[] F2 true;
11:w[] T 1;
12:do
13: r[] R3 F1;
14: r[] R4 T;
15:while R3 ∧ R4 ≠ 2;
16: f[p1] (* CS2 *)
17:w[] F2 false;

```

11 —co→ 2

case 4b: 10:F2,1:F1 (11 —co→ 2)

```

0:{ w F1 false; w F2 false; w T 0; }
P0:
1:w[] F1 true
2:w[] T 2
3:do
4: r[] R1 F2
5: r[] R2 T
6:while R1 ∧ R2 ≠ 1
7: f[p0] (* CS1 *)
8:w[] F1 false

```

10:F2,1:F1 (2 —co→ 11)

```

P1:
10:w[] F2 true;
11:w[] T 1;
12:do
13: r[] R3 F1;
14: r[] R4 T;
15:while R3 ∧ R4 ≠ 2;
16: f[p1] (* CS2 *)
17:w[] F2 false;

```

17:F2,8:F1 (4 —rf→ 13 —po→ 17 —rf→ 4)

# Example: Peterson

```

0:{ w F1 false; w F2 false; w T 0; }
P0:
1:w[] F1 true
2:w[] T 2
3:do
4: r[] R1 F2
5: r[] R2 T
po 6:while R1 ∧ R2 ≠ 1
7: f[p0] (* CS1 *)
8:w[] F1 false
rf
case 6: 8:F1

```

Diagram showing a sequence of operations:

- Initial state: 0:{ w F1 false; w F2 false; w T 0; }
- Step 1: P0: 1:w[] F1 true
- Step 2: 2:w[] T 2
- Step 3: 3:do
- Step 4: 4: r[] R1 F2
- Step 5: 5: r[] R2 T
- Step 6: po 6:while R1 ∧ R2 ≠ 1
- Step 7: 7: f[p0] (\* CS1 \*)
- Step 8: 8:w[] F1 false
- Step 9: rf
- Step 10: 10:w[] F2 true;
- Step 11: 11:w[] T 1;
- Step 12: 13:do
- Step 13: 13:r[] R3 F1; po
- Step 14: 14:r[] R4 T;
- Step 15: 15:while R3 ∧ R4 ≠ 2;
- Step 16: 16:f[p1] (\* CS2 \*)
- Step 17: 17:w[] F2 false;

Annotations:

- Red curved arrow from 7 to 13 labeled "rf".
- Blue curved arrow from 8 to 17 labeled "cut".
- Pink curved arrows from 13 to 14 and 14 to 15 labeled "po".

Sequence: 7 → po → 8 → rf → 13 → po → 17 → cut → 7

```

0:{ w F1 false; w F2 false; w T 0; }
P0:
1:w[] F1 true
2:w[] T 2
3:do
4: r[] R1 F2
5: r[] R2 T
rf 6:while R1 ∧ R2 ≠ 1
7: f[p0] (* CS1 *)
8:w[] F1 false
po 13: r[] R3 F1;
14: r[] R4 T;
15:while R3 ∧ R4 ≠ 2;
16:f[p1] (* CS2 *) po
17:w[] F2 false;

```

Diagram showing a sequence of operations:

- Initial state: 0:{ w F1 false; w F2 false; w T 0; }
- Step 1: P0: 1:w[] F1 true
- Step 2: 2:w[] T 2
- Step 3: 3:do
- Step 4: 4: r[] R1 F2
- Step 5: 5: r[] R2 T
- Step 6: rf 6:while R1 ∧ R2 ≠ 1
- Step 7: 7: f[p0] (\* CS1 \*)
- Step 8: 8:w[] F1 false
- Step 9: po 13: r[] R3 F1;
- Step 10: 14: r[] R4 T;
- Step 11: 15:while R3 ∧ R4 ≠ 2;
- Step 12: 16: f[p1] (\* CS2 \*) po
- Step 13: 17:w[] F2 false;

Annotations:

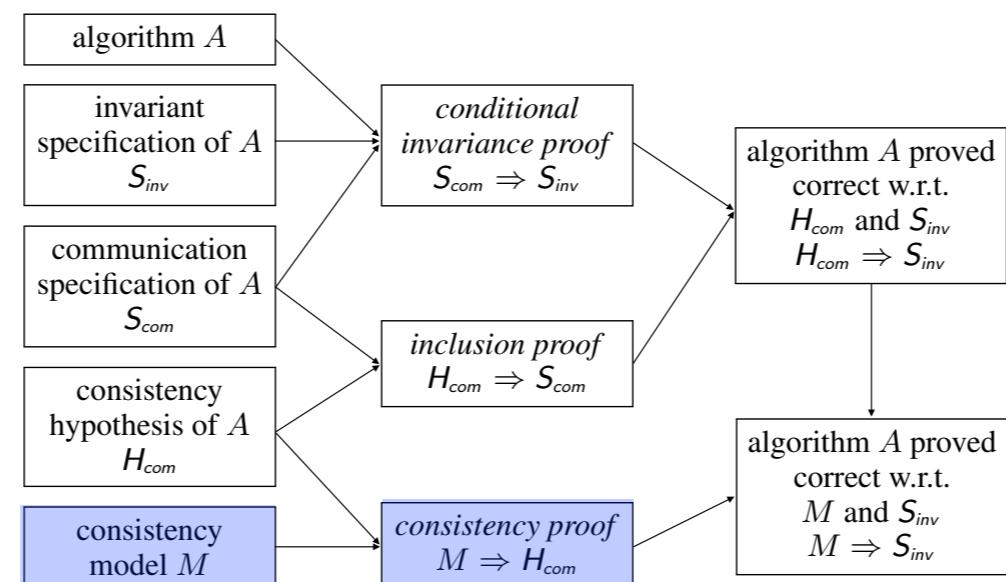
- Red curved arrow from 7 to 4 labeled "rf".
- Blue curved arrow from 8 to 17 labeled "cut".
- Pink curved arrows from 13 to 14 and 14 to 15 labeled "po".

Sequence: 16 → po → 17 → rf → 4 → po → 7 → cut → 16

case 7: 17:F2

- the cut relation can be expressed in cat using tags on fence markers  $f[p_0]$  and  $f[p_1]$
- $H_{com}$  is
  - irreflexive  $fr;po;fr;po$
  - irreflexive  $fr;po$
  - irreflexive  $co;po;fr;po$
  - irreflexive  $po;rf;po;rf$
  - irreflexive  $po;rf;po;cut$

# Consistency model and proof



# Example: Peterson in SC

- $H_{com}$  is
  - irreflexive  $fr;po;fr;po$
  - irreflexive  $fr;po$
  - irreflexive  $co;po;fr;po$
  - irreflexive  $po;rf;po;rf$
  - irreflexive  $po;rf;po;cut$*
- Sequential consistency in cat:  
let  $fr = (rf^\sim - 1 ; co)$   
acyclic  $po \mid rf \mid co \mid fr$  as sc
- Forbid all first 4 cases

# Example: Peterson in SC

```

0:{ w F1 false; w F2 false; w T 0; }
P0:
1:w[] F1 true
2:w[] T 2
3:do
4: r[] R1 F2
5: r[] R2 T
po 6:while R1 ∧ R2 ≠ 1
7: f[p0] (* CS1 *)
8:w[] F1 false
 cut
 13:r[] R3 F1; po
 14:r[] R4 T;
 15:while R3 ∧ R4 ≠ 2;
 16:f[p1] (* CS2 *)
 17:w[] F2 false;

```

7 →  
case 6: 8:F1

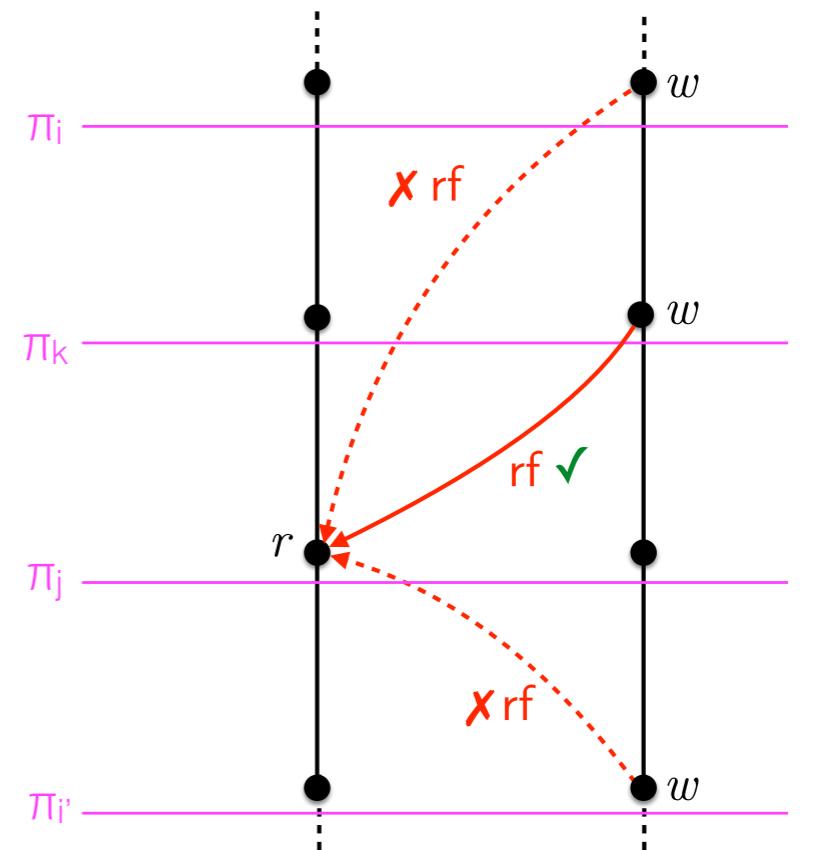
```

0:{ w F1 false; w F2 false; w T 0; }
P0:
1:w[] F1 true
2:w[] T 2
3:do
4: r[] R1 F2
5: r[] R2 T
Xrf 6:while R1 ∧ R2 ≠ 1
7: f[p0] (* CS1 *)
8:w[] F1 false
 cut
 13:r[] R3 F1; po
 14:r[] R4 T;
 15:while R3 ∧ R4 ≠ 2;
 16:f[p1] (* CS2 *)
 17:w[] F2 false;

```

16 →  
case 7: 17:F2

- The last case follows from the truly parallel execution trace semantics with cuts for sequential consistency



# Example: Peterson in TSO

- $H_{com}$  is not forbidden by TSO:

```
let fr = (rf^-1;co)
```

```
let po-loc = po & loc
```

```
acyclic po-loc | rf | co | fr as scpv
```

```
let ppo = po \ (W*R)
```

```
let rfe = rf & ext
```

```
acyclic ppo | rfe | co | fr as tso
```

- For example the case I,

$$\langle w_1, r_4 \rangle \in fr ; po ; fr ; po$$

is not forbidden by TSO since  $\langle w, r \rangle$  pairs on different variables are excluded from ppo.

# Implementation with (weak) cat fences

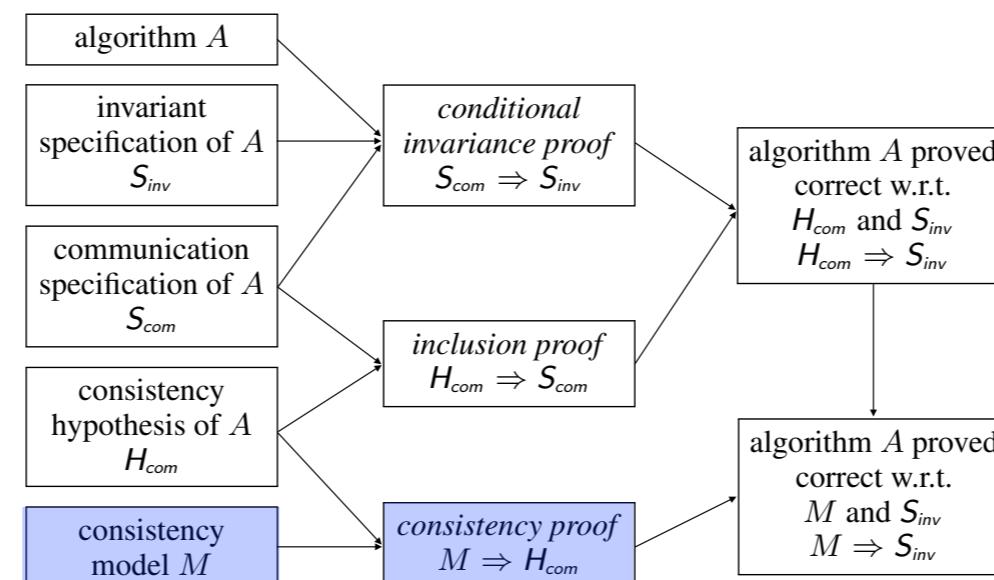
# Implementation with fences

```
0:{ F1 = 0; F2 = 0; T = 0; }
1: w[] F1 1 | 10: w[] F2 1 ;
2: w[] T 2 | 11: w[] T 1 ;
3: do | 12: do ;
 f[fhw] | f[fhw] ;
4: r[] r1 F2 | 13: r[] r3 F1 ;
5: r[] r2 T | 14: r[] r4 T ;
6: while r1 ∧ r2 ≠ 1 | 15: while r3 ∧ r4 ≠ 1 ;
7: (* CS1 *) | 16: (* CS2 *) ;
 f[fhw] | f[fhw] ;
8: w[] F1 0 | L17: w[] F2 0 ;
```

```
let fhw = (po & (_ * F)) ; po
let fre = (rf^-1;co) & ext
irreflexive fhw;fre; fhw;fre
...
```

- Invariance proof unchanged (fence = skip)
- Proved to imply the previous fenceless cat specification
- so  $S_{com}$  unchanged

# consistency proof



# Example: Peterson

- The proof is valid for the **virtual machine** defined by the cat specification Peterson
- Porting the algorithm to a **different machine  $M'$**  just need refencing (and redoing the proof  $M' \Rightarrow H_{cm}$ )
- On machine architecture stronger fences have to be used:
  - SC:  $fhw = \text{no fence}$
  - TSO:  $fhw = \text{mfence}$
  - ARM:  $fhw = \text{dbm} \mid \text{dsb}$

# Conclusion

# Algorithm design methodology

1. Design the algorithm  $A$  and its specification  $S$  in the sequential consistency model of parallelism
2. Consider the anarchic semantics of algorithm  $A$
3. Add communication specifications  $S_{\text{com}}$  to restrict anarchic communications and ensure the correctness of  $A$  with respect to specification  $S$
4. Do the invariance proof under WCM with  $S_{\text{com}}$
5. Infer  $H_{\text{cm}}$  in cat from  $S_{\text{com}}$
6. Prove that the machine memory model  $M$  in cat implies  $H_{\text{cm}}$

# Conclusion

- Modern machines have **complex memory models**
  - ⇒ **portability** has a price (refencing)
  - ⇒ **debugging** is very hard/quasi-impossible
  - ⇒ **proofs** are much harder than with sequential consistency (but still feasible?, mechanically?)
  - ⇒ **static analysis** parameterized by a WCM will be a challenge

# The End, Thank You