

Chair of Software Engineering at ETH Zurich, Switzerland
Software verification course
– Bertrand Meyer, Carlo A. Furia, Sebastian Nanz –
Fall semester 2009

An Informal Introduction to Static Analysis and Verification by Abstract Interpretation

Patrick Cousot

cousot@ens.fr pcousot@cs.nyu.edu
<http://www.di.ens.fr/~cousot> <http://cs.nyu.edu/~pcousot>

Guest course of November 25, 2009

Software Verification, ETH Zurich, Switzerland, 25 November 2009

1

© P. Cousot, <http://se.inf.ethz.ch/teaching/2009-H/itc-0239/index.html#slides>

Content

2.1	Mathematical Semantics
2.2	Mathematical Invariants
2.3	Mathematical Invariant Equations
2.4	Solutions to the Mathematical Invariant Equations
2.5	Solving the Fixpoint Equations by Infinite Iteration
2.6	Machine Invariants
2.7	Interval Abstraction
2.8	An Interval Abstract Interpreter
2.9	Finite but Slow Iteration
2.10	Convergence Speed Up
2.11	Convergence Acceleration
2.11.1	Convergence Acceleration with Widening
2.11.2	Convergence Acceleration with Narrowing
2.12	Chaotic and Structural Iteration
2.13	Verification

2

© P. Cousot, <http://se.inf.ethz.ch/teaching/2009-H/itc-0239/index.html#slides>

Mathematical Semantics

3

© P. Cousot, <http://se.inf.ethz.ch/teaching/2009-H/itc-0239/index.html#slides>

A sample program

Let us start with the following example program.

$$P \triangleq ^1x := 1 ; \text{while } ^2\text{true} \text{ do } ^3x := (x + 1); \text{od}^4.$$

The mathematical semantics of this program can be informally described as follows.

- Execution starts at program point ¹ by assigning 1 to program variable x and goes on at program point ².
- When at program point ² the evaluation of the loop test yields the value `true` so execution continues at program ³ where the value of variable x is incremented by 1 before coming back to ².
- Since the loop condition is never `false`, program point ⁴ is unreachable so program execution never ends.

4

© P. Cousot, <http://se.inf.ethz.ch/teaching/2009-H/itc-0239/index.html#slides>

States

$$P \triangleq {}^1x := 1 ; \text{while } {}^2\text{true} \text{ do } {}^3x := (x + 1) ; \text{od} {}^4.$$

we write $\langle \ell, x \rangle$ for the state of program execution where execution is at program point ℓ , $\ell = 1, 2, 3, 4$, and variable x has integer value $x \in \mathbb{Z}$ (where \mathbb{Z} is the set of all mathematical integers).

5

© P. Cousot, <http://se.inf.ethz.ch/teaching/2009-H/c-0239/index.html#slides>

Trace semantics

$$P \triangleq {}^1x := 1 ; \text{while } {}^2\text{true} \text{ do } {}^3x := (x + 1) ; \text{od} {}^4.$$

So the set of all such execution traces is

$$\{\langle 1, z \rangle \langle ^2, 1 \rangle \langle ^3, 1 \rangle \langle ^2, 2 \rangle \langle ^3, 2 \rangle \dots \langle ^2, i \rangle \langle ^3, i \rangle \langle ^2, i+1 \rangle \dots | z \in \mathbb{Z}\}$$

7

Software Verification, ETH Zurich, Switzerland, 25 November 2009

Execution trace

$$P \triangleq {}^1x := 1 ; \text{while } {}^2\text{true} \text{ do } {}^3x := (x + 1) ; \text{od} {}^4.$$

A complete program execution can be described by the following execution trace which is an infinite sequence of states

$$\langle 1, z \rangle \langle ^2, 1 \rangle \langle ^3, 1 \rangle \langle ^2, 2 \rangle \langle ^3, 2 \rangle \dots \langle ^2, i \rangle \langle ^3, i \rangle \langle ^2, i+1 \rangle \dots$$

where $z \in \mathbb{Z}$ can be any initial integer value of x .

6

© P. Cousot, <http://se.inf.ethz.ch/teaching/2009-H/c-0239/index.html#slides>

Mathematical Invariants

Software Verification, ETH Zurich, Switzerland, 25 November 2009

8

© P. Cousot, <http://se.inf.ethz.ch/teaching/2009-H/c-0239/index.html#slides>

Invariance abstraction

$P \triangleq {}^1x := 1 ; \text{while } {}^2\text{true} \text{ do } {}^3x := (x + 1) ; \text{od} {}^4.$

Let us now consider an abstraction of the set of all possible execution traces, which consists in remembering for each program point ℓ , $\ell = 1, 2, 3, 4$ the set I_ℓ of possible values that can be taken by variable x when execution reaches program point ℓ along any of these traces.

9

Software Verification, ETH Zurich, Switzerland, 25 November 2009
© P. Cousot, <http://se.inf.ethz.ch/teaching/2009-H/c-0239/index.html#slides>

Traces to invariants abstraction

$$\alpha(T) = \lambda l. \{ x \mid \exists \sigma, \sigma' : \sigma \langle l, x \rangle \sigma' \in T \}$$

11

Software Verification, ETH Zurich, Switzerland, 25 November 2009
© P. Cousot, <http://se.inf.ethz.ch/teaching/2009-H/c-0239/index.html#slides>

Invariance semantics

$P \triangleq {}^1x := 1 ; \text{while } {}^2\text{true} \text{ do } {}^3x := (x + 1) ; \text{od} {}^4.$

This set I_ℓ is called a **program local invariant** at program point ℓ . We have

$$\begin{aligned} I_1 &= \mathbb{Z} \\ I_2 &= \{z \in \mathbb{Z} \mid z > 0\} \\ I_3 &= \{z \in \mathbb{Z} \mid z > 0\} \\ I_4 &= \emptyset \end{aligned}$$

10

Software Verification, ETH Zurich, Switzerland, 25 November 2009
© P. Cousot, <http://se.inf.ethz.ch/teaching/2009-H/c-0239/index.html#slides>

Mathematical Invariant Equations

12

Software Verification, ETH Zurich, Switzerland, 25 November 2009
© P. Cousot, <http://se.inf.ethz.ch/teaching/2009-H/c-0239/index.html#slides>

Invariance Equations

$P \triangleq {}^1x := 1 ; \text{while } {}^2\text{true} \text{ do } {}^3x := (x + 1) ; \text{od} {}^4.$

Observe that the set I_ℓ of possible values of variable x at program point $\ell = 1, 2, 3, 4$ satisfies the following conditions.

$$\begin{cases} X_1 = \mathbb{Z} \\ X_2 = \{1\} \cup \{x + 1 \mid x \in X_3\} \\ X_3 = X_2 \cap \{x \in \mathbb{Z} \mid \text{true}\} \\ X_4 = X_2 \cap \{x \in \mathbb{Z} \mid \text{false}\} \end{cases} \quad (3.1)$$

13

© P. Cousot, <http://se.inf.ethz.ch/teaching/2009-Htc-0239/index.html#slides>

Invariance Equations

- At program point 1 the variable x can be initialized by any integer value $z \in \mathbb{Z}$ and so $X_1 = \mathbb{Z}$
- At program point 2 , either execution comes from program point 1 and so the value of variable x is 1 or execution comes from program point 2 and so the value of variable x is the value x that x had at this point 3 incremented by 1. So $X_2 = \{1\} \cup (\{x + 1 \mid x \in X_3\})$.
- At program point 3 , the possible values of x are those at point 2 for which the loop condition is true so $X_3 = X_2 \cap \{x \in \mathbb{Z} \mid \text{true}\} = X_2$.
- At program point 4 , the possible values of x are those at point 2 for which the loop condition is false so $X_4 = X_2 \cap \{x \in \mathbb{Z} \mid \text{false}\} = \emptyset$.

14

Software Verification, ETH Zurich, Switzerland, 25 November 2009

Fixpoint Equations

These conditions can be understood as a system of fixpoint equations $X = f(X)$ of the form

$$\begin{cases} X_i = f_i(X_1, \dots, X_4) \\ i = 1, \dots, 4 \end{cases}$$

with unknowns $X = \langle X_1, \dots, X_4 \rangle$.

15

© P. Cousot, <http://se.inf.ethz.ch/teaching/2009-Htc-0239/index.html#slides>

Fixpoint Solutions

$$\begin{cases} X_1 = \mathbb{Z} \\ X_2 = \{1\} \cup \{x + 1 \mid x \in X_3\} \\ X_3 = X_2 \cap \{x \in \mathbb{Z} \mid \text{true}\} \\ X_4 = X_2 \cap \{x \in \mathbb{Z} \mid \text{false}\} \end{cases} \quad (3.1)$$

- So solving this system of equations might lead to the desired invariant I .
- However these equations do not have a unique solution. For example $X_1 = X_2 = X_3 = \mathbb{Z}$ and $X_4 = \emptyset$ is another solution which is larger for componentwise set inclusion \subseteq .
- So we will prefer the smallest solution (called the *least fixpoint lfp f*), which is included in all other solutions¹ and turns out to be I .

¹by Tarski fixpoint theorem

16

Software Verification, ETH Zurich, Switzerland, 25 November 2009

© P. Cousot, <http://se.inf.ethz.ch/teaching/2009-Htc-0239/index.html#slides>

Solving the Equations by Exhaustive Enumeration

17

Software Verification, ETH Zurich, Switzerland, 25 November 2009

© P. Cousot, <http://se.inf.ethz.ch/teaching/2009-H/itc-0239/index.html#slides>

Solving the equations iteratively ...

The least solution $l = \text{lfp } f$ of $X = f(X)$ for \subseteq can be calculated iteratively, essentially by enumeration of all possible states reachable from the initial states.

18

Software Verification, ETH Zurich, Switzerland, 25 November 2009

Solving the equations iteratively ... (Cont'd)

- $X^0 = \langle X_1^0, X_2^0, X_3^0, X_4^0 \rangle = \langle \emptyset, \emptyset, \emptyset, \emptyset \rangle$ {starting with the smallest possible approximation}
- $X^1 = \langle X_1^1, X_2^1, X_3^1, X_4^1 \rangle = f(X^0) = \langle \mathbb{Z}, \{1\} \cup \{x + 1 \mid x \in X_3^0\}, X_2^0 \cap \{x \in \mathbb{Z} \mid \text{true}\}, X_2^0 \cap \{x \in \mathbb{Z} \mid \text{false}\} \rangle = \langle \mathbb{Z}, \{1\}, \emptyset, \emptyset \rangle$
- $X^2 = \langle X_1^2, X_2^2, X_3^2, X_4^2 \rangle = f(X^1) = \langle \mathbb{Z}, \{1\} \cup \{x + 1 \mid x \in X_3^1\}, X_2^1 \cap \{x \in \mathbb{Z} \mid \text{true}\}, X_2^1 \cap \{x \in \mathbb{Z} \mid \text{false}\} \rangle = \langle \mathbb{Z}, \{1\}, \{1\}, \emptyset \rangle$
- $X^3 = \langle X_1^3, X_2^3, X_3^3, X_4^3 \rangle = f(X^2) = \langle \mathbb{Z}, \{1\} \cup \{x + 1 \mid x \in X_3^2\}, X_2^2 \cap \{x \in \mathbb{Z} \mid \text{true}\}, X_2^2 \cap \{x \in \mathbb{Z} \mid \text{false}\} \rangle = \langle \mathbb{Z}, \{1, 2\}, \{1\}, \emptyset \rangle$

This calculation can go on like this ad infinitum since each iteration $X^{i+1} = f(X^i)$ of the equations corresponds to an iteration in the program loop and so adds one more possible value of variable x at program point ². The solution is to use mathematical induction which requires to invent the following inductive hypothesis

19

Software Verification, ETH Zurich, Switzerland, 25 November 2009

© P. Cousot, <http://se.inf.ethz.ch/teaching/2009-H/itc-0239/index.html#slides>

Solving the equations iteratively ... (Cont'd)

- $X^{2n} = \langle X_1^{2n}, X_2^{2n}, X_3^{2n}, X_4^{2n} \rangle = \langle \mathbb{Z}, \{1, \dots, n\}, \{1, \dots, n\}, \emptyset \rangle$ {induction hypothesis which holds for the basis $n = 1$ }
- $X^{2n+1} = \langle X_1^{2n+1}, X_2^{2n+1}, X_3^{2n+1}, X_4^{2n+1} \rangle = f(X^{2n}) = \langle \mathbb{Z}, \{1\} \cup \{x + 1 \mid x \in X_3^{2n}\}, X_2^{2n} \cap \{x \in \mathbb{Z} \mid \text{true}\}, X_2^{2n} \cap \{x \in \mathbb{Z} \mid \text{false}\} \rangle = \langle \mathbb{Z}, \{1, \dots, n+1\}, \{1, \dots, n\}, \emptyset \rangle$
- $X^{2n+2} = \langle X_1^{2n+2}, X_2^{2n+2}, X_3^{2n+2}, X_4^{2n+2} \rangle = f(X^{2n+1}) = \langle \mathbb{Z}, \{1\} \cup \{x + 1 \mid x \in X_3^{2n+1}\}, X_2^{2n+1} \cap \{x \in \mathbb{Z} \mid \text{true}\}, X_2^{2n+1} \cap \{x \in \mathbb{Z} \mid \text{false}\} \rangle = \langle \mathbb{Z}, \{1, \dots, n+1\}, \{1, \dots, n+1\}, \emptyset \rangle$
- By recurrence on n , we have proved that

$$\forall n : X^{2n} = \langle X_1^{2n}, X_2^{2n}, X_3^{2n}, X_4^{2n} \rangle = \langle \mathbb{Z}, \{1, \dots, n\}, \{1, \dots, n\}, \emptyset \rangle$$

20

Software Verification, ETH Zurich, Switzerland, 25 November 2009

© P. Cousot, <http://se.inf.ethz.ch/teaching/2009-H/itc-0239/index.html#slides>

Solving the equations iteratively ... (Cont'd)

- Passing to the limit, we get the desired strongest invariant

$$\begin{aligned} I &= \langle I_1, I_2, I_3, I_4 \rangle && \{\text{invariants}\} \\ &= \lim_{n \rightarrow +\infty} X^{2^n} \\ &= \langle \mathbb{Z}, \{n \in \mathbb{Z} \mid n > 0\}, \{n \in \mathbb{Z} \mid n > 0\}, \emptyset \rangle \end{aligned}$$

21

© P. Cousot, <http://se.inf.ethz.ch/teaching/2009-H/c-0239/index.html#slides>

f is increasing

- A fundamental property of the invariants equations $X = f(X)$ is that f is *increasing*.
- This means that if $X \dot{\subseteq} Y$ then $f(X) \dot{\subseteq} f(Y)$ where $\langle X_1, \dots, X_n \rangle \dot{\subseteq} \langle Y_1, \dots, Y_n \rangle$ if and only if $\forall i \in [1, n] : X_i \subseteq Y_i$.
- The intuition is that if more states can be reached at some program point then more states will be reachable at next program point.
- It follows that the iterates form an ascending chain meaning $X^0 \dot{\subseteq} X^1 \dot{\subseteq} \dots \dot{\subseteq} X^n \dot{\subseteq} X^{n+1} \dot{\subseteq} \dots \dot{\subseteq} \lim_{n \rightarrow +\infty} X^n = \text{lfp } f$.

22

© P. Cousot, <http://se.inf.ethz.ch/teaching/2009-H/c-0239/index.html#slides>

Machine Invariants

23

© P. Cousot, <http://se.inf.ethz.ch/teaching/2009-H/c-0239/index.html#slides>

Machine Integers

- No computer can represent any, arbitrary large, integer. In practice integer variables like x take their values in an interval $[\min_int, \max_int]$ where $\min_int < 0 < \max_int$ are machine dependant².
- It follows that we have to decide what happens in case of overflow when evaluating expression $(x + 1)$.
- We will assume that execution immediately stops in case of integer overflow³.

²e.g. in two's complement representation on 64 bits, we have generally have $\min_int = -2147483648$ and $\max_int = 2147483647$.

³Which is a rather simplifying hypothesis since most computers will go on providing a result modulo \max_int so that e.g. $\max_int + 1 = \min_int$ in two's complement representation.

24

© P. Cousot, <http://se.inf.ethz.ch/teaching/2009-H/c-0239/index.html#slides>

Machine states and execution traces

Hence the set of program states $\mathcal{S} \triangleq \{1, 2, 3, 4\} \times [\text{min_int}, \text{max_int}]$ is now finite and the execution traces are now finite of the form

$$\{\langle^1, z\rangle\langle^2, 1\rangle \dots \langle^2, i\rangle\langle^3, i\rangle\langle^2, i+1\rangle \dots \langle^3, \text{max_int}\rangle \mid z \in [\text{min_int}, \text{max_int}]\}.$$

Convergence

- Now the convergence of the iterations is guaranteed but is so slow that it cannot be of any practical use, but for programs with very few program variables.
- Moreover, mathematical sets of integers can be arbitrarily complex hence very expensive to represent in computer memory which is likely to produce memory overflows after lengthy computations, a flaw of all program verification methods based upon the exhaustive enumeration of all possible cases.

Machine Invariant Equations

$$P \triangleq {}^1x := 1 ; \text{while } {}^2\text{true} \text{ do } {}^3x := (x + 1); \text{od} {}^4.$$

It follows that the machine invariant satisfies the following equations

$$\begin{cases} X_1 &= [\text{min_int}, \text{max_int}] \\ X_2 &= \{1\} \cup \{x + 1 \in [\text{min_int}, \text{max_int}] \mid x \in X_3\} \\ X_3 &= X_2 \cap \{x \in [\text{min_int}, \text{max_int}] \mid \text{true}\} \\ X_4 &= X_2 \cap \{x \in [\text{min_int}, \text{max_int}] \mid \text{false}\} \end{cases} \quad (3.2)$$

Interval Abstraction

Interval Abstraction

- A further abstraction must be used to solve the machine invariant computer representation problem.
- We will use intervals $[l, h] \triangleq \{x \in \mathbb{Z} \mid l \leq x \leq h\}$ with the convention that $[l, h] = \emptyset$ whenever $h < l$.
- In doing so we perform an approximation of a non-empty set $X \subseteq [\min_int, \max_int]$ by the interval $[\min X, \max X]$.
- This approximation is sound in that whenever the value of variable x belongs to a set X_i whenever execution reaches program point i , it definitely also belongs to the set $[\min X_i, \max X_i]$.
- This information is certainly correct but just less precise.

29

© P. Cousot, <http://se.inf.ethz.ch/teaching/2009-H/c-0239/index.html#slides>

Traces to intervals abstraction

$$\alpha(T) = \lambda l. \{ [\min x, \max x] \mid \exists \sigma, \sigma': \sigma \langle l, x \rangle \sigma' \in T \}$$

30

© P. Cousot, <http://se.inf.ethz.ch/teaching/2009-H/c-0239/index.html#slides>

Interval Invariance Equations

$$P \triangleq {}^1x := 1 ; \text{while } {}^2\text{true} \text{ do } {}^3x := (x + 1); \text{od} {}^4.$$

The interval invariance equations are now

$$\left\{ \begin{array}{l} X_1 = [\min_int, \max_int] \\ X_2 = [1, 1] \sqcup (\exists X_3 = \emptyset \text{ ? } \emptyset : \text{let } [a, b] = X_3 \text{ in } [\min(a + 1, \max_int), \min(b + 1, \max_int)])) \\ X_3 = X_2 \sqcap [\min_int, \max_int] \\ X_4 = X_2 \sqcap \emptyset \end{array} \right.$$

31

Software Verification, ETH Zurich, Switzerland, 25 November 2009

Interval Operations

- the interval join is $\emptyset \sqcup \emptyset \triangleq \emptyset$, $\emptyset \sqcup [l, h] \triangleq [l, h] \sqcup \emptyset \triangleq [l, h]$, and
$$[a, b] \sqcup [c, d] \triangleq [\min(a, c), \max(b, d)]$$
- and the interval meet is $\emptyset \sqcap \emptyset \triangleq \emptyset$, $\emptyset \sqcap [l, h] \triangleq [l, h] \sqcap \emptyset \triangleq \emptyset$, and
$$\begin{aligned} [a, b] \sqcap [c, d] &\triangleq [\max(a, c), \min(b, d)] && \text{when } b \geq c \wedge d \geq a \\ [a, b] \sqcap [c, d] &\triangleq \emptyset && \text{when } b < c \vee d < a \end{aligned}$$

32

Software Verification, ETH Zurich, Switzerland, 25 November 2009

Over-approximation

- The interval equations over-estimate the machine invariant in than they will provide in general more states that possible in actual program executions.
- For example the set $\{1, 2, 5\}$ will be overapproximated by $[1, 5]$ which introduces the spurious values 3 and 4.
- Notice that overapproximation preserve invariance. For example if the values of variable x are always greater than one at some program point then they are certainly positive (although the value 0 is spurious).

33

© P. Cousot, <http://se.inf.ethz.ch/teaching/2009-H/c-0239/index.html#slides>

An Interval Abstract Interpreter

35

© P. Cousot, <http://se.inf.ethz.ch/teaching/2009-H/c-0239/index.html#slides>

Example of incorrect approximations

For $x \in \{1, 2, 5\}$

- Underapproximations (such as x are always greater than 10) would be incorrect.
- Similarly, incomparable approximations (such as x is negative) are also unsound.

In particular the interval join \sqcup overapproximates the interval union \cup and the interval meet \sqcap overapproximates the interval intersection \cap .

Objective

- We now briefly sketch the design and functional encoding in OCAML of the interval abstract interpreter.
- Such an interval abstract interpreter reads any program, builds the interval invariance equations, and then solve them.
- For simplicity, we concentrate on the second part and will provide encodings of the interval invariance equations manually.

34

Software Verification, ETH Zurich, Switzerland, 25 November 2009

36

© P. Cousot, <http://se.inf.ethz.ch/teaching/2009-H/c-0239/index.html#slides>

The Interval Abstract Domain

- We first encode the interval abstract domain, implementing a computer representation of abstract interval properties with a type `interval` (where `EMPTY` encodes the empty set \emptyset). In OCaml, we have `max_int = 1073741823` and `min_int = -1073741824`⁴.
- We also encode the basic interval operations \sqsubseteq (less, interval inclusion), \sqcup (interval join), \sqcap (interval meet), interval printing (`print`) and interval incrementation (`add1`).
- Of course many more interval operations are needed to handle a full language, but we aim at extreme simplicity.

⁴One of the 64 bits is used for garbage collection.

or `max_int = 4611686018427387903` depending on the machine/compiler

```
(* interval.ml, interval abstract domain *)
type interval = EMPTY | INT of (int * int);;
let less x y = match x,y with
| EMPTY, _ -> true
| _, EMPTY -> false
| INT (a,b), INT (c,d) -> (c<=a)&&(b<=d);;
let greater x y = less y x;;
let join x y = match x,y with
| EMPTY, _ -> y
| _, EMPTY -> x
| INT (a,b), INT (c,d) -> INT (min a c,max b d);;
let meet x y = match x,y with
| EMPTY, _ -> EMPTY
| _, EMPTY -> EMPTY
| INT (a,b), INT (c,d) ->
  if (b<c) or (d<a) then EMPTY
  else INT (max a c,min b d);;
let add1 x = match x with
| EMPTY -> EMPTY
| INT (a,b) ->
  (INT ((if a<max_int then a+1 else max_int),
        (if b<max_int then b+1 else max_int))),;
let print x = match x with
| EMPTY -> print_string "_|_"
| INT (a,b) -> print_string "("; print_int a;
                  print_string ","; print_int b; print_string ")" ";";
```

Abstract Environments

- For programs with more than one variable, we would have to encode an abstract environment assigning intervals to program variables.
- Writing $X = \{x_1 \leftarrow v_1, \dots, x_n \leftarrow v_n\}$ for the function X mapping x_i to v_i such that $X(x_i) = v_i, i = 1, \dots, n$, the interval invariance equations would be

$$\begin{cases} X_1 &= \{x \leftarrow [\min_int, \max_int]\} \\ X_2 &= \{x \leftarrow [1, 1] \sqcup (\exists X_3(x) = \emptyset \Rightarrow \emptyset : \text{let } [a, b] = X_3(x) \text{ in } [\min(a + 1, \max_int), \min(b + 1, \max_int)])\} \\ X_3 &= X_2 \sqcap \{x \leftarrow [\min_int, \max_int]\} \\ X_4 &= X_2 \sqcap \{x \leftarrow \emptyset\} \end{cases}$$

where the abstract operations are extended pointwise such as $\{x_1 \leftarrow v_1, \dots, x_n \leftarrow v_n\} \sqcap \{x_1 \leftarrow v'_1, \dots, x_n \leftarrow v'_n\} \triangleq \{x_1 \leftarrow v_1 \sqcap v'_1, \dots, x_n \leftarrow v_n \sqcap v'_n\}$.

- Since our example has only one variable, this boils down to using the interval abstract domain (and leaving implicit the variable name `x`).

Abstract Invariants

- Then we have to encode an abstract domain for representing abstract invariants $\langle X^1, X^2, X^3, X^4 \rangle$ which attach to each program point i an abstract local invariant X^i which holds whenever controls reaches program point i .
- Each abstract local invariant X^i is represented by an abstract environment (abstract intervals in our simplified case).
- The encoding is very simple as a 4-tuple specifying the value of program variable x at each program point $(^1, ^2, ^3, ^4)$.

Abstract Invariants (Cont'd)

We essentially have to represent the logical structure, which boils down to

- the partial order $\dot{\sqsubseteq}$ (pless), encoding abstract implication (\sqsubseteq in set theory and \Rightarrow in logic);
- $\dot{\sqsupseteq}$ (pgreater), the abstract inverse implication (\sqsupseteq in set theory and \Leftarrow in logic);
- the pointwise infimum (\emptyset)⁴ (pbot), the abstract encoding of `false`,
- the pointwise meet (for later use), and
- the printing of local abstract invariants attached to program points (pprint).

```
(* invariant.ml, interval invariant abstract domain *)
open Interval
type invariant = interval*interval*interval*interval;;
let cless (x1,x2,x3,x4) (x'1,x'2,x'3,x'4) =
  (less x1 x'1, less x2 x'2, less x3 x'3, less x4 x'4);;
let pless x x' =
  let (b1, b2, b3, b4) = cless x x' in
  b1 && b2 && b3 && b4;;
let pgreater x x' = pless x' x;;
let pbot = (EMPTY, EMPTY, EMPTY, EMPTY);;
let pmeet (x1,x2,x3,x4) (x'1,x'2,x'3,x'4) =
  (meet x1 x'1, meet x2 x'2, meet x3 x'3, meet x4 x'4);;
let pprint (x1,x2,x3,x4) =
  print_string " 1:";print x1; print_string " 2:";
  print x2; print_string " 3:";print x3;
  print_string " 4:";print x4; print_newline ();;
```

The Iterator

- Next the iterator module implements the iterative computation of the least solution of the invariance equations (lfp⁵).
- It is parameterized by the order (leq), the starting point (a) and the abstract transformer (f) so as to compute a, f(a), f²(a), ..., fⁿ(a), ..., until reaching the limit f^ℓ(a) such that f(f^ℓ(a)) ⊑ f^ℓ(a).
- Of course, convergence may not be guaranteed in which case lfp does not terminate (or terminates with a runtime error, e.g. out of memory).

⁵least fixpoint.

```
(* iterator.ml, iteration of f from a to x >= f(x) *)
let lfp leq a f =
  let rec iterate x =
    let y = f x in
    if leq y x then x
    else iterate y
  in iterate a;;
```

Jacobi versus chaotic iteration strategies

Of course the Jacobi iteration strategy

$$\begin{cases} X_i^{k+1} = f_i(X_1^k, \dots, X_4^k) & k = 1, 2, 3, \dots \\ i = 1, \dots, 4 \end{cases}$$

is simplistic, more elaborate ones would use e.g. a working list

The Abstract Interpreter

The abstract interpreter performs the iterative abstract reachability fixpoint computation and prints the least fixpoint result.

```
(* reachability interval analysis *)
open Invariant
open TransformerUnbounded
open Iterator
let analyzer () = pprint (lfp pless pbot f);
analyzer ();;
```

Abstract Invariant Equations $X=f(X)$

Then we encode the abstract reachable state transformer $f(X) = f(\langle X_1, \dots, X_4 \rangle)$ using the environment abstract domain (the intervals in our simplified case).

```
(* transformerUnbounded.ml, abstract transformer *)
open Interval
open Invariant
let f1 () = INT (min_int,max_int);;
let f2 x1 x3 = join (INT (1,1)) (add1 x3);;
let f3 x2 = meet x2 (INT (min_int,max_int));;
let f4 x2 = meet x2 EMPTY;;
let f (x1,x2,x3,x4) = (f1 (), f2 x1 x3, f3 x2, f4 x2);;
```

encoding: $\begin{cases} X_1 = [\min_int, \max_int] \\ X_2 = [1, 1] \sqcup (\exists X_3 = \emptyset \text{ ? } \emptyset : \text{let } [a, b] = X_3 \text{ in } [\min(a + 1, \max_int), \min(b + 1, \max_int)] \sqcup) \\ X_3 = X_2 \sqcap [\min_int, \max_int] \\ X_4 = X_2 \sqcap \emptyset \end{cases}$

Infinitary Iteration

Iterative Resolution of the Interval Equations

Because the abstract domains are finite, the static analysis will always terminate. In our case, after more than 40mn of computation⁶, we get

```
% ocamlc interval.ml invariant.ml transformeUnbounded.ml iterator.ml \
? reachability_unbounded.ml
% time ./a.out
1:(-1073741824,1073741823) 2:(1,1073741823) 3:(1,1073741823) 4:_|_
2977.460u 9.632s 50:43.46 98.1% 0+0k 0+0io 0pf+0w
%
```

⁶On a MacBook Pro with Intel Core 2 Duo at 2.6 GHz.

A look at the iterates...

The Jacobi iterates are as follows

```
% ocamlc interval.ml invariant.ml transformerUnbounded.ml \
? iteratorPartialUnboundedTrace.ml reachability_unbounded_trace.ml
% time ./a.out
1:_|_ 2:_|_ 3:_|_ 4:_|_
1:(-1073741824,1073741823) 2:(1,1) 3:_|_ 4:_|_
1:(-1073741824,1073741823) 2:(1,1) 3:(1,1) 4:_|_
1:(-1073741824,1073741823) 2:(1,2) 3:(1,1) 4:_|_
1:(-1073741824,1073741823) 2:(1,2) 3:(1,2) 4:_|_
1:(-1073741824,1073741823) 2:(1,3) 3:(1,2) 4:_|_
1:(-1073741824,1073741823) 2:(1,3) 3:(1,3) 4:_|_
...
...
1:(-1073741824,1073741823) 2:(1,1073741821) 3:(1,1073741820) 4:_|_
1:(-1073741824,1073741823) 2:(1,1073741821) 3:(1,1073741821) 4:_|_
1:(-1073741824,1073741823) 2:(1,1073741822) 3:(1,1073741821) 4:_|_
1:(-1073741824,1073741823) 2:(1,1073741822) 3:(1,1073741822) 4:_|_
1:(-1073741824,1073741823) 2:(1,1073741823) 3:(1,1073741822) 4:_|_
1:(-1073741824,1073741823) 2:(1,1073741823) 3:(1,1073741823) 4:_|_
converged to lfp.
1:(-1073741824,1073741823) 2:(1,1073741823) 3:(1,1073741823) 4:_|_
3115.012u 7.706s 52:49.34 98.5% 0+0k 0+0io 0pf+0w
%
```

On the Convergence Criterion

- Notice that the abstract invariance equations $X = f(X)$ are increasing, if $X \dot{\subseteq} Y$ then $f(X) \dot{\subseteq} f(Y)$.
 - The intuition is that the interval of possible value of a variable is larger at a program point, it should be also larger at the next program point.
 - It follows that the iterates $X^0 \dot{\subseteq} \dots \dot{\subseteq} X^n \dot{\subseteq} \dots \dot{\subseteq} \lim_{n \rightarrow +\infty} X^n$ are increasing.
 - Since the abstract interpreter stops iterating when reaching of postfixpoint $f(\lim_{n \rightarrow +\infty} X^n) \dot{\subseteq} \lim_{n \rightarrow +\infty} X^n$, the limit satisfies $f(\lim_{n \rightarrow +\infty} X^n) = \lim_{n \rightarrow +\infty} X^n$ by antisymmetry.

On Slow Convergence !

Of course the convergence is extremely slow and in practice must be accelerated.

Convergence Speed Up

Software Verification, ETH Zurich, Switzerland, 25 November 2009

53

© P. Cousot, <http://se.inf.ethz.ch/teaching/2009-H/c-0239/index.html#slides>

Convergence speed-up

Instead of iterating a function f , one can iterate its powers like f^2 or f^{64} .

If the compiler can generate code for computing $f^n(x)$ which is faster than applying n times f to x , one can hope for a convergence speed up.

This would be the case for compilers able to implement f^n by expanding n times the code of f thus saving the intrinsic cost of $n - 1$ function calls.

Software Verification, ETH Zurich, Switzerland, 25 November 2009

54

© P. Cousot, <http://se.inf.ethz.ch/teaching/2009-H/c-0239/index.html#slides>

Inconclusive experiment

The following experiment

in OCAML

```
(* reachability interval analysis with speed up *)
open Invariant
open TransformerUnbounded
open Iterator
let o f g x = f (g x);;
let i x = x;;
let even n = n / 2 * 2 = n;;
let rec fn n = if n = 1 then f
               else if (even n) then o (fn (n / 2)) (fn (n / 2))
               else o (fn (n / 2)) (fn ((n + 1) / 2));;
let analyzer () = pprint (lfp pless pbot (fn 64));
analyzer ();;
```

is inconclusive.

```
% ocamlc interval.ml invariant.ml transformerUnbounded.ml \
? iterator.ml reachability_unboundedx2.ml
% time ./a.out
1:(-1073741824,1073741823) 2:(1,1073741823) 3:(1,1073741823) 4:_I_
2429.727u 5.740s 40:59.82 99.0% 0+0k 0+0io 0pf+0w
%
```

Software Verification, ETH Zurich, Switzerland, 25 November 2009

55

© P. Cousot, <http://se.inf.ethz.ch/teaching/2009-H/c-0239/index.html#slides>

Convergence Acceleration

Software Verification, ETH Zurich, Switzerland, 25 November 2009

56

© P. Cousot, <http://se.inf.ethz.ch/teaching/2009-H/c-0239/index.html#slides>

Objective

- When convergence requires infinitely many steps or is very slow, it may not be possible, due to undecidability or high complexity, to exactly calculate the least solution to the abstract system of equations.^(*)
- The only sound solution is then to have overapproximations of the desired result.
- We have already exploited the overapproximation idea when replacing sets of integer values in the invariant equations by interval of values.
- We now exploit the approximation idea a second time now while computing the solution of the invariance equations.
- The possibility of computing sound but approximate solutions to the invariant equations leads to powerful sound and fast static program analysis methods.

(*) Of course direct solutions do sometimes exist e.g. linear equations on regular languages

57

Software Verification, ETH Zurich, Switzerland, 25 November 2009

© P. Cousot, <http://se.inf.ethz.ch/teaching/2009-H/itc-0239/index.html#slides>

Convergence Acceleration by Widening

Convergence Acceleration by Widening

- The intuition for convergence acceleration is to speed up the increasing iteration $X^0 = \perp, \dots, X^{n+1} = f(X^n), \dots, \lim_{n \rightarrow +\infty} X^n$ so as to reach an overapproximation \hat{A} of the least solution $\lim_{n \rightarrow +\infty} X^n$ of the fixpoint equation $X = f(X)$.⁷
- Convergence acceleration means that X^{n+1} will be a function of X^n and $f(X^n)$.⁸ and so $X^{n+1} = X^n \nabla f(X^n)$ where ∇ is called a *widening*.⁹

⁷The justification is again by Tarski theorem since $f(\hat{A}) \sqsubseteq \hat{A}$ implies $\text{lfp } f \sqsubseteq \hat{A}$.

⁸and more generally X^{n+1} could depend on the sequence of previous iterates $X^0, f(X^0), \dots, X^n, f(X^n)$

⁹We use a binary operator notation rather than a functional notation because of the analogy between widenings ∇ and joins \vee, \sqcup , etc

59

Software Verification, ETH Zurich, Switzerland, 25 November 2009

© P. Cousot, <http://se.inf.ethz.ch/teaching/2009-H/itc-0239/index.html#slides>

Soundness

- For soundness, the widening must perform over-approximations, that is $x \sqsubseteq x \nabla y$ and $y \sqsubseteq x \nabla y$.

Software Verification, ETH Zurich, Switzerland, 25 November 2009

58

© P. Cousot, <http://se.inf.ethz.ch/teaching/2009-H/itc-0239/index.html#slides>

60

Software Verification, ETH Zurich, Switzerland, 25 November 2009

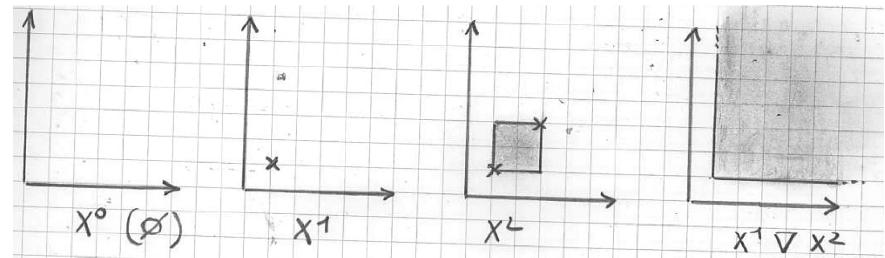
© P. Cousot, <http://se.inf.ethz.ch/teaching/2009-H/itc-0239/index.html#slides>

Convergence enforcement

- For convergence, the widening must ensure termination with an overapproximation of the desired solution.

Example: Interval Widening (Cont'd)

The extrapolation of bounds to infinity is illustrated on the following iteration (for two variables).



Example: Interval Widening

For example, a widening for intervals could be

$$\emptyset \nabla y \triangleq y$$

$$x \nabla \emptyset \triangleq x$$

$$[a, b] \nabla [c, d] \triangleq [\{c < a \Rightarrow -\infty : a\}, \{d > b \Rightarrow +\infty : b\}]$$

- Recall than in $x \nabla y$ the x is an iterate and y is the next iterate $f(x)$. So in $[a, b] \nabla [c, d]$ if $c < a$ the next iterate decreases the lower limit of the interval so widening to $-\infty$ ensures this cannot happen infinitely often.

- Similarly, if $d > b$ then the next iterate increases the upper limit of the interval so widening to $+\infty$ ensures this cannot happen infinitely often. Moreover the widened interval is larger which ensures that we perform an overapproximation.

Widenings are not increasing!

- Observe than the interval widening is not increasing. For example $[0, 1] \sqsubseteq [0, 2]$ but $[0, 1] \nabla [0, 2] = [0, +\infty] \not\sqsubseteq [0, 2] = [0, 2] \nabla [0, 2]$
- It can be shown that if the widening stops loosing information when a solution is found and is increasing then it cannot enforce termination

Encoding the interval widening

A functional encoding in of the widening in OCAML could be

```
(* intervalWidening.ml, interval widening *)
open Interval
let widen x y = match x,y with
| EMPTY, _ -> y
| _, EMPTY -> x
| INT (a,b), INT (c,d) ->
    let a' = if c<a then min_int else a in
    let b' = if d>b then max_int else b in
        INT (a',b');
```

65

© P. Cousot, <http://se.inf.ethz.ch/teaching/2009-H/c-0239/index.html#slides>

Environment Widening

If we had abstract environments to handle several variables, the widening would have to be applied individually for each of these variables.

Software Verification, ETH Zurich, Switzerland, 25 November 2009

66

© P. Cousot, <http://se.inf.ethz.ch/teaching/2009-H/c-0239/index.html#slides>

Invariant widening

$P \triangleq ^1x := 1 ; \text{while } ^2\text{true} \text{ do } ^3x := (x + 1); \text{ od}^4.$

We must also extend the widening to local invariants attached to program points. In our example, the widening is applied once around the loop at program point ² as follows.

```
(* invariantWidening.ml, invariant widening *)
open IntervalWidening
let pwiden (x1,x2,x3,x4) (x'1,x'2,x'3,x'4) =
    (x'1,widen x2 x'2,x'3,x'4);;
```

67

Software Verification, ETH Zurich, Switzerland, 25 November 2009

Abstract Interpreter with Widening

The abstract interpreter now calls the iterator using the invariant widening.

```
(* reachability analysis with widening *)
open Invariant
open InvariantWidening
open TransformerUnbounded
open Iterator
let analyzer () =
    let fw x = pwiden x (f x) in
        pprint (lfp pless pbot fw);;
    analyzer ();;
```

68

Software Verification, ETH Zurich, Switzerland, 25 November 2009

© P. Cousot, <http://se.inf.ethz.ch/teaching/2009-H/c-0239/index.html#slides>

Static Analysis with Widening

The result is now almost instantaneous.

```
% ocamlc interval.ml intervalWidening.ml invariant.ml \
? invariantWidening.ml transformerUnbounded.ml iterator.ml \
? reachability_widening.ml
% time ./a.out
1:(-1073741824,1073741823) 2:(1,1073741823) 3:(1,1073741823) 4:_I_
0.000u 0.000s 0:00.00 0.0%      0+0k 0+0io 0pf+0w
%
```

Trace of the Iterations with Widening

The Jacobi iterates with widening are extremely fast as shown below.

```
% ocamlc interval.ml intervalWidening.ml invariant.ml \
? invariantWidening.ml transformerUnbounded.ml \
? iteratorTrace.ml reachability_widening_trace.ml
% time ./a.out
1:_I_ 2:_I_ 3:_I_ 4:_I_
1:(-1073741824,1073741823) 2:(1,1) 3:_I_ 4:_I_
1:(-1073741824,1073741823) 2:(1,1) 3:(1,1) 4:_I_
1:(-1073741824,1073741823) 2:(1,1073741823) 3:(1,1) 4:_I_
1:(-1073741824,1073741823) 2:(1,1073741823) 3:(1,1073741823) 4:_I_
converged to lfp.
1:(-1073741824,1073741823) 2:(1,1073741823) 3:(1,1073741823) 4:_I_
0.000u 0.000s 0:00.00 0.0%      0+0k 0+0io 0pf+0w
%
```

Imprecision of the Widening

Of course, the widening cannot, in general, provide the exact result! To see that, consider the bounded iteration

$$P \triangleq {}^1x := 1 ; \text{while } {}^2(x \leq 100) \text{ do } {}^3x := (x + 1); \text{ od} {}^4.$$

so that the abstract interval equations become

$$\begin{cases} X_1 &= \{x \leftarrow [\min_int, \max_int]\} \\ X_2 &= \{x \leftarrow [1, 1] \sqcup (\exists X_3(x) = \emptyset \Rightarrow \emptyset : \text{let } [a, b] = X_3(x) \text{ in } [\min(a + 1, \max_int), \min(b + 1, \max_int)]\}) \\ X_3 &= X_2 \dot{\cap} \{x \leftarrow [\min_int, 100]\} \\ X_4 &= X_2 \dot{\cap} \{x \leftarrow [101, \max_int]\} \end{cases}$$

I) Direct iteration (without widening)

A direct iteration

```
(* reachability interval analysis *)
open Invariant
open TransformerBounded
open Iterator
let analyzer () = pprint (lfp pless pbot f);;
analyzer ();;
```

yields

```
% ocamlc interval.ml invariant.ml transformerBounded.ml \
? iterator.ml reachability_bounded.ml
% time ./a.out
1:(-1073741824,1073741823) 2:(1,101) 3:(1,100) 4:(101,101)
0.001u 0.000s 0:00.00 0.0%      0+0k 0+0io 0pf+0w
%
```

in more details

```
% ocamlc interval.ml invariant.ml transformerBounded.ml \
? iteratorPartialBoundedTrace.ml reachability_bounded_trace.ml
% time ./a.out
1:_|_ 2:_|_ 3:_|_ 4:_|_
1:(-1073741824,1073741823) 2:(1,1) 3:_|_ 4:_|_
1:(-1073741824,1073741823) 2:(1,1) 3:(1,1) 4:_|_
1:(-1073741824,1073741823) 2:(1,2) 3:(1,1) 4:_|_
1:(-1073741824,1073741823) 2:(1,2) 3:(1,2) 4:_|_
1:(-1073741824,1073741823) 2:(1,3) 3:(1,2) 4:_|_
...
1:(-1073741824,1073741823) 2:(1,99) 3:(1,99) 4:_|_
1:(-1073741824,1073741823) 2:(1,100) 3:(1,99) 4:_|_
1:(-1073741824,1073741823) 2:(1,100) 3:(1,100) 4:_|_
1:(-1073741824,1073741823) 2:(1,101) 3:(1,100) 4:_|_
1:(-1073741824,1073741823) 2:(1,101) 3:(1,100) 4:(101,101)
converged to lfp.
1:(-1073741824,1073741823) 2:(1,101) 3:(1,100) 4:(101,101)
0.001u 0.001s 0:00.00 0.0%      0+0k 0+0io 0pf+0w
%
```

Again convergence is guaranteed but slow.

73

© P. Cousot, <http://se.inf.ethz.ch/teaching/2009-H/itc-0239/index.html#slides>

In more details the widening effect is not compensated by the test on loop exit.

```
? reachability_widening_bounded_trace.ml
% time ./a.out
1:_|_ 2:_|_ 3:_|_ 4:_|_
1:(-1073741824,1073741823) 2:(1,1) 3:_|_ 4:_|_
1:(-1073741824,1073741823) 2:(1,1) 3:(1,1) 4:_|_
1:(-1073741824,1073741823) 2:(1,1073741823) 3:(1,1) 4:_|_
1:(-1073741824,1073741823) 2:(1,1073741823) 3:(1,100) 4:(101,1073741823)
converged to lfp.
1:(-1073741824,1073741823) 2:(1,1073741823) 3:(1,100) 4:(101,1073741823)
0.000u 0.000s 0:00.00 0.0%      0+0k 0+0io 0pf+0w
%
```

75

© P. Cousot, <http://se.inf.ethz.ch/teaching/2009-H/itc-0239/index.html#slides>

II) Iteration with widening

```
(* reachability analysis with widening *)
open Invariant
open InvariantWidening
open TransformerBounded
open Iterator
let analyzer () =
  let fw x = pwiden x (f x) in
  pprint (lfp pless pbot fw);;
analyzer ();;
```

we rapidly get a strictly less precise result.

```
% ocamlc interval.ml intervalWidening.ml invariant.ml \
? invariantWidening.ml transformerBounded.ml iterator.ml \
? reachability_widening_bounded.ml
% time ./a.out
1:(-1073741824,1073741823) 2:(1,1073741823) 3:(1,100) 4:(101,1073741823)
0.000u 0.000s 0:00.00 0.0%      0+0k 0+0io 0pf+0w
%
% ocamlc interval.ml intervalWidening.ml invariant.ml \
? invariantWidening.ml transformerBounded.ml iteratorTrace.ml \
```

74

© P. Cousot, <http://se.inf.ethz.ch/teaching/2009-H/itc-0239/index.html#slides>

Convergence Acceleration by Narrowing

76

© P. Cousot, <http://se.inf.ethz.ch/teaching/2009-H/itc-0239/index.html#slides>

Intuition for Convergence Acceleration with Narrowing

- Because the upward iteration sequence with widening converges to a postfixpoint \hat{A} of f such that $\text{lfp } f \sqsubseteq \hat{A} \wedge f(\hat{A}) \sqsubseteq \hat{A}$, we have, by recurrence and since f is increasing, that $\text{lfp } f \sqsubseteq f^n(\hat{A}) \sqsubseteq \hat{A}$.
- When \hat{A} is not a fixpoint of f , any iterate in the sequence $Y^0 = \hat{A}, \dots, Y^{n+1} = f(Y^n) = f^n(\hat{A})$ is an overapproximation of the unknown $\text{lfp } f$ more precise than \hat{A} .
- However, this downward iteration $\langle Y^n, n \in \mathbb{N} \rangle$ might be infinite or converging slowly.
- It is therefore necessary to ensure its fast convergence. Convergence acceleration means that Y^{n+1} will be a function of Y^n and $f(Y^n)$ ¹⁰ and so $Y^{n+1} = Y^n \Delta f(Y^n)$ where Δ is called a *narrowing*¹¹.

¹⁰and more generally Y^{n+1} could depend on the sequence of previous iterates $Y^0, f(Y^0), \dots, Y^n, f(Y^n)$, as was also the case for widening.

¹¹We use a binary operator notation rather than a functional notation because of the analogy between narrowing Δ and meets \wedge, \sqcap , etc

Soundness

- For soundness, the narrowing must perform over-approximations, that is $y \sqsubseteq x \Delta y$, so as to stay above the unknown least fixpoint, which requires remaining above any fixpoint (which we have no way to distinguish from the least one)¹².

¹²By recurrence, if $X = f(X)$ is any fixpoint of f such that $X \sqsubseteq Y^n$ then $X = f(X) \sqsubseteq f(Y^n)$ since f is increasing so $X \sqsubseteq Y^n \sqsubseteq Y^n \Delta f(Y^n) = Y^{n+1}$ by the overapproximation hypothesis.

Convergence

- For convergence, the narrowing must ensure termination with a fixpoint.

Example: Interval Narrowing

For example, a narrowing for intervals could be

$$\emptyset \Delta y \triangleq \emptyset$$

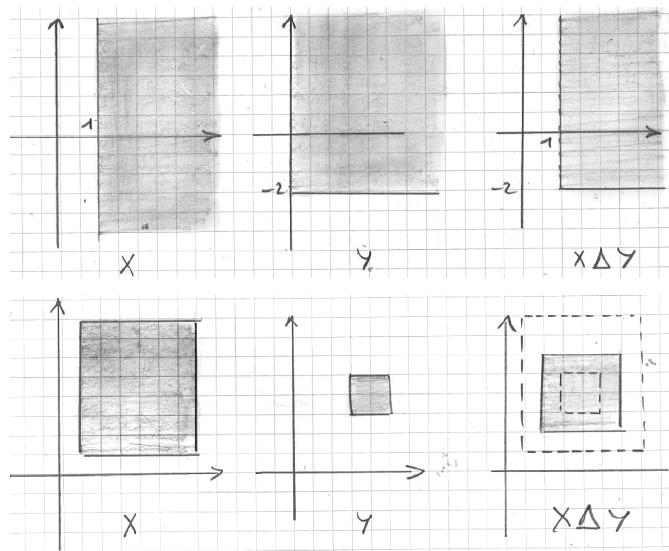
$$x \Delta \emptyset \triangleq \emptyset$$

$$[a, b] \Delta [c, d] \triangleq [\langle a = -\infty ? c : a \rangle, \langle b = +\infty ? d : b \rangle]$$

Recall that in $x \Delta y$ the x is an iterate and y is the next iterate $f(x)$. So $[a, b] \Delta [c, d]$ will just eliminate the infinite bounds in $[a, b]$ and replace them by the bounds of the next iterate $[c, d]$.

So the narrowed interval is larger than $[c, d]$ that is $f(x)$ which ensures that we perform an overapproximation. Because only finitely many bounds can be infinite hence potentially removed, termination is guaranteed.

Example: Interval Narrowing (Cont'd)



Software Verification, ETH Zurich, Switzerland, 25 November 2009

81

© P. Cousot, <http://se.inf.ethz.ch/teaching/2009-Htc-0239/index.html#slides>

Encoding the Interval Narrowing

A functional encoding in OCAML could be

```
(* interval narrowing *)
open Interval
let narrow x y = match x,y with
| EMPTY, _ -> EMPTY
| _, EMPTY -> EMPTY
| INT (a,b), INT (c,d) ->
    let a' = if a=min_int then c else a in
    let b' = if b=max_int then d else b in
    INT (a',b');;
```

Software Verification, ETH Zurich, Switzerland, 25 November 2009

82

© P. Cousot, <http://se.inf.ethz.ch/teaching/2009-Htc-0239/index.html#slides>

Invariant Narrowing

$$P \triangleq {}^1x := 1 ; \text{while } {}^2\text{true} \text{ do } {}^3x := (x + 1); \text{ od} {}^4.$$

In our example, the narrowing is applied once around the loop at program point 2 , like the widening.

```
(* invariantNarrowing.ml, invariant narrowing *)
open IntervalNarrowing
let pnarrow (x1,x2,x3,x4) (x'1,x'2,x'3,x'4) =
  (x'1,narrow x2 x'2,x'3,x'4);;
```

Software Verification, ETH Zurich, Switzerland, 25 November 2009

83

© P. Cousot, <http://se.inf.ethz.ch/teaching/2009-Htc-0239/index.html#slides>

Abstract Interpreter with Widening/Narrowing

The abstract interpreter now calls the iterator using the invariant widening until reaching a postfixpoint and then calls the iterator using the invariant narrowing until reaching a fixpoint.

```
(* reachability analysis with widening and narrowing *)
open Invariant
open InvariantWidening
open InvariantNarrowing
open TransformerBounded
open Iterator
let analyzer () =
  let fw x = pwiden x (f x) in
  let w = (lfp pless pbot fw) in
    let fn x = pnarrow x (f x) in
      pprint (lfp pgreater w fn);;
analyzer ();;
```

Software Verification, ETH Zurich, Switzerland, 25 November 2009

84

© P. Cousot, <http://se.inf.ethz.ch/teaching/2009-Htc-0239/index.html#slides>

Example of convergence acceleration by widening/narrowing

The result is now almost instantaneous.

```
% ocamlc interval.ml intervalWidening.ml intervalNarrowing.ml \
? invariant.ml invariantWidening.ml invariantNarrowing.ml \
? transformerBounded.ml iterator.ml \
? reachability_narrowing_boundeds.ml
% time ./a.out
1:(-1073741824,1073741823) 2:(1,101) 3:(1,100) 4:(101,101)
0.000u 0.000s 0:00.00 0.0%      0+0k 0+0io 0pf+0w
%
```

On the (im)precision of the analysis...

Of course the narrowing cannot always recover all information lost by the widening, in particular because it is blocked by fixpoints jumped over by the widening.

Details of the iteration with Narrowing/Widening

When compared to the Jacobi iterations, the chaotic iterates with widening and narrowing are extremely fast as shown below.

```
% ocamlc interval.ml intervalWidening.ml intervalNarrowing.ml \
? invariant.ml invariantWidening.ml invariantNarrowing.ml \
? transformerBounded.ml iteratorTrace.ml \
? reachability_narrowing_boundeds_trace.ml
% time ./a.out
1:_|_ 2:_|_ 3:_|_ 4:_|_
1:(-1073741824,1073741823) 2:(1,1) 3:_|_ 4:_|_
1:(-1073741824,1073741823) 2:(1,1) 3:(1,1) 4:_|_
1:(-1073741824,1073741823) 2:(1,1073741823) 3:(1,1) 4:_|_
1:(-1073741824,1073741823) 2:(1,1073741823) 3:(1,100) 4:(101,1073741823)
converged to lfp.
1:(-1073741824,1073741823) 2:(1,1073741823) 3:(1,100) 4:(101,1073741823)
1:(-1073741824,1073741823) 2:(1,101) 3:(1,100) 4:(101,1073741823)
1:(-1073741824,1073741823) 2:(1,101) 3:(1,100) 4:(101,101)
converged to lfp.
1:(-1073741824,1073741823) 2:(1,101) 3:(1,100) 4:(101,101)
0.000u 0.000s 0:00.00 0.0%      0+0k 0+0io 0pf+0w
%
```

Widening/Narrowing are not duals

	Iteration starts from	Iteration stabilizes
Widening ∇	below	above
Narrowing Δ	above	above
Dual widening $\tilde{\nabla}$	above	below
Dual narrowing $\tilde{\Delta}$	below	below

No dual widening $\tilde{\nabla}$ has ever been found but trivial ones such as bounded execution (bounded model-checking), execution on a few cases (debugging), etc.

On actual abstract interpreters

- For simplicity, we have designed a specific abstract interpreter for a specific program.
- In practice, abstract interpreters are parameterized by the program they have to analyze, and by the abstraction which should be used for the analysis.
- Observe that the code defining the transformer could be directly generated from the program text and so we have a model of an abstract compiler.
- (An alternative would use a computer representation of the equations and an abstract interpreter would be used to evaluate the transformer by calls to the interval abstract domain). □

Chaotic Iterations: A Structural Instance (optional)

Chaotic iterations

We will see in section 6.3 that the iteration of the abstract equations need not follow the Jacobi iteration strategy and can be done in any chaotic order provided no equation is forgotten forever (or equivalently every equation is evaluated infinitely often) until it is stabilized.

A particular instance of such an efficient chaotic iteration follows program execution as defined by induction on its syntax (see chapter 13).

Starting from the entry condition at program point ¹, we can stabilize the loop ^{2—3} before computing the invariant at program point ⁴.

We define

Structural iterations

```
(* structural reachability analysis with widening and
   narrowing *)
open Interval
open IntervalWidening
open IntervalNarrowing
open Invariant
open TransformerBounded
open Iterator
let analyzer () =
  let p1 = f1 () in
  let p2 = let f x2 = f2 p1 (f3 x2) in
          let fw x2 = widen x2 (f x2) in
          let w = (lfp less EMPTY fw) in
          let fn x2 = narrow x2 (f x2) in
          (lfp greater w fn) in
  let p3 = f3 p2 in
  let p4 = f4 p2 in
  pprint (p1, p2, p3, p4);;
  analyzer ();;
```

Structural iterations (cont'd)

and get exactly the same global result (the trace shows the iteration with widening and then the iteration with narrowing for the loop²—³)

```
% ocamlc interval.ml intervalWidening.ml intervalNarrowing.ml \
? invariant.ml invariantWidening.ml invariantNarrowing.ml \
? transformerBounded.ml iteratorTrace.ml \
? structural_reachability_narrowing_bounded_trace.ml
% time ./a.out
_1_ (1,1) (1,1073741823) converged to fixpoint.
(1,1073741823) (1,101) converged to fixpoint.
1:(-1073741824,1073741823) 2:(1,101) 3:(1,100) 4:(101,101)
0.000u 0.000s 0:00.00 0.0%      0+0k 0+0io 0pf+0w
%
```

Verification

Verifier

- The abstract interpreter that we have designed is a sound *static analyzer*. Given a program it produces interval information always valid at runtime.
- We can turn it into a *Verifier* checking an interval specification.
- The specification can be provided by the user or remain implicit (e.g. absence of runtime errors such as overflows).
- One kind of user specification is a type declaration, for example an interval declaration for integer variables like `var x : 1..100;`.
- Let us understand this declaration as: “only values between 1 and 100 can be assigned to x, otherwise execution stops” (with a runtime error).
- Observe that this does not mean that x always has a value between 1 and 100 because it can be initialized with any integer value.¹³.

¹³This interpretation of the interval declaration is that of the PASCAL programming language, see K. Jensen and N. Wirth: Pascal User Manual and Report, Second Edition, Springer, 1975.

Example of Interval Verification

For the following example

$$P' \triangleq \text{var } x : 1..100 ; ^1x := 1 ; \text{while } ^2(x \leq 100) \text{ do } ^3x := (x + 1); \text{ od}^4 .$$

the abstract interval equations become

$$\begin{cases} X_1 &= \{x \leftarrow [\min_int, \max_int]\} \\ X_2 &= \{x \leftarrow ([1, 1] \sqcup (\exists X_3(x) = \emptyset \Rightarrow \emptyset) : \text{let } [a, b] = X_3(x) \text{ in } [\min(a + 1, \max_int), \min(b + 1, \max_int)]\}) \sqcap [1, 100]\} \\ X_3 &= X_2 \dot{\sqcap} \{x \leftarrow [\min_int, 100]\} \\ X_4 &= X_2 \dot{\sqcap} \{x \leftarrow [101, \max_int]\} \end{cases}$$

since execution stops if and when a value outside [1, 100] is going to be assigned to x.

Encoding the Declaration

This declaration is encoded in OCAML as follows

```
(* declaration.ml *)
open Interval
open Invariant
let d =
  (INT (min_int,max_int),
   INT (1,100),
   INT (min_int,max_int),
   INT (min_int,max_int));;
```

Software Verification, ETH Zurich, Switzerland, 25 November 2009

97

© P. Cousot, <http://se.inf.ethz.ch/teaching/2009-H/c-0239/index.html#slides>

Encoding the Verifier

The abstract interpreter performs the iterative abstract reachability fixpoint overapproximation with widening/narrowing and intersection with the declaration, then prints the least fixpoint result, and finally checks for errors.

```
(* reachability verification with widening and narrowing *)
open Invariant
open InvariantWidening
open InvariantNarrowing
open TransformerBounded
open Iterator
open Declaration
open Verifier
let verifier () =
  let fw x = (pmeet (pwiden x (f x)) d) in
  let w = (lfp pless pbot fw) in
    let fn x = pnarrow x (f x) in
      let a = (lfp pgreater w fn) in
        pprint a; pverify cless f a d;;
verifier ();;
```

Software Verification, ETH Zurich, Switzerland, 25 November 2009

99

© P. Cousot, <http://se.inf.ethz.ch/teaching/2009-H/c-0239/index.html#slides>

Encoding the Verification Phase

The verification of absence of errors checks that at any point during an execution without error up to some point in the computation will not have an error at the next execution step.

```
(* verifier.ml, interval invariant abstract domain *)
let pwarning (b1, b2, b3, b4) =
  let m = "Potential error at line " in
  if not b1 then print_string (m^"1\n");
  if not b2 then print_string (m^"2\n");
  if not b3 then print_string (m^"3\n");
  if not b4 then print_string (m^"4\n");
let pverify leq f a d =
  let b = leq (f a) d in
    pwarning b;
```

Software Verification, ETH Zurich, Switzerland, 25 November 2009

98

© P. Cousot, <http://se.inf.ethz.ch/teaching/2009-H/c-0239/index.html#slides>

Result of the Analysis

```
% ocamlc interval.ml intervalWidening.ml intervalNarrowing.ml \
? invariant.ml invariantWidening.ml invariantNarrowing.ml \
? transformerBounded.ml iterator.ml declaration.ml \
? verifier.ml reachability_narrowing_declaration.ml
% time ./a.out
1:(-1073741824,1073741823) 2:(1,100) 3:(1,100) 4:_-
Potential error at line 2
0.000u 0.000s 0:00.00 0.0%          0+0k 0+0io 0pf+0w
%
```

- Observe that the program execution always stops at program point ³ with an overflow outside the range [1, 100] so program point ⁴ is now unreachable (with an overapproximation we can prove the presence of dead code but not its absence).
- Notice that the error is signaled as potential (with an overapproximation we can prove the values to definitely be within given bounds but not to prove that execution ever assigns a given value to a variable). Here is a trace of the analysis.

Software Verification, ETH Zurich, Switzerland, 25 November 2009

100

© P. Cousot, <http://se.inf.ethz.ch/teaching/2009-H/c-0239/index.html#slides>

Details of the Analysis

```
% ocamlc interval.ml intervalWidening.ml intervalNarrowing.ml \
? invariant.ml invariantWidening.ml invariantNarrowing.ml \
? transformerBounded.ml iteratorTrace.ml declaration.ml \
? verifier.ml reachability_narrowing_declaration_trace.ml
% time ./a.out
1:_|_ 2:_|_ 3:_|_ 4:_|_
1:(-1073741824,1073741823) 2:(1,1) 3:_|_ 4:_|_
1:(-1073741824,1073741823) 2:(1,1) 3:(1,1) 4:_|_
1:(-1073741824,1073741823) 2:(1,100) 3:(1,1) 4:_|_
1:(-1073741824,1073741823) 2:(1,100) 3:(1,100) 4:_|_
converged to lfp.
1:(-1073741824,1073741823) 2:(1,100) 3:(1,100) 4:_|_
converged to lfp.
1:(-1073741824,1073741823) 2:(1,100) 3:(1,100) 4:_|_
Potential error at line 2
0.000u 0.000s 0:00.00 0.0%      0+0k 0+0io 0pf+0w
%
```

101

© P. Cousot, <http://se.inf.ethz.ch/teaching/2009-H/c-0239/index.html#slides>

Correcting the Declaration

```
(* declarationCorrect.ml *)
open Interval
open Invariant
let d =
  (INT (min_int,max_int),
   INT (1,101),
   INT (min_int,max_int),
   INT (min_int,max_int));;
```

yields no error, the verification is completed.

```
% ocamlc interval.ml intervalWidening.ml intervalNarrowing.ml \
? invariant.ml invariantWidening.ml invariantNarrowing.ml \
? transformerBounded.ml iterator.ml declarationCorrect.ml \
? verifier.ml reachability_narrowing_declaration_correct.ml
% time ./a.out
1:(-1073741824,1073741823) 2:(1,101) 3:(1,100) 4:(101,101)
0.000u 0.000s 0:00.00 0.0%      0+0k 0+0io 0pf+0w
%
```

102

© P. Cousot, <http://se.inf.ethz.ch/teaching/2009-H/c-0239/index.html#slides>

When to do the verification?

Notice that in general the verification cannot be done during the analysis since a widening may cause an overapproximation potentially raising a potential error while the narrowing may refine the analysis well enough to that this potential error disappears.

103

© P. Cousot, <http://se.inf.ethz.ch/teaching/2009-H/c-0239/index.html#slides>

Varieties of abstract interpreters

(optional)

104

© P. Cousot, <http://se.inf.ethz.ch/teaching/2009-H/c-0239/index.html#slides>

Abstract interpreters

An *abstract interpreter* computes an abstract semantics or an approximation of this abstract semantics for a language.

- Abstract interpreters are parameterized by abstract domains which composition specifies the abstraction.
- So the abstraction can be changed, in particular refined without redesigning the whole abstract interpreter.
- There are many different possible styles of abstract interpreters¹.

¹We exclude here debuggers and simulators which involve executing the program and restrict ourselves to static analysis and verification methods involving no execution of the program, in which case we would use the dynamic qualifier.

Analyzers/Verifiers/Bug finders

— A first distinction come from the use of the abstract semantics and its soundness.

Analyzers: the objective is to automatically and statically determine abstract information about the program dynamic/runtime concrete behavior. A typical use of this information is in program transformation or compiler optimization, in which case loss of precision means loss of performance but not of correctness which is often preferred to high analysis costs.

Verifiers: The automatically and statically determined information about the program dynamic/runtime concrete behavior is used to prove a program specification. By definition verifiers are sound² (but in general incomplete). They provide presumptions of presence as well as guarantees of absence of some categories of bugs with respect to a class of specifications (i.e. true or false positives and true but never false negatives so that an error may be reported in the abstract while absent in the concrete while absence of error in the abstract implies absence of error in the concrete).

Bug-finding checkers: Bug-finders are both unsound and incomplete. They provide presumptions of presence of some categories of bugs (i.e. both true or false positives and negatives are possible since not all concrete errors are reported in the abstract while reported errors might be spurious).

²Although the word is sometimes clichéd and hackneyed through overuse for provably unsound tools. For example checking liveness on a program model may be unsound for the concrete semantics.

Scope of abstract interpreters

— A second distinction is between the scope of applicability of the analyzer/verifier/checker.

Program specific interpreters: The *program specific interpreters* are built for a given program P . For example a specific model of the program execution is built manually and this model is then checked in this abstract with respect to given specification.

Language-specific interpreters: The *language-specific interpreters* are built for all programs $P \in \mathcal{P}$ of a given language \mathcal{P} .

Domain-specific interpreters: The *domain-specific interpreters* are language-specific interpreters specifically built for a given infinite family of programs in the language \mathcal{P} (e.g. corresponding to a specific domain of application).

Language-independant interpreters: The *language-independant interpreters* are built for all languages, which generally means for a given family of languages which program abstract semantics are transformed into a common representation of their abstract semantics.

Modularity

— A similar third distinction is between analyzable/verifiable/checkable units.

Global interpreters: Global interpreters can analyze/verify/check only executable programs, as a whole (may be with stubbed libraries). So a program that is not compilable, linkable and executable cannot be analyzed in general because its concrete semantics would be undefined.

Local/modular interpreters: Local/modular interpreters can analyze/verify/check program parts (may be with hypotheses on the execution environment of the part). So a program part may be analyzed even if it does not execute or even compile.

Specifications

— A forth distinction is between the styles of specifications of program properties to be analyzed, verified or checked. We assume that this specification includes the hypothesis to be done on the runtime environment, if any.

Implicit specification: The specification to be verified is defined once for all for the language (e.g. absence of runtime errors for a given language which verification conditions can be generated automatically for each program in the language);

Explicit specification: The specification to be verified is provided by the user using a specification language.

Mixed specification: Part of the specification is provided by the user and part automatically. For example, the user can specify hypothesis on the execution environment while the property to be checked is generated automatically from the program text.

109

© P. Cousot, <http://se.inf.ethz.ch/teaching/2009-Htc-0239/index.html#slides>

Abstract semantics

— A fifth distinction is between the styles of abstract semantics of programs.

Transitional interpreters: The *transitional abstract interpreters* are based on the transitional semantics of the language and so iterate the application of transformers for individual program execution steps until global stabilization.

Structural interpreters: The *structural abstract interpreters* are based on the structural semantics of the language and so proceed by induction on the program syntax with iteration of the transformers of the body of loops until local stabilization.

110

© P. Cousot, <http://se.inf.ethz.ch/teaching/2009-Htc-0239/index.html#slides>

User interaction

— A sixth difference is on the user interaction requirements.

Automatic interpreters: Once given a program and a specification, the abstract interpreter works completely automatically, without any need (and possibility) of user intervention.

Interactive interpreters: The execution of the interpreter is in interaction with the user e.g. to provide inductive arguments (such as inductive invariants when their fixpoint definition is not amenable to automatization) or to help in guiding proofs.

Automatic controllable interpreters: The execution of the interpreter is automatic but can depend on offline parameters or directives to help taking decisions during the online analysis (as opposed to automatic interpreters never depending on the user understanding of the abstraction and interactive interpreters where users can observe the déroulement of the analysis/verification/check online).

111

© P. Cousot, <http://se.inf.ethz.ch/teaching/2009-Htc-0239/index.html#slides>

Convergence

— A seventh difference is between the choice of convergence.

Terminating interpreters: always terminate with finite resources (but may be an imprecise result)

Finitary interpreters: termination is due to the choice of a finitary abstraction (e.g. finite abstract domain or increasing iteration with ACC which involve an offline³ loss of information, at design time);

Infinitary interpreters: termination is obtained by enforcing convergence of the iterates through extrapolation which involves an online⁴ loss of information, at analysis time);

Non-terminating interpreters: have the possibility to never terminate for some programs or specifications or terminate by memory overflow or time out.

³Offline is sometimes called static meaning before the analysis or verification, which can be confusing with static meaning at compile time, before program execution.

⁴Online is sometimes called dynamic meaning during the analysis or verification, which can be confusing with dynamic meaning at runtime, during program execution.

112

© P. Cousot, <http://se.inf.ethz.ch/teaching/2009-Htc-0239/index.html#slides>

Extensibility

- A eighth difference is between the choice of the abstraction.

Fixed interpreters: the abstraction used by the interpreter is fixed and cannot be changed without redesigning the abstract interpreter. This is typically the case of dataflow analyses embedded in an optimizing compiler.

Extensible interpreters: the abstraction used by the interpreter can be modified without redesigning the whole interpreter e.g. by combining abstract domains.

Refinement consistency

- A tenth distinction concerns the refinement process.

Increasing interpreters: Increasing abstract interpreters ensure that a refinement of the abstraction will always produce a more precise result.

Non-increasing interpreters: Non-increasing abstract interpreters have the property that a refinement of the abstraction may sometime yield less precise results.

Refinement-time

- A ninth difference is between the choice of the refinement of the abstraction (e.g. in case of false alarm).

Offline refining interpreters: The abstraction of the concrete semantics is chosen before the computation of the abstract semantics so that the abstraction is fixed during the analysis phase. The refinement must therefore be done offline by changing the abstract domains.

Online refining interpreters: The abstraction of the concrete semantics is chosen during the computation of the abstract semantics so that the abstraction can be refined online, during the analysis phase.

Encoding of the transformer

- A eleventh distinction concerns the encoding and computation of the abstract transitional or structural semantics.

Interpretation: In a first phase, the abstract semantics are encoded in an intermediate equational, constraint, etc. form using some intermediate language. The encoding is generally in the form of a term encoded in some data structure. In a second phase, an interpreter designed for this intermediate language is executed to compute the abstract semantics iteratively.

Compilation: In a first phase, the abstract semantics is compiled into executable code written in some programming language and including an iterator. In a second phase, the execution of the compiled program directly yields the abstract semantics.

Abstract transformer computation time

— A twelfth distinction involves the abstract transformers.

Offline abstract transformers: The *offline abstract transformers* are computed during the first phase while generating the intermediate encoding or code (so that $\alpha \circ f \circ \gamma$ is encoded directly into some \bar{f} which evaluation in the second phase does not involve that of α , f , or γ);

Online abstract transformers: The *online abstract transformers* are encoded during the first phase but effectively computed during the second interpretation or execution phase (for example $\alpha \circ f \circ \gamma(x)$ is computed for a given parameter x e.g. using a theorem prover to discover a minimal y such that $\alpha \circ f \circ \gamma(x) \sqsubseteq y$ or directly computing $\alpha \circ f \circ \gamma(x)$ by referring to the concrete domain);

Mixed abstract transformers: The abstract transformers are computed partly during the first generation phase and partly during the second interpretation phase. (e.g. $\alpha_2 \circ \alpha_1 \circ f \circ \gamma_1 \circ \gamma_2(x)$ where $\bar{f} = \alpha_1 \circ f \circ \gamma_1$ is computed in the first phase while $\alpha_2 \circ \bar{f} \circ \gamma_2(x)$ is computed in the second).

117

© P. Cousot, <http://se.inf.ethz.ch/teaching/2009-H/c-0239/index.html#slides>

Software Verification, ETH Zurich, Switzerland, 25 November 2009

Cost/precision ratio

— A thirteenth distinction involves the cost/precision ratio.

Precise interpreters: *precise interpreters* single out precision of the analysis to the detriment of the cost, typically hours of computations per hundreds of thousands of program lines.

Rapid interpreters: *rapid interpreters* favour fast responses (typically seconds or minutes of computations per hundreds of thousands of program lines) often to the detriment of their precision

Of course precise and rapid abstract interpreters are an hardly reachable goal.

118

© P. Cousot, <http://se.inf.ethz.ch/teaching/2009-H/c-0239/index.html#slides>

Software Verification, ETH Zurich, Switzerland, 25 November 2009

Error reporting

— A fourteenth distinction involves error reporting (in particular for rapid verifiers often subject to many false alarms).

Total error reporting: All potential errors found are reported.

Partial error reporting: Only some of the potential errors are reported (e.g. only the most common ones to avoid discouraging the user in case of imprecise or unsound analysis).

119

© P. Cousot, <http://se.inf.ethz.ch/teaching/2009-H/c-0239/index.html#slides>

Software Verification, ETH Zurich, Switzerland, 25 November 2009

etc

And the main difference is certainly on the choice of the abstractions!

120

© P. Cousot, <http://se.inf.ethz.ch/teaching/2009-H/c-0239/index.html#slides>

Software Verification, ETH Zurich, Switzerland, 25 November 2009

Conclusion & references

121

Software Verification, ETH Zurich, Switzerland, 25 November 2009

© P. Cousot, <http://se.inf.ethz.ch/teaching/2009-Htc-0239/index.html#slides>

Conclusion

- The presentation relied purely on intuition, can be made formal (see references)
- The abstraction ideas can **scale up** with enough precision, e.g.
 - **ASTRÉE:**
 - <http://www.astree.ens.fr/>
 - <http://www.absint.de/astree/>
 - **Clousot:**
 - MSR, Redmond

122

Software Verification, ETH Zurich, Switzerland, 25 November 2009

Bibliography

The very first report on static analysis in infinite abstract domains not satisfying the ACC with widening/narrowing (Cousot and Cousot, 1975) was published in (Cousot and Cousot, 1976). The most cited reference is (Cousot and Cousot, 1977a). It is extended in (Cousot and Cousot, 1977b; Cousot, 1978) to handle procedures, see also (Bourdoncle, 1993). (Cousot, 1978) contains a presentation of reachability analysis using transition systems (i.e. language independent semantics and equational analyzers) later published in (Cousot, 1981).

An online introduction (in French) : <http://www.di.ens.fr/~cousot/COUSOTtalks/CollegeDeFrance08.shtml>

An online course : <http://web.mit.edu/afs/athena.mit.edu/course/16/16.399/www/>

123

Software Verification, ETH Zurich, Switzerland, 25 November 2009

© P. Cousot, <http://se.inf.ethz.ch/teaching/2009-Htc-0239/index.html#slides>

References

- Cousot, P. and Cousot, R. (1976). Static determination of dynamic properties of programs, *Proceedings of the Second International Symposium on Programming*, Dunod, Paris, Paris, pp. 106–130.
- Cousot, P. and Cousot, R. (1977a). Abstract interpretation: a unified lattice model for static analysis of programs by construction or approximation of fixpoints, *Conference Record of the Fourth Annual ACM SIGPLAN-SIGACT Symposium on Principles of Programming Languages*, ACM Press, New York, Los Angeles, pp. 238–252.
- Cousot, P. and Cousot, R. (1977b). Static determination of dynamic properties of recursive procedures, in E. Neuhold (ed.), *IFIP Conference on Formal Description of Programming Concepts*, St-Andrews, N.B., North-Holland Pub. Co., Amsterdam, pp. 237–277.

124

§

88

Cousot, P. and Cousot, R. (1992). Comparing the Galois connection and widening/narrowing approaches to abstract interpretation, invited paper, in M. Bruynooghe and M. Wirsing (eds), *Proceedings of the Fourth International Symposium on Programming Language Implementation and Logic Programming, PLILP'92*, Leuven, 26–28 August 1992, Lecture Notes in Computer Science 631, Springer, Berlin, pp. 269–295.

Leroy, X., Doligez, D., Garrigue, J., Rémy, D. and Vouillon, J. (2009). The Objective Caml system release 3.11, documentation and user's manual, *Technical report*, INRIA, Rocquencourt, France. <http://caml.inria.fr/pub/docs/manual-ocaml/>.